



“Biais politiques dans les LLMs: comprendre, analyser et atténuer”



Akram Elbouanani

PhD student @ CEA



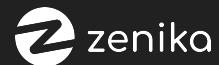
23 Septembre 2025 à 19h15



4 rue Voltaire, 44100 Nantes

chez

icilundi



[sfɛir]

lonestone

N M L M

Nantes Machine Learning

Meetup

December 4th : Entity alignment in Knowledge Graph using LLMs

techstars_ **Startup Weekend** **Nantes**

Du 21/11 au 23/11
Le Palace _jclundi



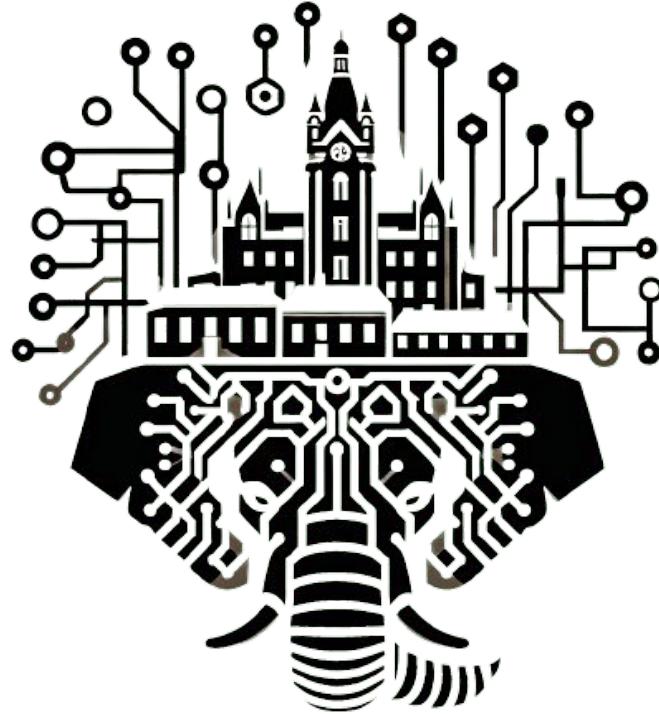
Réserve ta place !



Qui n'est jamais venu au meetup Gen AI Nantes ?

GenAI Nantes

- 15 événements / an
- 1 hackathon (Shift)
- 1 workshops
- 1 communauté de 900p*



* 9.000p selon le syndicat des llamas 

Qui vient pour la première fois? 🏅

Qui est tech ?



Qui est MLE ?



Qui code assisté par une IA ? 

Qui code dans un fork de VSCode ? 

(Cursor, Windsurf...)

Qui utilise un agent AI dans le terminal ? 

(Claude Code, Gemini Cli, Codex...)

**Qui utilise un agent AI dans une extension
de l'IDE ?** 

(Cline, Copilot, Roo code, Aider, Continue...)

Qui recrute dans la GenAI ? 

Qui cherche un job dans la GenAI ?



Qui souhaite s'associer dans la GenAI ? 

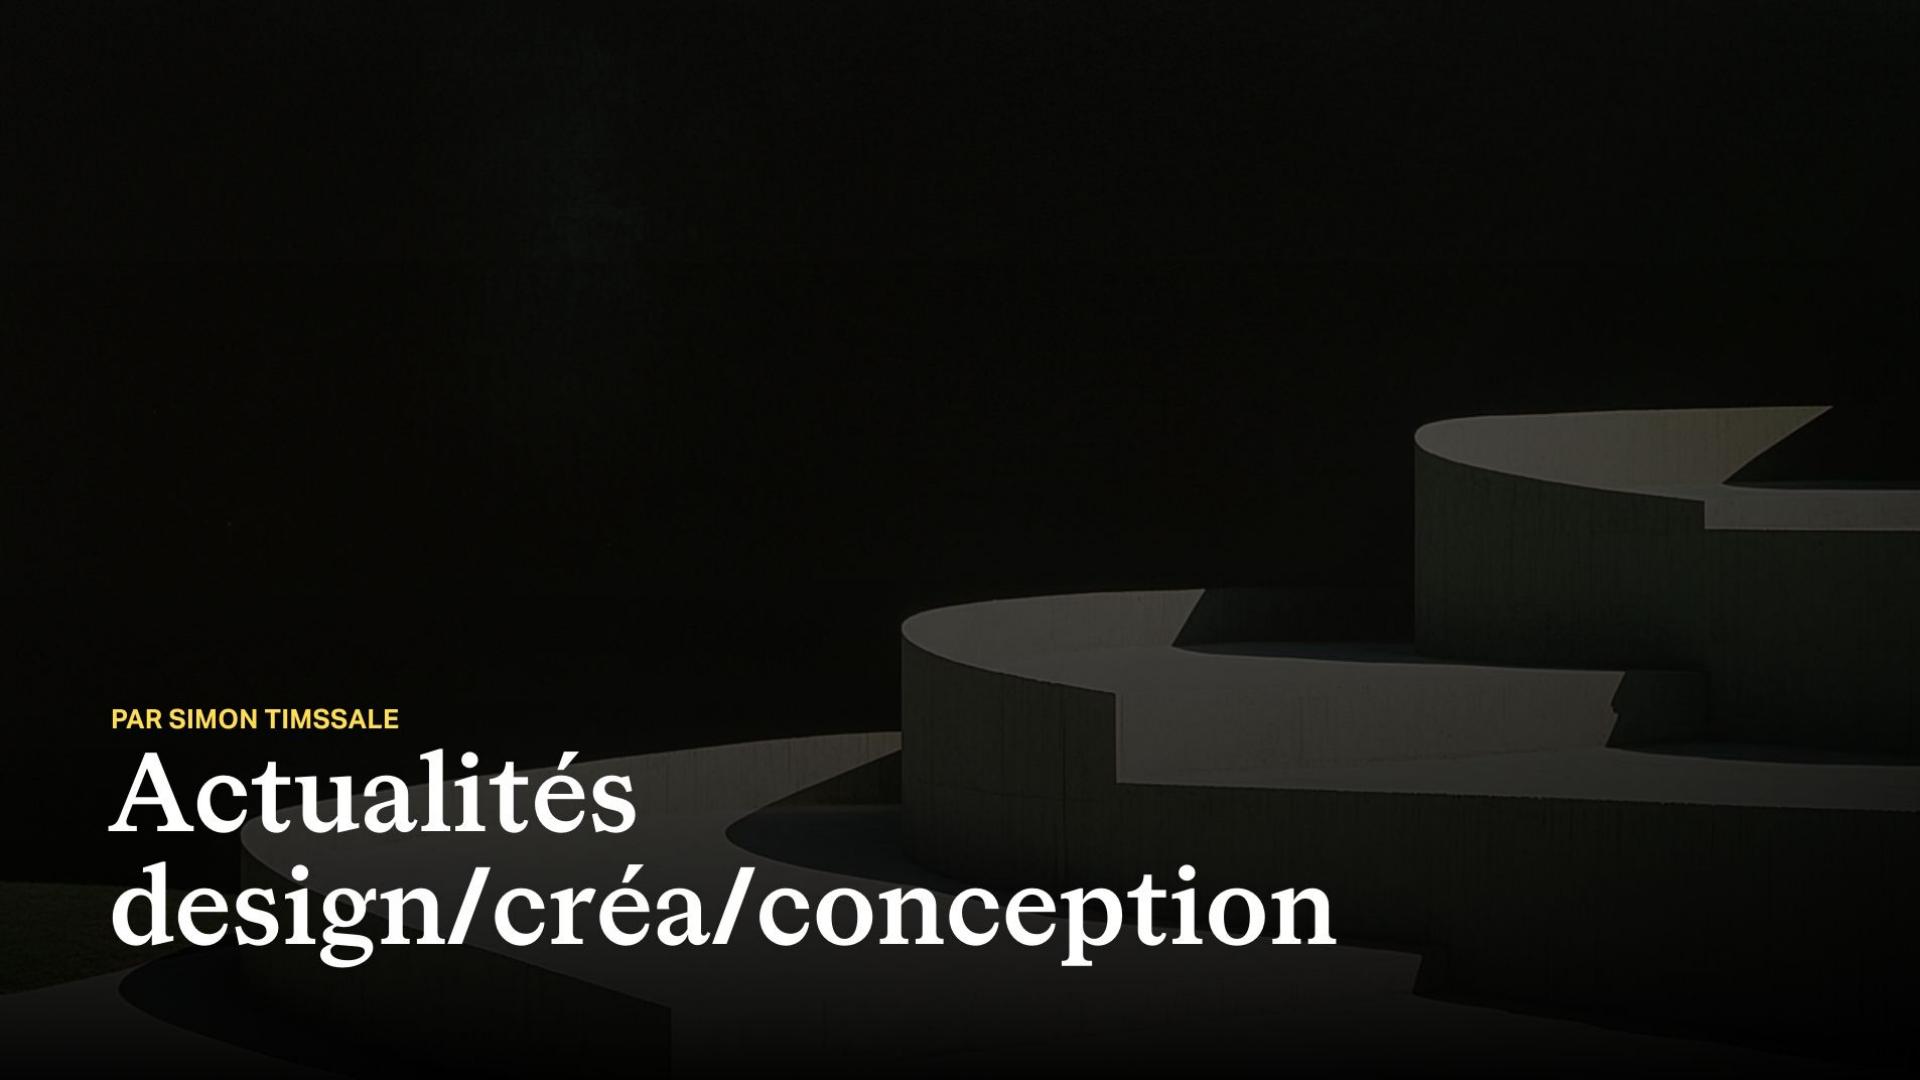
Schedule

👀 1- News

🔍 2- Biais dans les LLMs

🍺 3- Enjoy

News créa



PAR SIMON TIMSSALE

Actualités design/créa/conception

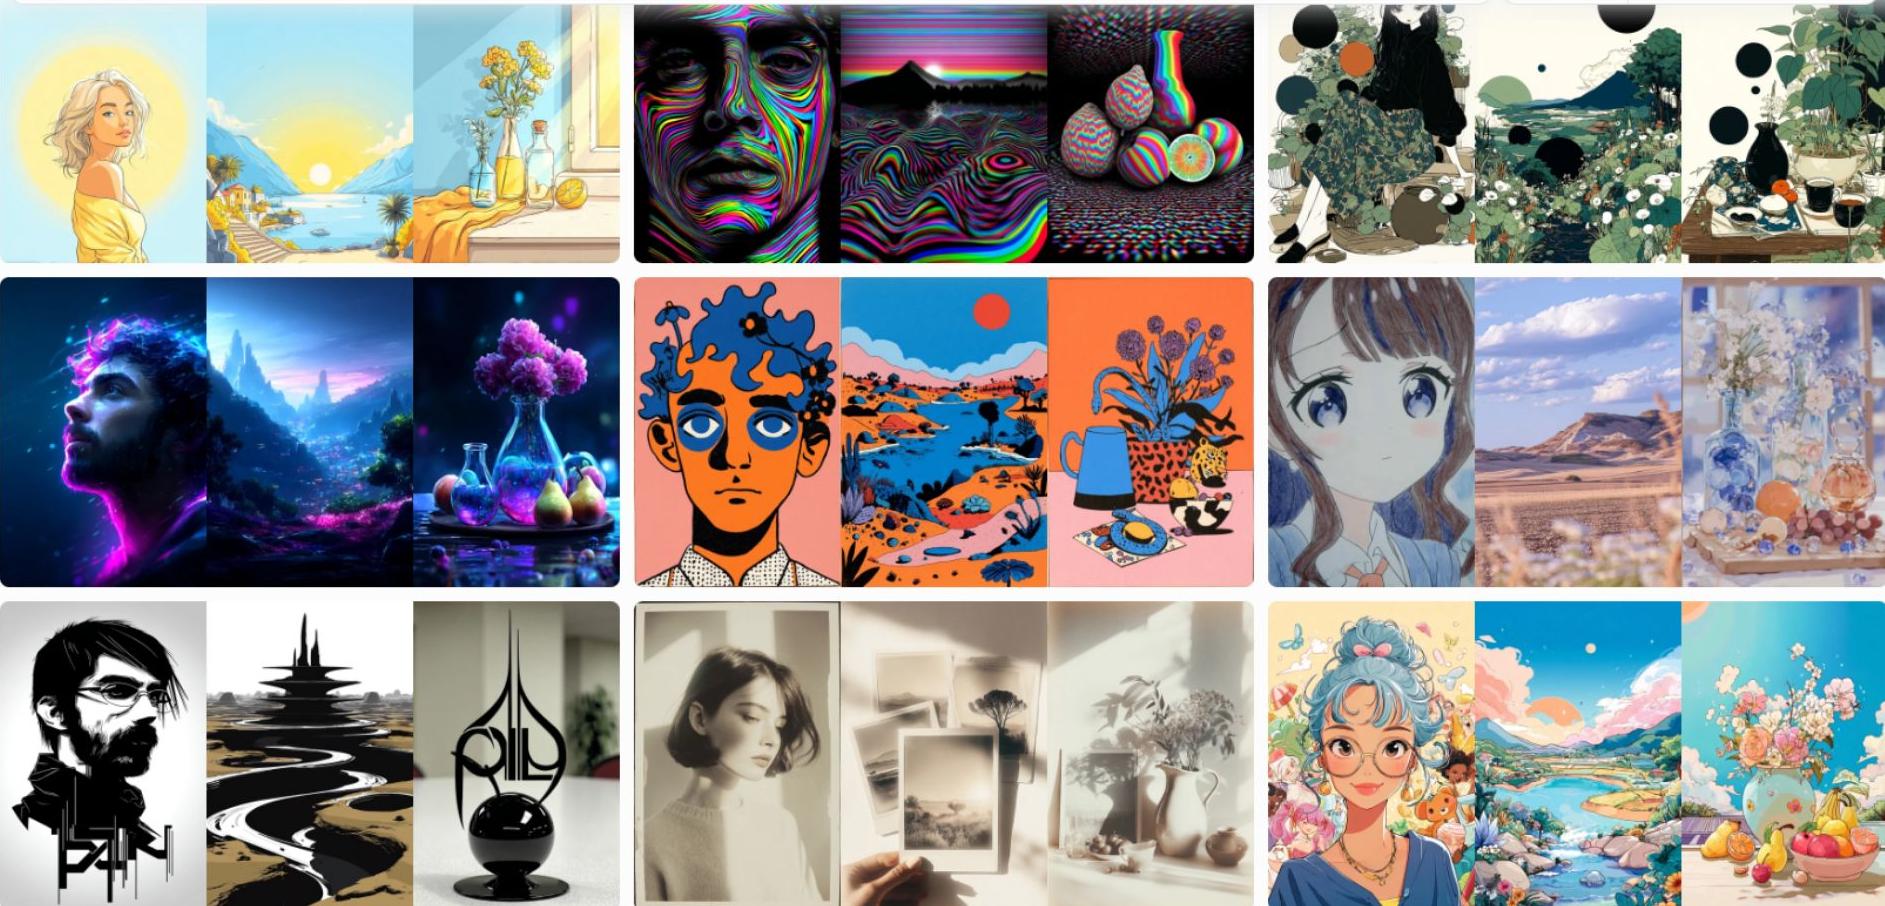


What will you imagine?



P ↻ ⚡

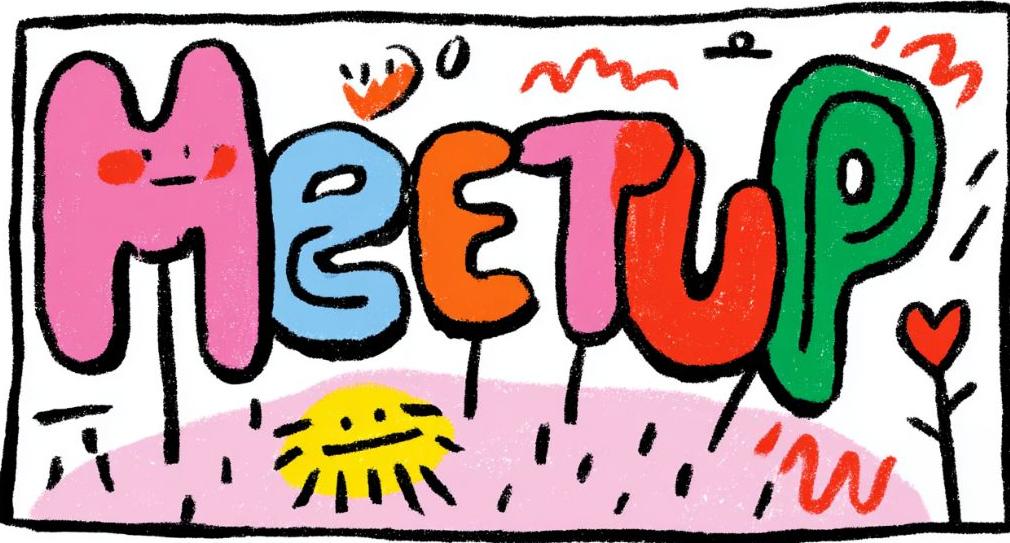
Search





01
MIDJOURNEY





"Meetup" logo written in bold font on a white background

nombre de styles initial x 2 x 7



GEMINI-FLASH IMAGE 2.5

Nano-banana

02

NANO-BANANA

Google AI Studio

Chat prompt

Run settings Get code

Nano Banana
gemini-2.5-flash-image-preview
(aka Gemini 2.5 Flash Image) State-of-the-art image generation and editing model.

System instructions
Optional tone and style instructions for the model

Temperature

Advanced settings

Show me different logos and brand swag ideas for my startup called Avurna →
Run Ctrl ↵

What's new

- Try Nano Banana
Gemini 2.5 Flash Image, state of the art image generation and editing
- Fetch information with URL context
Fetch real-time information from web links
- Generate native speech with Gemini
Generate high quality text to speech with Gemini
- Talk to Gemini live
Try Gemini's natural, real-time dialog experience, with audio and video inputs

This model is not stable and may not be suitable for production use. [Learn more](#).

Get API key

Settings

no prompting toussa

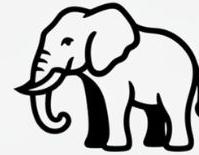
User



Crée une image avec ce produit posé sur une table dans une cuisine traditionnelle italienne. Un fromage posé à coté. Un visuel qui fasse pub.

Model









Modifie cette image en améliorant le parquet au sol et en le rendant neuf (même parquet mais dans un état neuf)

Model

Bien sûr, voici l'image avec un parquet neuf et amélioré :





Modifie cette image en améliorant le parquet au sol et en le rendant neuf (même parquet mais dans un état neuf)

Model

Bien sûr, voici l'image avec un parquet neuf et amélioré :



User



Déplace le canapé bleu vers la droite. Il doit être dos à la fenêtre de droite et faire face à fenêtre de gauche.

Model





Ajoute ce tapis au sol en face du canapé. Ce tapis gradient va du canapé à droite jusqu'au même niveau à gauche de l'appartement. Il s'étend jusqu'à l'extrémité gauche de la fenêtre d'en face mais le tapis reste devant le canapé bleu.

Model













tonifier
légèrement le
tronc des arbres.

remplir la piscine
et renover le bord.

tonifier
légèrement le
tronc des arbres.

l'herbe verte avec
des pâquerettes.



02

NANO-BANANA



02

NANO-BANANA

Limites professionnelles

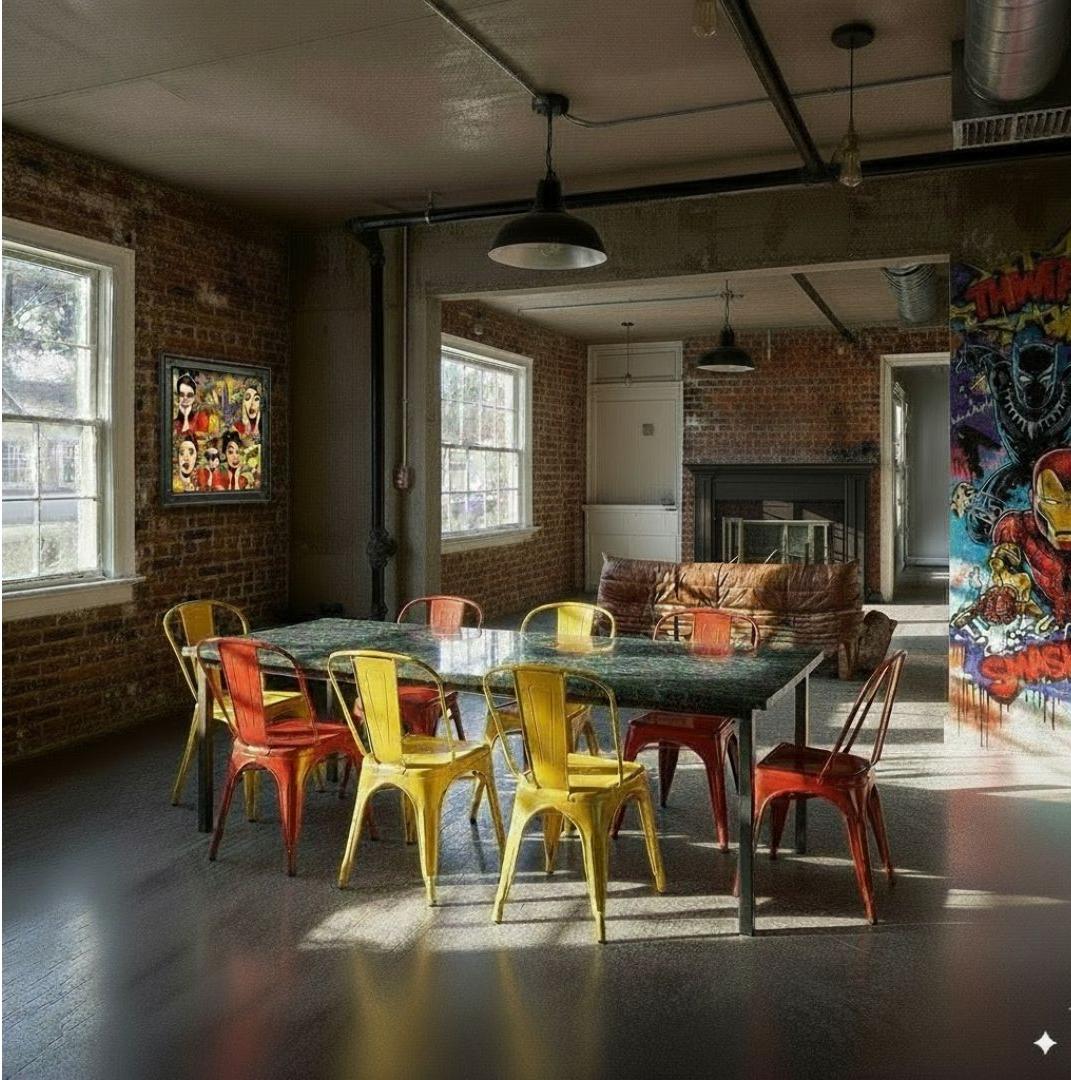
Conflits IA & droits

Résolution des visuels

03

PROCESS PRO





Seedream 4.0

Nano-banana

03

PROCESS PRO

Weavy Artistic Intelligence

Turn your creative vision into scalable workflows.
Access all AI models and professional editing tools
in one node based platform.

3D RODIN 2.0



COLOR REFERENCE



IMAGE STABLE DIFFUSION



TEXT

a Great-Tailed Grackle bird is flying from the background and seating on the model's shoulder slowly and barely moves, the model looks at the camera, then bird flies away, cinematic.

IMAGE FLUX PRO 1.1



VIDEO MINIMAX VIDEO



Usage > waouh effect

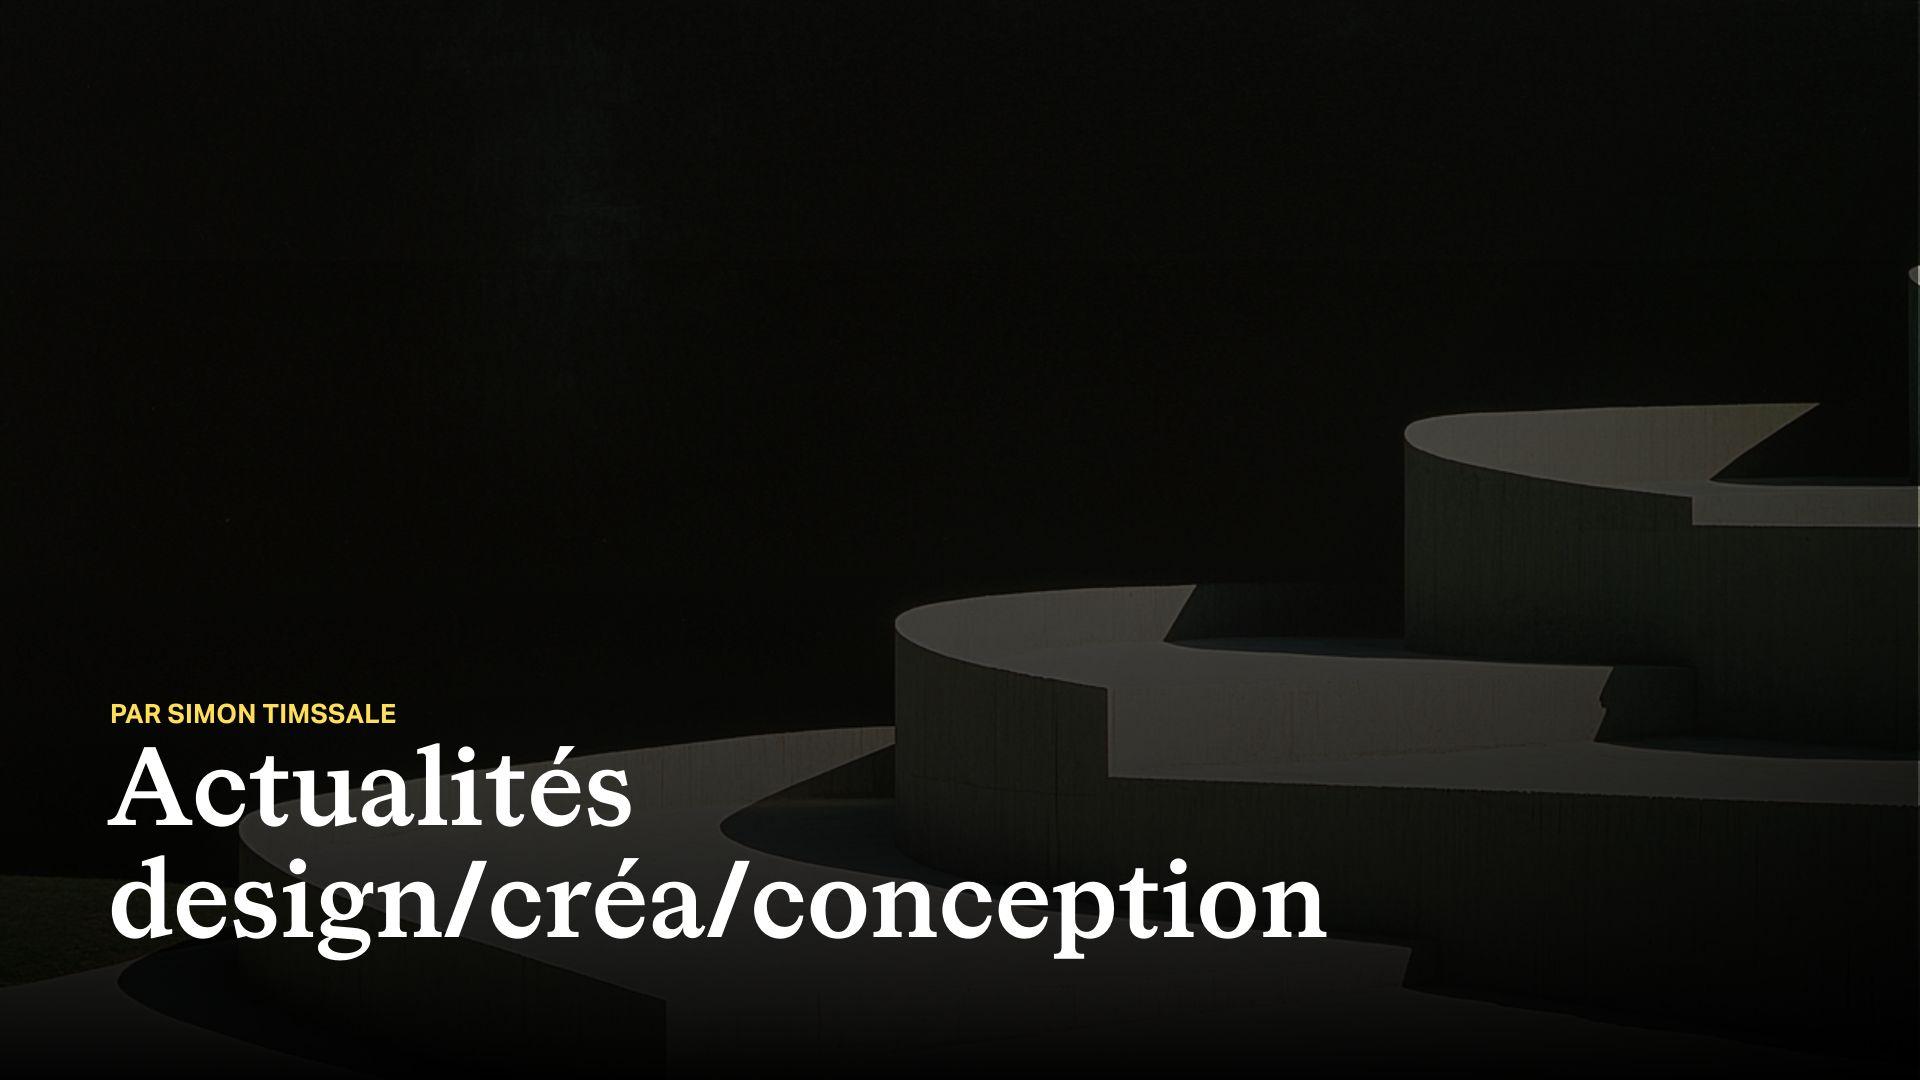


style : Midjourney

modèles perso : Seelab ou Freepik

correction : Nano-banana ou Seedream 4.0

upscaling et retouches : Seelab ou Freepik



PAR SIMON TIMSSALE

Actualités design/créa/conception

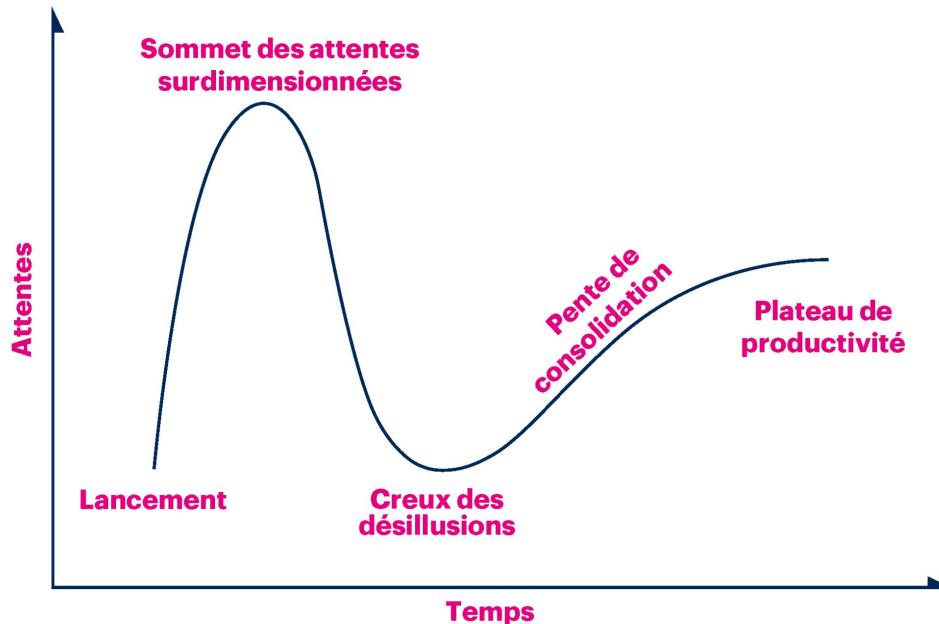
News non-tech

Tour d'horizon

- Modèles
- Outils créatifs
- Ecosystème
- Impact sociaux

GPT5

On nous a promis la singularité, on a eu une mise à jour.



La hype plafonne, le progrès continue

- **OpenAI GPT-5** (August 7, 2025) - Multimodal model with 1M+ token context, 30% reduction in hallucinations, unified reasoning architecture, 94.6% AIME 2025 score
- **OpenAI gpt-oss-120b and gpt-oss-20b** (August 4, 2025) - First open-weight models from OpenAI since GPT-2, MoE architecture optimized for edge deployment
- **Anthropic Claude Opus 4.1** (August 5, 2025) - 74.5% SWE-bench accuracy, 200k context, hybrid instant/extended reasoning modes
- **xAI Grok 4** (July 2025) - Perfect 100% score on AIME 2025 mathematical reasoning, 256k context window
- **Meta Llama 4 Scout** - 10M token context, 17B active parameters, fits on single H100 GPU
- **Meta Llama 4 Maverick** - 1M context, beats GPT-4o and Gemini 2.0 Flash on benchmarks
- **Alibaba Qwen3-Next** - Revolutionary 80B/3B hybrid architecture with 90% training cost reduction
- **Alibaba Qwen 3 Max** (September 2025) - 80.6% AIME 2025, #6 on LMarena leaderboard
- **ByteDance Seed-OSS-36B-Instruct** (August 21, 2025) - Mobile-optimized open-source model
- **DeepSeek R1 and R1-Zero** - Open-sourced reasoning models comparable to o1 at fraction of cost
- **Mistral Magistral Small and Medium** (June 10, 2025) - First reasoning models from Mistral
- **Google Jules Coding Agent** (August 2025) - Autonomous asynchronous coding agent with GitHub integration



Modèles - En bref

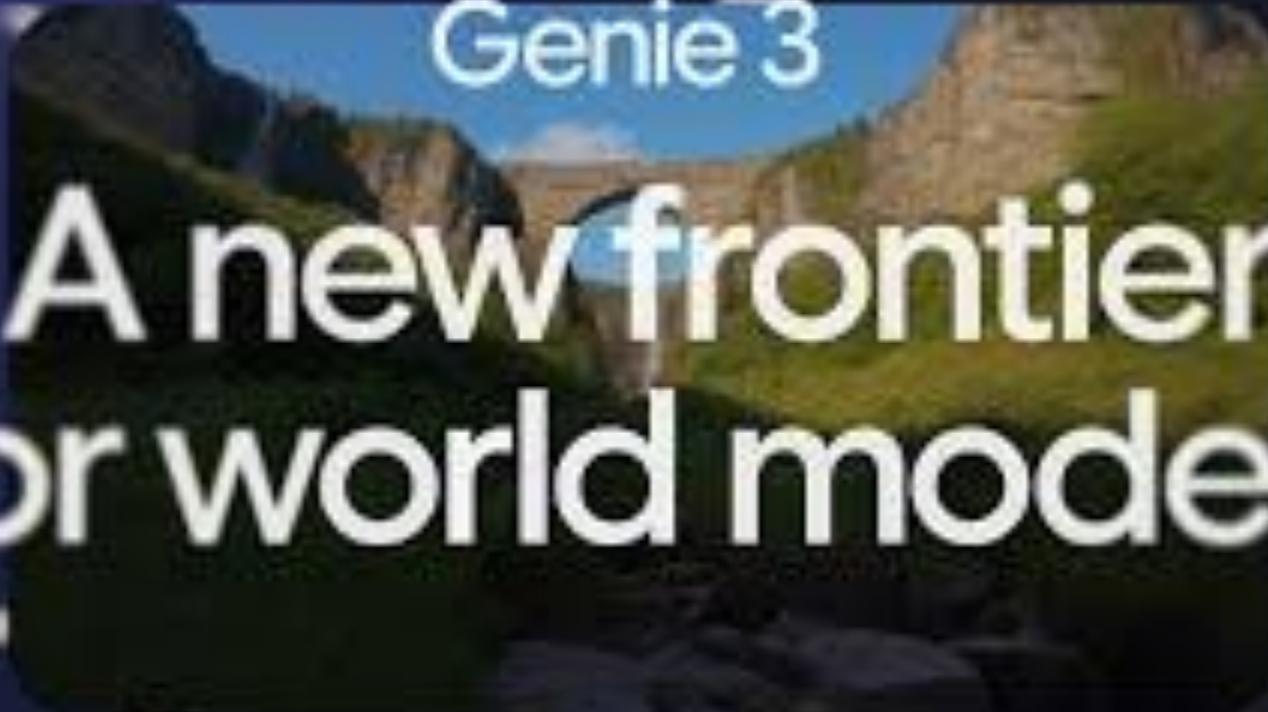
- **Nouvelle phase** : Optimisation logicielle > Puissance brute + Reasoning avancé + Approches agentiques
- **Convergence** : Les modèles fusionnent (fast + reasoning = unifié) et simplification de l'offre d'OpenAI
- **Open Source Revival** : OpenAI retour au source + la Chine mise sur la stratégie open source
- **La Chine avance, l'Occident l'ignore** : DeepSeek, Qwen, ByteDance performant sur les benchmarks dans le silence
- **Integration > Performance** : Gemini dans Chrome touche 3B d'utilisateurs

Meet Gemini in Chrome



Création sans limite

L'IA ne génère plus des vidéos, elle génère des mondes



Genie 3

A new frontier
for world models



RAY3

Aperçu des sorties estivales

VIDEO & MONDES

- **Veo 3 (Google)** - Vidéo + audio natif
- **Veo 3 Fast** - Génération mobile temps réel
- **Genie 3 (Google)** - Mondes interactifs explorables

IMAGE

- **Midjourney v7**
- **Google Nano Banana** - Domine les benchmarks sur l'édition d'image

AUDIO/MUSIQUE

- **Suno v4.5+** - Add Vocals, VST plugins
- **ElevenLabs Music** - Multilingue

Création - En bref

- **La barrière créateur disparaît** - Idée → Création sans compétences techniques
- **Multimodal par défaut** - Plus de silos : vidéo = image + son + mouvement
- **Les modèles-monde arrivent** - Genie 3, NVIDIA Isaac, World Labs → L'IA comprend la physique
- **Le pont vers les robots** - Vidéo IA = données d'entraînement pour robots physiques

L'argent ne doute pas



Des chiffres qui donnent le tourni

LEVÉES RECORDS

- Mistral : 1.7B€ (ASML lead) - val: 11.7B€ 🇫🇷
- Thinking Machines : \$2B seed (record mondial)
- Open AI : \$40B - val: \$300B

GUERRE DES TALENTS

- Google recrute Character.AI et Windsurf pour plusieurs milliards
- Meta acquière 49% de Scale pour \$14.3B
- "Acqui-hire" : jusqu'à 500M\$ par personne clé



ChatGPT

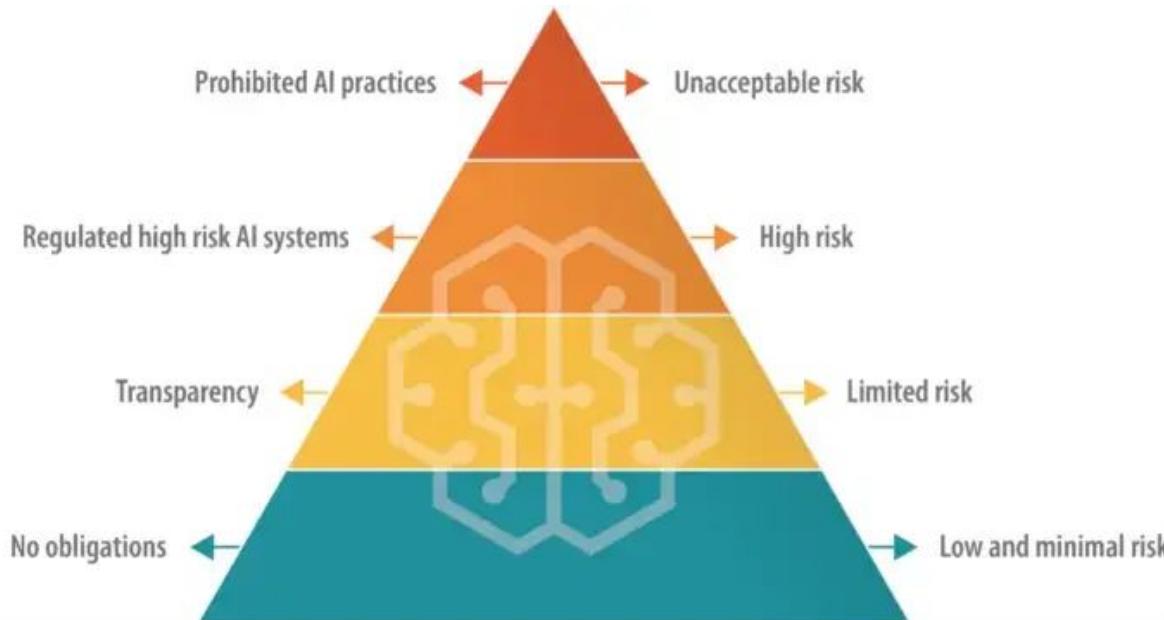
Le
Chat

Des chiffres qui donnent le tourni

- **Infrastructure & Innovation** - Financement massive pour les infrastructures et les talents
- **Consolidation accélère** - Meta-Scale = nouveaux géants verticaux (\$14.3B pour 49%)
- **Guerre des talents = guerre des millions** Meta paie 100-500M\$ par chercheur star
- **Paradoxe du ROI** - MIT : 95% sans retour MAIS investissements x10
- **Géopolitique de l'IA : USA vs Chine** L'Europe spectatrice malgré 40B\$ UK et Mistral

La société s'adapte

2 AOÛT - EU AI Act effectif L'Europe régule enfin l'IA



La société s'adapte

7 SEPTEMBRE - Anthropic paie 1.5B\$ Plus gros settlement copyright de l'histoire

PIXELS • INTELLIGENCES ARTIFICIELLES GÉNÉRATIVES

IA générative : Anthropic paiera 1,5 milliard de dollars pour éviter un procès sur le téléchargement illégal de livres

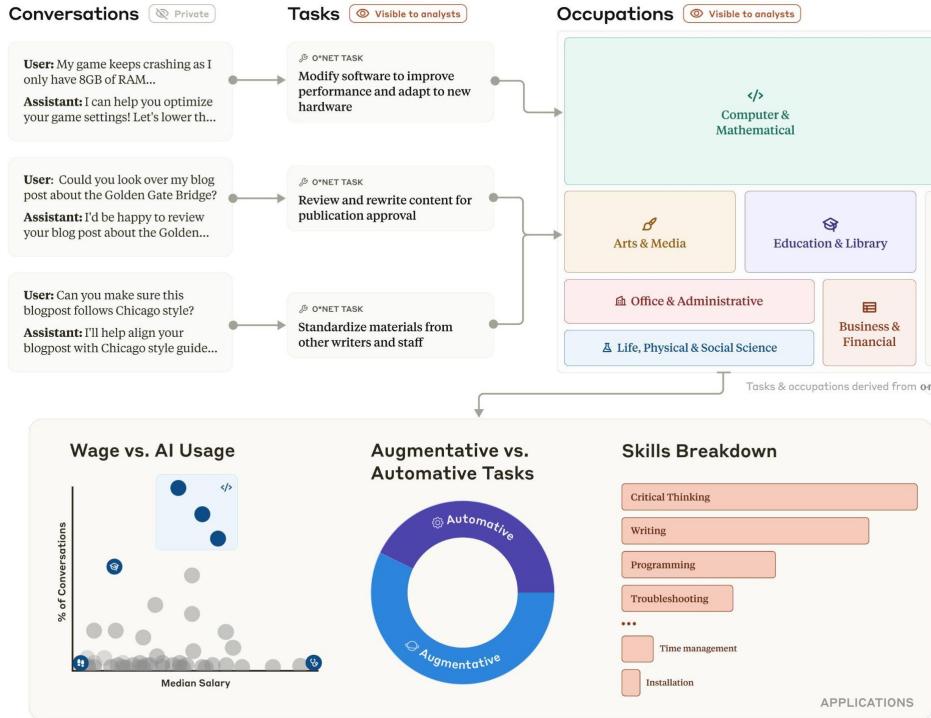
La start-up américaine d'intelligence artificielle a choisi la voie de l'accord financier avec les auteurs, leurs ayants droit et les éditeurs. Elle est poursuivie pour avoir téléchargé des œuvres à partir de sites pirates, afin d'alimenter ses modèles d'IA.

Le Monde avec AP et AFP

Publié le 06 septembre 2025 à 00h15, modifié le 06 septembre 2025 à 11h23 · ⏱ Lecture 2 min.

La société s'adapte

15 SEPTEMBRE - 40% des US workers De 3.7% à 40% en 24 mois (plus rapide qu'Internet)



News tech

Focus sur l'Agentic Coding

MCP-UI

Interactive UI Components for MCP

Build rich, dynamic user interfaces for your MCP applications with SDKs that bring UI to AI interactions.

[Get Started](#)[GitHub](#)[Live Demo](#)[About](#)

Client SDK

Provides a React component and Web Component for easy frontend integration. Render interactive UI resources and handle UI actions effortlessly.

Server SDKs

Powerful utilities to construct interactive UI for MCP servers. Create HTML, React, Web Components, and external app UI with ergonomic APIs for Typescript and Ruby.

Secure

All remote code executes in sandboxed iframes, ensuring host and user security while maintaining rich interactivity.

Flexible

Support for multiple content types, including iframes and Remote DOM components that match your host's look-and-feel.

MCP-UI Playground





```
{  
  "id": "bob",  
  "name": "Bob",  
  "avatarUrl": "/avatar2.png"  
}
```

→ Result

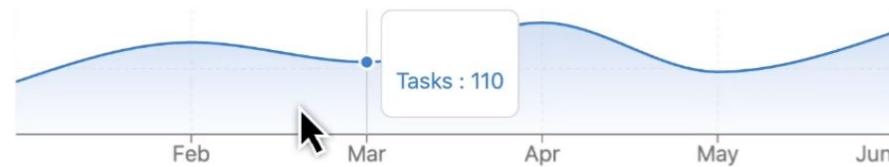


Bob

Performance Overview

Average 125

Reach 91%



1243

Completed

289

Pending

67

Blocked

Nudge

Send a message...



AI Agent in the terminal

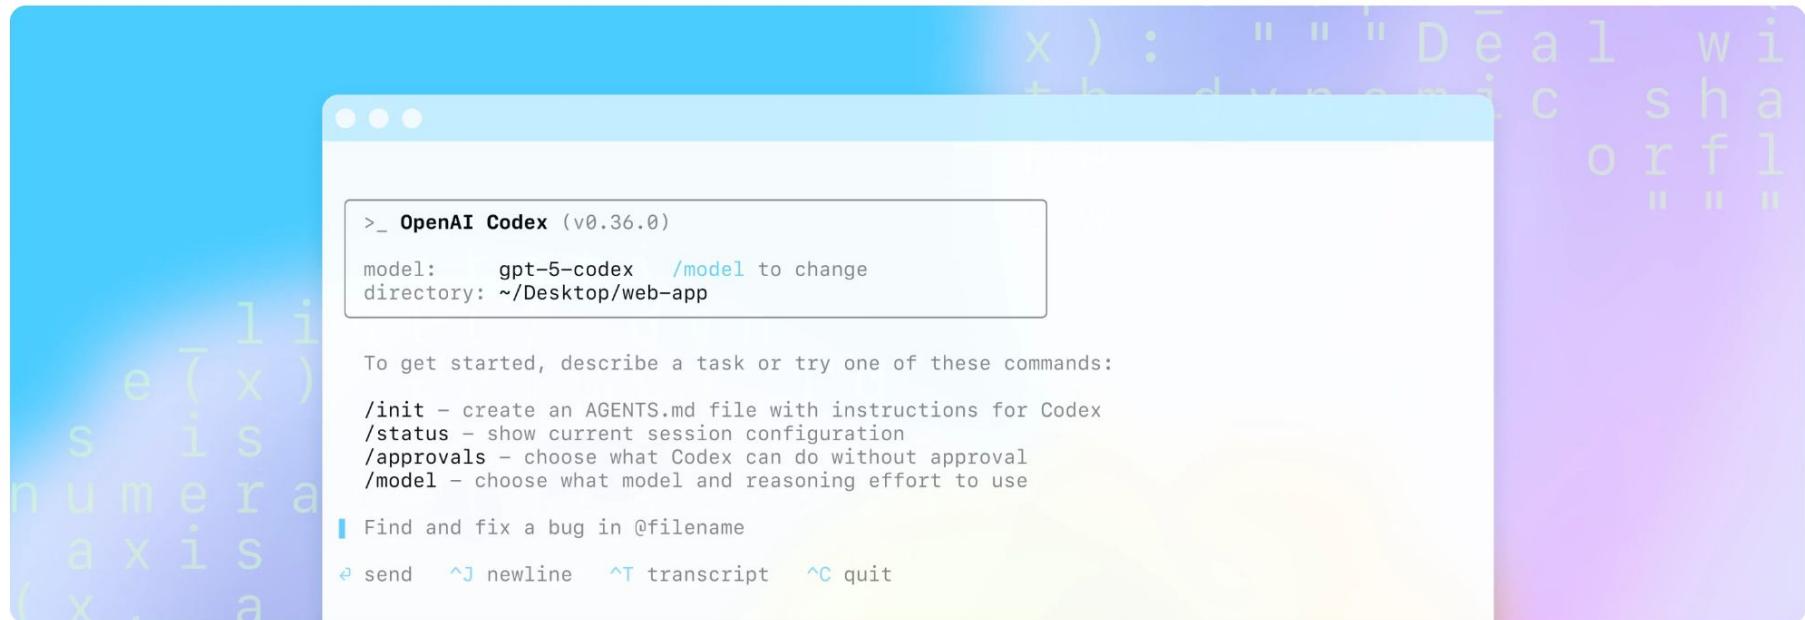
* Welcome to **Claude Code**

CLAUDE
CODE

Press **Enter** to continue

Codex CLI

Pair with Codex in your terminal





Anthropic ✨

@AnthropicAI

...

We're rolling out new weekly rate limits for Claude Pro and Max in late August. We estimate they'll apply to less than 5% of subscribers based on current usage.



Anthropic ✨ @AnthropicAI · Jul 28

...

Some of the biggest Claude Code fans are running it continuously in the background, 24/7.

These uses are remarkable and we want to enable them. But a few outlying cases are very costly to support. For example, one user consumed tens of thousands in model usage on a \$200 plan.

141

412

2.6K

2.3M



CLI vs IDE ?

Claude code ==



Pro

Claude Code is included in your Pro plan. Perfect for short coding sprints in smaller codebases with Sonnet 4.

\$17

Per month with annual subscription discount (\$200 billed up front). \$20 if billed monthly.

[Try Claude](#)



Max 5x

Claude Code is included in your Max plan. Great value for everyday use in larger codebases with access to both Sonnet 4 and Opus 4.1.

\$100

Per person billed monthly

[Try Claude](#)



Max 20x

Even more Claude Code included in your Max plan. Great value for power users with the most access to Opus 4.1.

\$200

Per person billed monthly

[Try Claude](#)

Additional usage limits apply. Prices shown don't include applicable tax.

```
~ » bunx ccusage@latest -s 20250610  
WARN Fetching latest model pricing from LiteLLM... ccusage 12:08:53 AM  
i Loaded pricing for 1111 models ccusage 12:08:53 AM  
  
Claude Code Token Usage Report - Daily
```

Date	Models	Input	Output	Cache Create	Cache Read	Total Tokens	Cost (USD)
2025 06-10	opus-4, sonnet-4	1,257	47,773	705,365	15,800,...	16,555,...	\$14.17
2025 06-11	opus-4, sonnet-4	9,358	26,086	947,715	16,997,...	17,980,...	\$11.48
2025 06-12	opus-4, sonnet-4	4,585	17,794	3,016,0...	12,514,...	15,552,...	\$24.40
2025 06-13	opus-4, sonnet-4	18,134	69,184	2,475,5...	61,119,...	63,682,...	\$36.72
2025 06-14	opus-4, sonnet-4	2,500	25,604	1,178,2...	33,471,...	34,678,...	\$22.03
2025 06-16	opus-4, sonnet-4	3,366	40,206	1,534,7...	51,222,...	52,800,...	\$40.14
2025 06-17	opus-4, sonnet-4	37,168	78,068	3,297,8...	78,011,...	81,424,...	\$63.89
Total		76,368	304,715	13,155,...	269,138...	282,675...	\$212.82

~ 1.1s » █

ccusage

GLM Coding Plan

The smartest investment for your coding productivity. [Learn More](#)

Monthly **-\$30**

Quarterly **-\$90**

Yearly **-\$360**

GLM Coding Lite

For lightweight workloads

☆ 50% off 1st Month

\$3 /month \$6/month

\$6 / month from 2nd month

Subscribe

✓ Powered by GLM-4.5

✓ Compatible with over 10 coding tools,
including Claude Code, Roo Code, Cline,
Kilo Code, OpenCode, Crush, and Goose,
with more being continuously added

GLM Coding Pro

For professional workloads

☆ 50% off 1st Month

\$15 /month \$30/month

\$30 / month from 2nd month

Subscribe

✓ All Lite plan benefits

✓ 5x Lite plan usage
✓ 40%-60% faster compared to Lite
✓ Access image & video understanding and
web search MCP

GLM Coding Max

For high-volume workloads

☆ 50% off 1st Month

\$30 /month \$60/month

\$60 / month from 2nd month

Subscribe

✓ All Pro plan benefits

✓ 4x Pro plan usage
✓ Guaranteed peak-hour performance
✓ Early access to new features

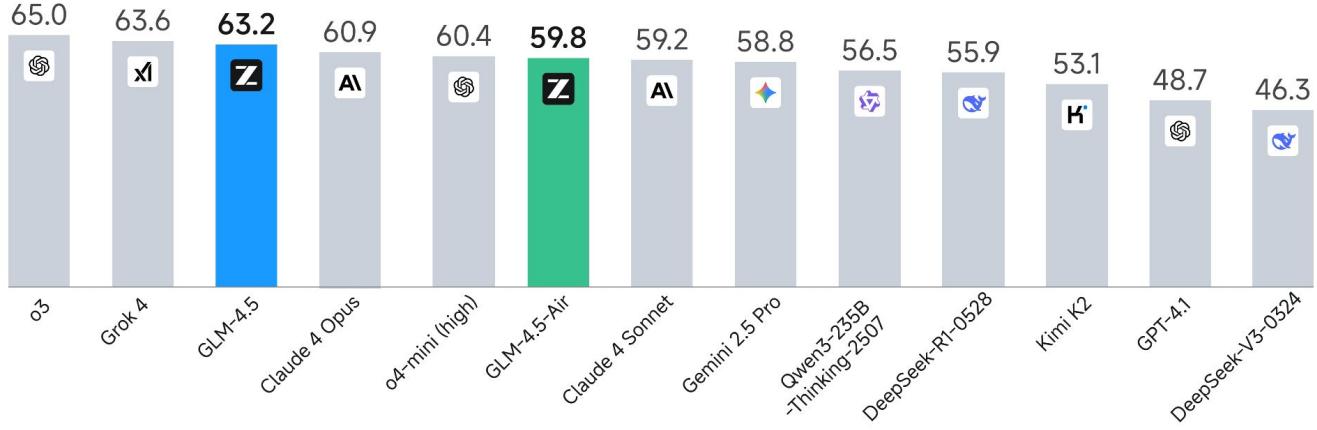
```
export ANTHROPIC_BASE_URL=https://api.z.ai/api/anthropic
```

```
export ANTHROPIC_AUTH_TOKEN=YOUR_API_KEY
```

LLM Performance Evaluation: Agentic, Reasoning, And Coding Benchmarks

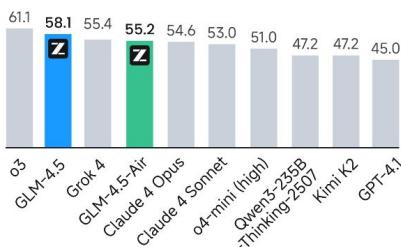


12 benchmarks: MMLU-Pro, AIME 24, MATH-500, SciCode, GPQA, HLE, LCB (2407-2501), SWE-Bench Verified, Terminal-Bench, TAU-Bench, BFCL V3, BrowseComp



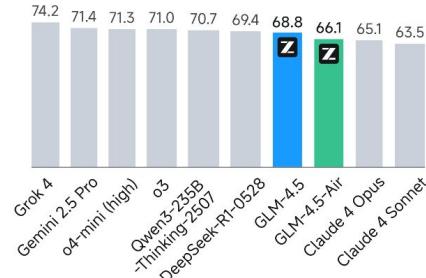
Agentic

Agentic Benchmarks: TAU-Bench, BFCL V3 (Full),
BrowseComp



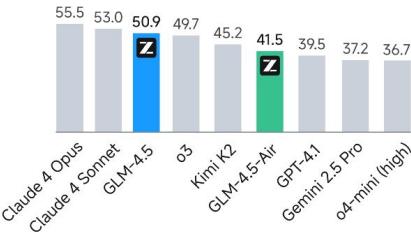
Reasoning

Reasoning Benchmarks: MMLU-Pro, AIME 24, MATH 500,
SciCode, GPQA, HLE, LCB (2407-2501)



Coding

Coding Benchmarks: SWE-Bench Verified, Terminal-Bench





Track Every Request. Go Low-Level. Never Hit Rate Limits Again.

The ultimate Claude API proxy with intelligent load balancing across multiple accounts. Full visibility into every request, response, and rate limit.

video_compressed.mp4 ▾



github.com/snipeship/ccflare

Claude code Toolchain

YOLO mode with safety

RchGrav / **claudebox**

Code Issues 21 Pull requests 9 Actions Projects ...



The Ultimate Claude Code Docker Development Environment - Run Claude AI's coding assistant in a fully containerized, reproducible environment with pre-configured development profiles.

MIT license

532 stars 62 forks 8 watching 6 Branches 7 Tags Activity

Public repository

CLAUDE CODE TEMPLATES

Ready-to-use configurations for your Claude Code projects

Claude Code

Blog

Github



Add New Agent

Create a new AI specialist agent

Installation Command

```
npx claude-code-  
templates@latest --  
agent=development-  
team/frontend-developer --yes
```

[View Details](#)

[Add to Stack](#)



1.6K

Backend Architect

Backend system architecture and API design specialist. Use PROACTIVELY for RESTful APIs, microservice boundaries, dat...



1.5K

Code Reviewer

Expert code review specialist for quality, security, and maintainability. Use PROACTIVELY after writing or modifying ...

DEVELOPMENT TEAM



1.4K

Ui Ux Designer

UI/UX design specialist for user-centered design and interface systems. Use PROACTIVELY for user research, wireframes...



1.3K

Prompt Engineer

Expert prompt optimization for LLMs and AI systems. Use PROACTIVELY when building AI features, improving agent perfor...



936

Fullstack Developer

Full-stack development specialist covering frontend, backend, and database technologies. Use PROACTIVELY for end-to-e...



859

Debugger

Debugging specialist for errors, test failures, and unexpected behavior. Use PROACTIVELY when encountering issues, an...

PROGRAMMING LANGUAGES



794

Python Pro

Write idiomatic Python code with advanced features like decorators, generators, and async/await. Optimizes performanc...



713

Task Decomposition Expert

Complex goal breakdown specialist. Use PROACTIVELY for multi-step projects requiring different capabilities. Masters ...



682

Context Manager

Context management specialist for multi-agent workflows and long-running tasks. Use PROACTIVELY for complex projects,...



635

Test Engineer

Test automation and quality assurance specialist. Use PROACTIVELY for test strategy, test automation, coverage analys...

EXPERT ADVISORS



610

DATABASE



582

PROGRAMMING LANGUAGES



531

PROGRAMMING LANGUAGES



527

CLAUDE.md
GEMINI.md
CONTEXT.md
DEVIN.md
ASSISTANT.md

...

<https://agents.md/>

AGENTS.md

A simple, open format for guiding coding agents,
used by over [20k open-source projects](#).

Think of AGENTS.md as a **README for agents**: a dedicated,
predictable place to provide the context and instructions to help
AI coding agents work on your project.

[Explore Examples](#)[View on GitHub](#)

AGENTS.md

Setup commands

- Install deps: `pnpm install`
- Start dev server: `pnpm dev`
- Run tests: `pnpm test`

Code style

- TypeScript strict mode
- Single quotes, no semicolons
- Use functional patterns where possible

One AGENTS.md works across many agents

Your agent definitions are compatible with a growing ecosystem of AI coding agents and tools:



Codex
from OpenAI



Amp



Jules
from Google



Cursor



Factory



RooCode



Aider



Gemini CLI
from Google



Kilo Code



opencode



Phoenix



Zed



Semgrep



Warp



Coding agent
from GitHub Copilot



VS Code



Ona



Devin
from Cognition

...except Claude Code #\$_!@

One AGENTS.md works across many agents

Your agent definitions are compatible with a growing ecosystem of AI coding agents and tools:



Codex
from OpenAI



Amp



Jules
from Google



Cursor



Factory



RooCode



Aider



Gemini CLI
from Google



Kilo Code



opencode



Phoenix



Zed



Semgrep



Warp



Coding agent
from GitHub Copilot



VS Code



Ona



Devin
from Cognition

In -s CLAUDE.md AGENTS.md

Spec-driven development

Spec-kit (by Github)

The screenshot shows a dark-themed code editor interface. On the left is the Explorer sidebar with the following file structure:

- HOMEBRIDGE_MANAGER...
 - .claude
 - memory
 - constitution_update...
 - constitution.md
 - scripts
 - specs/001-ok...
 - contracts
 - data-model.md
 - plan.md
 - quickstart.md
 - research.md
 - spec.md
 - tasks.md
 - src
 - templates
 - tests
 - CLAUDE.md
 - pyproject.toml
 - requirements.txt

The main editor area displays the contents of tasks.md, which includes:

- Validation, error handling, authentication
- Phase 3.4 - Integration (5 tasks):**
 - MCP server integration
 - Configuration and logging
 - Production readiness
- Phase 3.5 - Polish (9 tasks):**
 - Unit and performance tests [P]
 - Documentation updates
 - Quickstart validation
 - Code review
- Key Features**
 - **TDD Compliance:** All tests (T004-T014) before implementation
 - **Parallel Execution:** 23 tasks marked [P] for concurrent execution
 - **Constitutional Adherence:** Library-first architecture, structured logging
 - **Production Ready:** Performance targets, error handling, monitoring
 - **Immediate Executability:** Each task includes exact file paths and specifications
- Ready for implementation execution following Spec-Driven Development methodology.

Below the main text, there is a "Context Usage" section showing token distribution:

Category	Tokens	Percentage
claude-sonnet-4-20250514	88k	44%
System prompt	3.1k	1.5%
System tools	11.4k	5.7%
MCP tools	927	0.5%
Messages	72.9k	36.5%
Free space	111.6k	55.8%

At the bottom, there is a "CP tools" section with token counts:

- /mcp: 428 tokens
- mcp_ide_getDiagnostics (ide): 499 tokens
- mcp_ide_executeCode (ide): 499 tokens

github.com/github/spec-kit/

TaskMaster

Step 1/7: Initialize a new Taskmaster project

```
$ task-master init
```

- ✓ Project initialized: Taskmaster
- ✓ Created tasks directory
- ✓ Created scripts directory
- ✓ Added Cursor rules
- ✓ Ready to parse your PRD!

Step 2/7: Parse the Product Requirements Document to generate tasks

```
$ task-master parse-prd scripts/prd.txt
```

Analyzing PRD...

Identifying core features...

Breaking down implementation steps...

Establishing dependencies...

Generated 10 tasks with logical dependencies.

Step 3/7: List all generated tasks with their dependencies

```
$ task-master list
```

ID	Status	Priority	Title	Dependencies
1	done	high	Initialize Next.js Project	None
2	pending	high	Implement User Authentication	1
3	pending	high	Create Database Schema	1
4	pending	medium	Implement API Routes	1, 2, 3
5	pending	high	Build UI Components	1

Step 4/7: Analyze the complexity of each task

```
$ task-master analyze-complexity
```

Analyzing task complexity...

Using Claude to evaluate implementation difficulty...

Checking for external dependencies...

Identifying potential roadblocks...

Complexity analysis complete!

Step 5/7: View the complexity report to identify tasks needing breakdown

```
$ task-master complexity-report
```

ID	Task	Complexity	Recommendation
2	Implement User Authentication	8/10	Expand into subtasks
3	Create Database Schema	6/10	Consider expansion
4	Implement API Routes	7/10	Expand into subtasks
5	Build UI Components	5/10	Manageable as-is
6	Implement State Management	4/10	Manageable as-is

Step 6/7: Expand the complex authentication task into subtasks

```
$ task-master expand --id=2
```

Expanding task 2: Implement User Authentication

Breaking down into subtasks...

- ✓ Created subtask 2.1: Set up authentication provider
- ✓ Created subtask 2.2: Implement sign-up functionality

Spec on Github instead of .md

The screenshot shows a GitHub issue page for a task titled "Create Secret Usage Tracker". The URL is <https://github.com/automazeio/ccpm/issues/19>. The page has a sidebar on the left with an "Epic: Secrets Example" card and a main content area for the task.

Epic: Secrets Example (Card):

- Open
- Feature
- 0 / 10

Create Secret Usage Tracker #19 (Issue Card):

- Open
- Parent: Epic: Secrets Example File Generation

Description: Implement `SecretsExampleGenerator` to create the `.env-format secrets.example` file with usage comments. This component transforms the tracked usage data into a developer-friendly example file.

Acceptance Criteria:

- Track secret key → set of usage paths mapping
- Support adding new usage: `track_usage(secret_key, context_path)`
- Get sorted usage map: `get_usage_map()` returns `Dict[str, List[str]]`
- Clear tracking with timestamp: `clear()` method
- Thread-safe operations for concurrent access
- Unit tests with 100% coverage

Technical Details:

- Use `collections.defaultdict` for efficient mapping
- Implement thread safety with `threading.RLock()`
- Store scan timestamp for tracking when data was collected
- Convert to sorted lists when returning via `get_usage_map()`

Dependencies: None - this is a foundational component

Assignees: No one - Assign yourself

Labels: task

Type: No type

Projects: CCPM - Engineering, Status: Backlog

Milestone: No milestone

Relationships: Parent issue: Epic: Secrets Example File Generation

Development: Code with agent mode

[Create a branch](#) for this issue or link a pull request.

[Notifications](#) [Customize](#)

<https://github.com/automazeio/ccpm>

Background agents



Lovable

Build apps and websites by chatting with AI

Build something lovable |



Agents Dashboard

Ask Cursor to build, fix bugs, explore



Claude 4.1 Opus **MAX** ▾



samber/headercheck ▾ ⚙ main ▾

Try these examples to get started



Write documentation



Optimize performance



Find and fix 3 bugs



What are we coding next?

We're working on moving off of the /api/wham/tasks endpoint and onto the new faster condensed payload /api/wham/tasks/list api. Can you update the tests in test_tasks_api.py to make sure they all still pass on the new api?

monorepo ▾

main ▾

Ask

Code

Tasks Archive

Diagnose issues with status updates

+76 -47

Just now · monorepo

Convert non-critical components to React.lazy with Suspense fallbacks

+142 -12

2h ago · monorepo

Merged

Fix uncontrolled to controlled warning in ToggleGroup

+158 -25

3h ago · monorepo

Merged



Spark



Dream it. See it. Ship it.

Transform ideas into full-stack intelligent apps in a snap.

Publish with a click.

Create web apps with React and TypeScript to prototype ideas, build tools, and more



Attach image



To-do list app

Recipe collection app

Time zone converter

Spark uses AI. Check for mistakes. [Terms of service](#).

Looking for speaker 🤘



AI-powered engineering

 multigres / **multigres**         

<> Code Issues 6 Pull requests 2 Discussions Actions Projects Security ...

Add Postgres Parser #109

 sougou merged 130 commits into `multigres:main` from `GuptaManan100:postgres-parser-port` 3 days ago

Conversation 0 Commits 130 Checks 4 Files changed 300+ +286,786 -11

 GuptaManan100 commented 3 days ago · edited Collaborator ...

Description

This PR adds the Postgres parser for multigres to be able to parse and understand the Postgres queries. The parser is written completely in go, but is inspired by the original Postgres parser.

Reviewers
 sougou ✓

Assignees
 GuptaManan100

<https://github.com/multigres/multigres/pull/109>

Focus sur les infra d'inférence

Snowflake: Speculative decoding

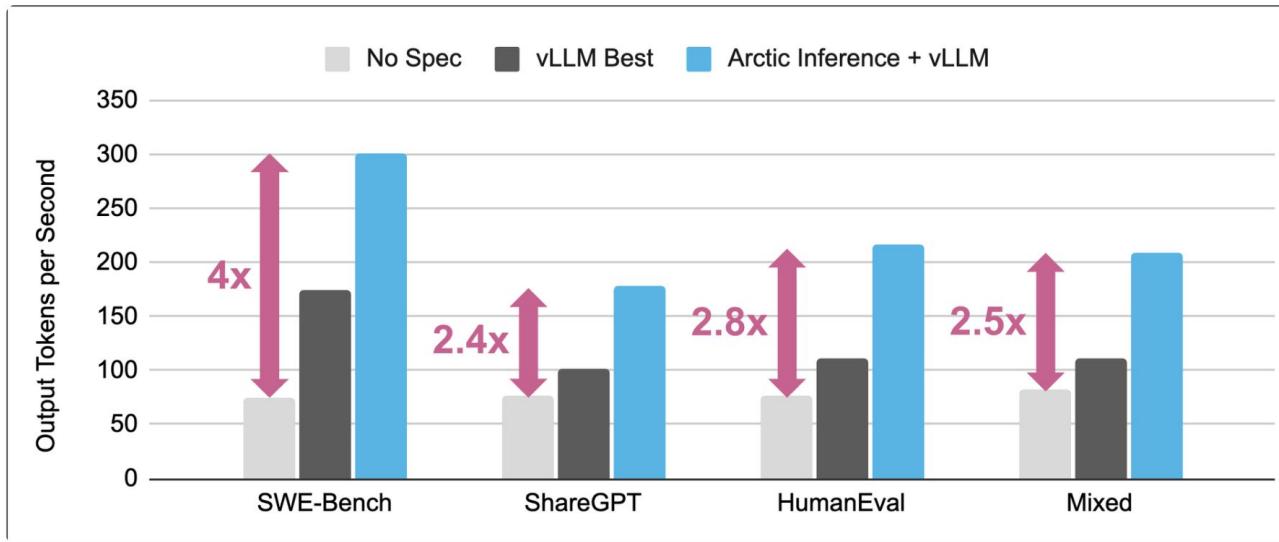


Figure 1: Generation speed for Llama-3.1-70B-Instruct, using speculative decoding in vLLM running on 8xH100 GPUs.

<https://www.snowflake.com/en/engineering-blog/fast-speculative-decoding-vllm-arctic/>

Snowflake: Speculative decoding

Serving

```
vllm serve Snowflake/Llama-3.1-SwiftKV-8B-Instruct \
--quantization "fp8" \
--tensor-parallel-size 1 \
--ulysses-sequence-parallel-size 2 \
--enable-shift-parallel \
--speculative-config '{
    "method": "arctic",
    "model": "Snowflake/Arctic-LSTM-Speculator-Llama-3.1-8B-Instruct",
    "num_speculative_tokens": 3,
    "enable_suffix_decoding": true,
    "disable_by_batch_size": 64
}'
```



<https://github.com/snowflakedb/ArcticInference>

Tokenomics

Mixture of Expert (MoE):

<https://www.tensoreconomics.com/p/moe-inference-economics-from-first>

Hardware + inference engines:

<https://blog.ovhcloud.com/gpu-for-llm-inferencing-guide/>



“Biais politiques dans les LLMs: comprendre, analyser et atténuer”



Akram Elbouanani

PhD student @ CEA



23 Septembre 2025 à 19h15



4 rue Voltaire, 44100 Nantes

chez

icilundi



[sfɛir]

lonestone

Slides dispo sur:

<https://github.com/genai-nantes-meetup/meetups/>





Presented by
Cursor Community >

[Subscribe](#)

Cursor community hosted meetups, hackathons, workshops taking place around the world. Learn more here: cursor.com/community

X ⊕

Cursor Meetup Nantes

LIVE

SEPT

25

Thursday 25 September

19:00 - 20:30 CEST



Google Meet

Registration

Welcome! To join the event, please register below.

[Register](#)

[About Event](#)

Ready to take your Cursor game to the next level? 🚀