



“Prompt to Video: comment l'IA redéfinit la publicité”



Yannis Sulfourt
Designer et formateur

Maxime Bodereau
Art Director @ Seelab



19 Novembre 2025 à 19h00
54 Bd de Berlin, 44000 Nantes

chez



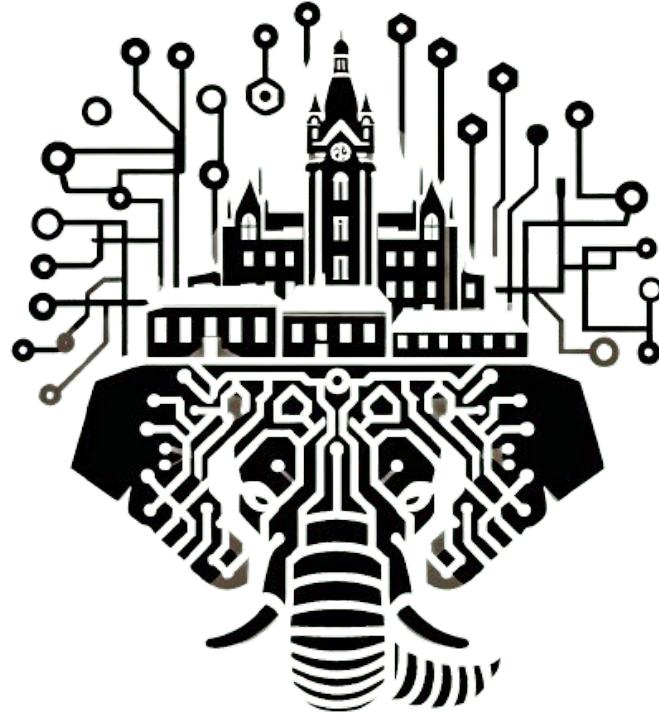
[sfɛir]

lonestone

Qui n'est jamais venu au meetup Gen AI Nantes ?

GenAI Nantes

- Microphone icon: 15 événements / an
- Pirate flag icon: 1 hackathon (Shift)
- Construction worker icon: 2 workshops
- Heart icon: 1 communauté de 900p*



* 9.000p selon le syndicat des llamas 

SHIFT

Le Hackathon Gen AI

27 au 29 Mars



techstars_ **Startup Weekend** **Nantes**

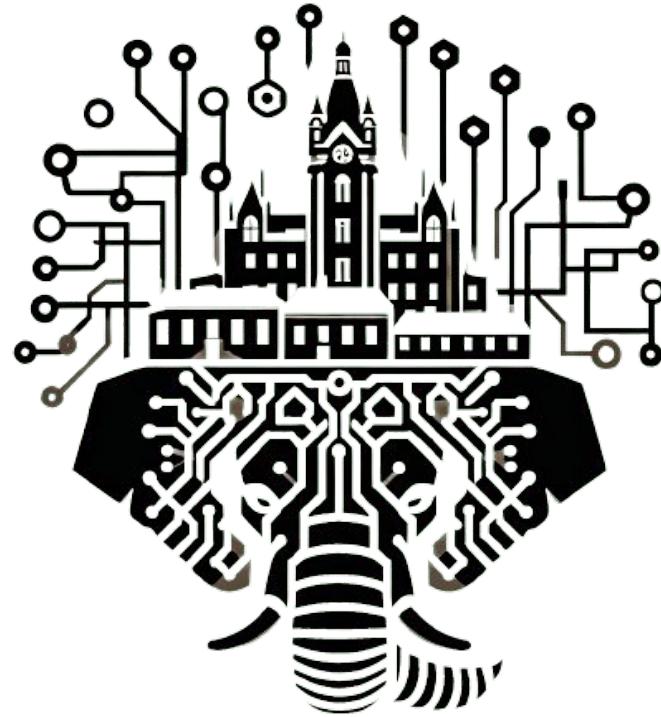
Du 21/11 au 23/11
Le Palace _jclundi



Réserve ta place !



Support us! ❤



Qui est designer ?



Qui est tech ?



Qui est [autre] ?



Qui utilise Midjourney ? 

Qui utilise Imagen/Veo (Google) ? 🍪

Qui utilise Seelab ? 

Qui recrute dans la GenAI ? 

Qui cherche un job dans la GenAI ?



Qui souhaite s'associer dans la GenAI ? 

Schedule

👀 1- News

🔍 2- Prompt to Video

🍻 3- Enjoy

News non-tech



Yann LeCun quitte Meta pour fonder sa startup

Contexte : Meta a réduit le rôle de son labo historique FAIR (Fundamental AI Research) après l'échec de son modèle Llama 4 à rivaliser avec les modèles rivaux les plus avancés pour fonder MSL (Meta Superintelligence Lab).





Cybersécurité : l'IA, nouvelle arme des hackers

Anthropic révèle l'usage des LLMs dans toutes les étapes des cyberattaques

De nouveaux kits de ransomware

« no-code » générés par IA en vente sur le dark web.

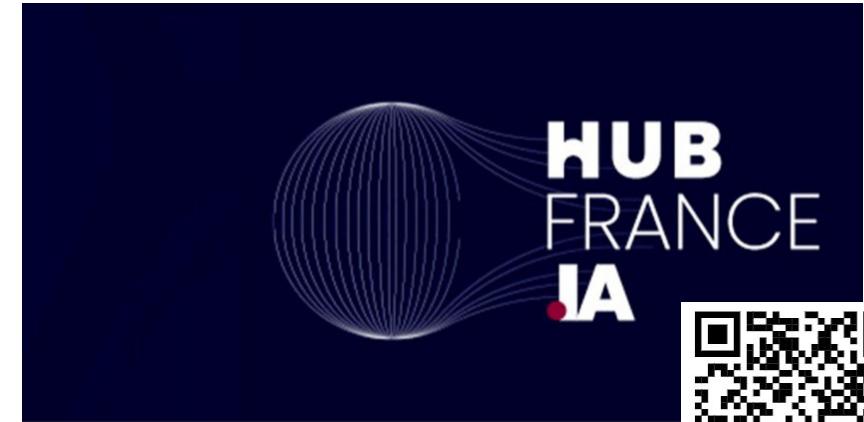
Les nords coréens se servent d'IA génératives pour produire des faux CVs, passer des entretiens techs et se faire embaucher à distance.





Digital Omnibus

- Fin novembre 2025
- Collecte des données (anonymisées) sans consentement explicite



France & AI Act

Partage de la surveillance de l'IA entre régulateurs existants, chacun compétent dans son domaine :

- CNIL pour les données personnelles
- ARCOM pour certains contenus
- ANSM pour les dispositifs médicaux, etc.)

Le tout sous la coordination de la DGCCRF (Concurrence & Consommation) et de la DGE (Entreprises).

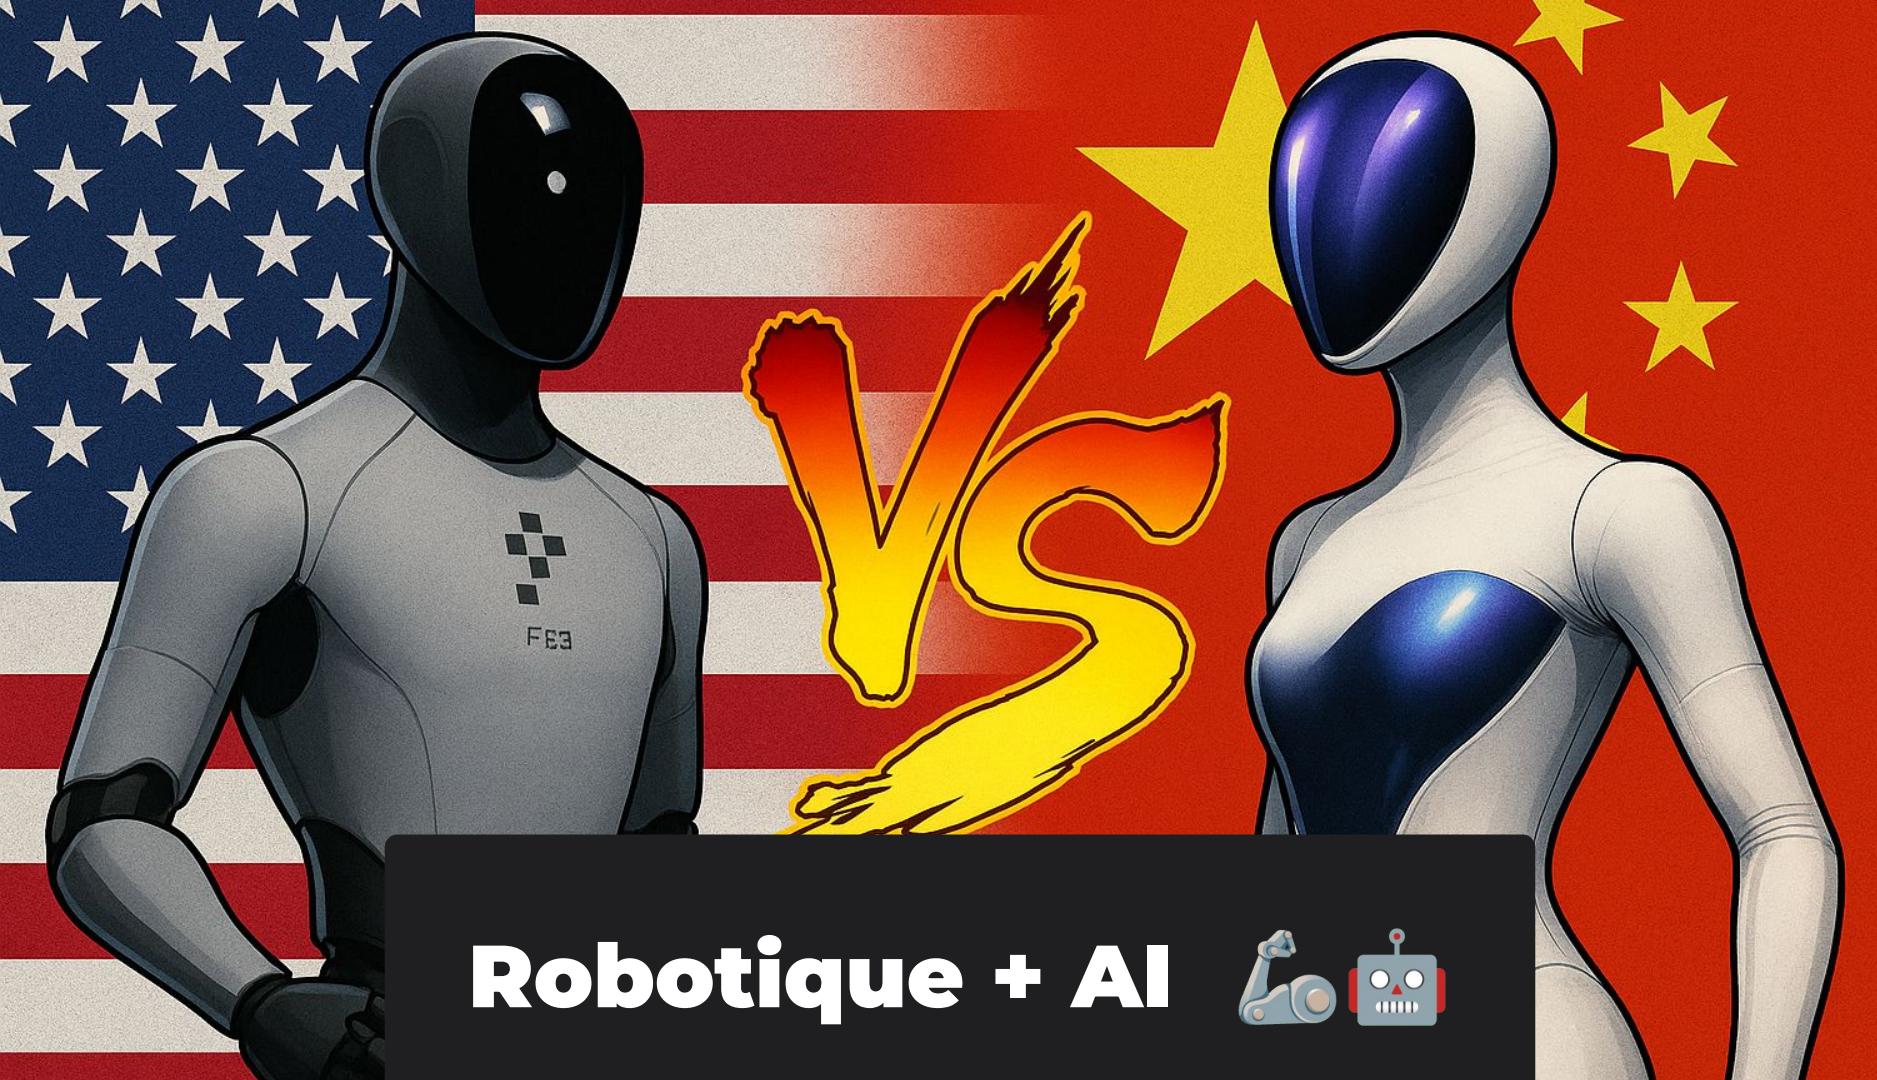


Hébergeurs de contenus & IA : Scrapping VS GEO

- robot.txt & CAPTCHA
- Labyrinthes
- Deals (Grands groupes / Petits)
- Droits de péage (Cloudflare Pay Per Crawl)
- License RSL (Really Simple Licensing)
- BNF : le contenu 
- GEO







Robotique + AI



Mais aussi ...

Grem, sage et au
grand cœur.



Grok, la fusée fougueuse.



Gabbo, le robot
sympathique et fidèle.



Doudou Chatbots

heycurio.com

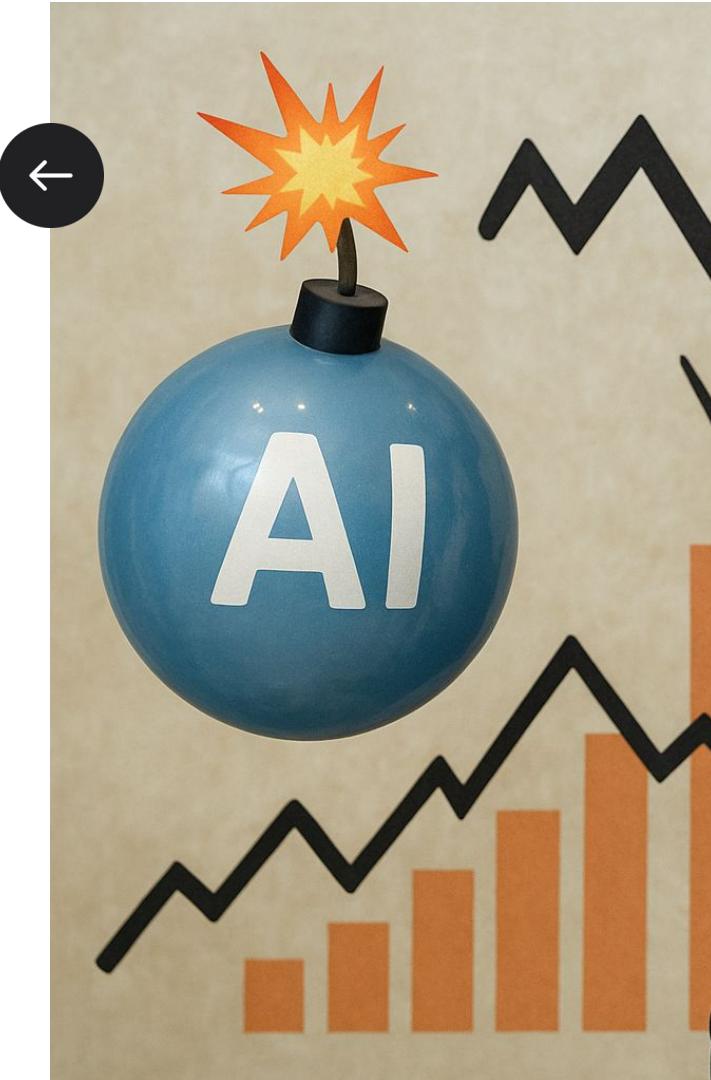


Craintes de bulle financière de l'IA

FMI + banque d'Angleterre

Le PDG de Goldman Sachs anticipe une correction boursière de 10 à 20 % dans les deux prochaines années, et le gouverneur de la BoE a souligné que si les gains de productivité liés à l'IA sont bien réels, la rentabilité effective du secteur reste très incertaine.

Michael Burry (le monsieur qui a anticipé la crise des subprimes) a misé 1Md\$ de dollar sur l'explosion de la bulle.





OpenAI freine les rumeurs d'IPO à 1000 Mds\$

Focus sur la croissance, la R&D et l'expansion de l'infrastructure avant d'envisager une rentrée en bourse.

... et se restructure en PBC

Par ailleurs Microsoft voit ses parts diluées de 32,5 à 27% du capital.



GPT-5



Gemini 3



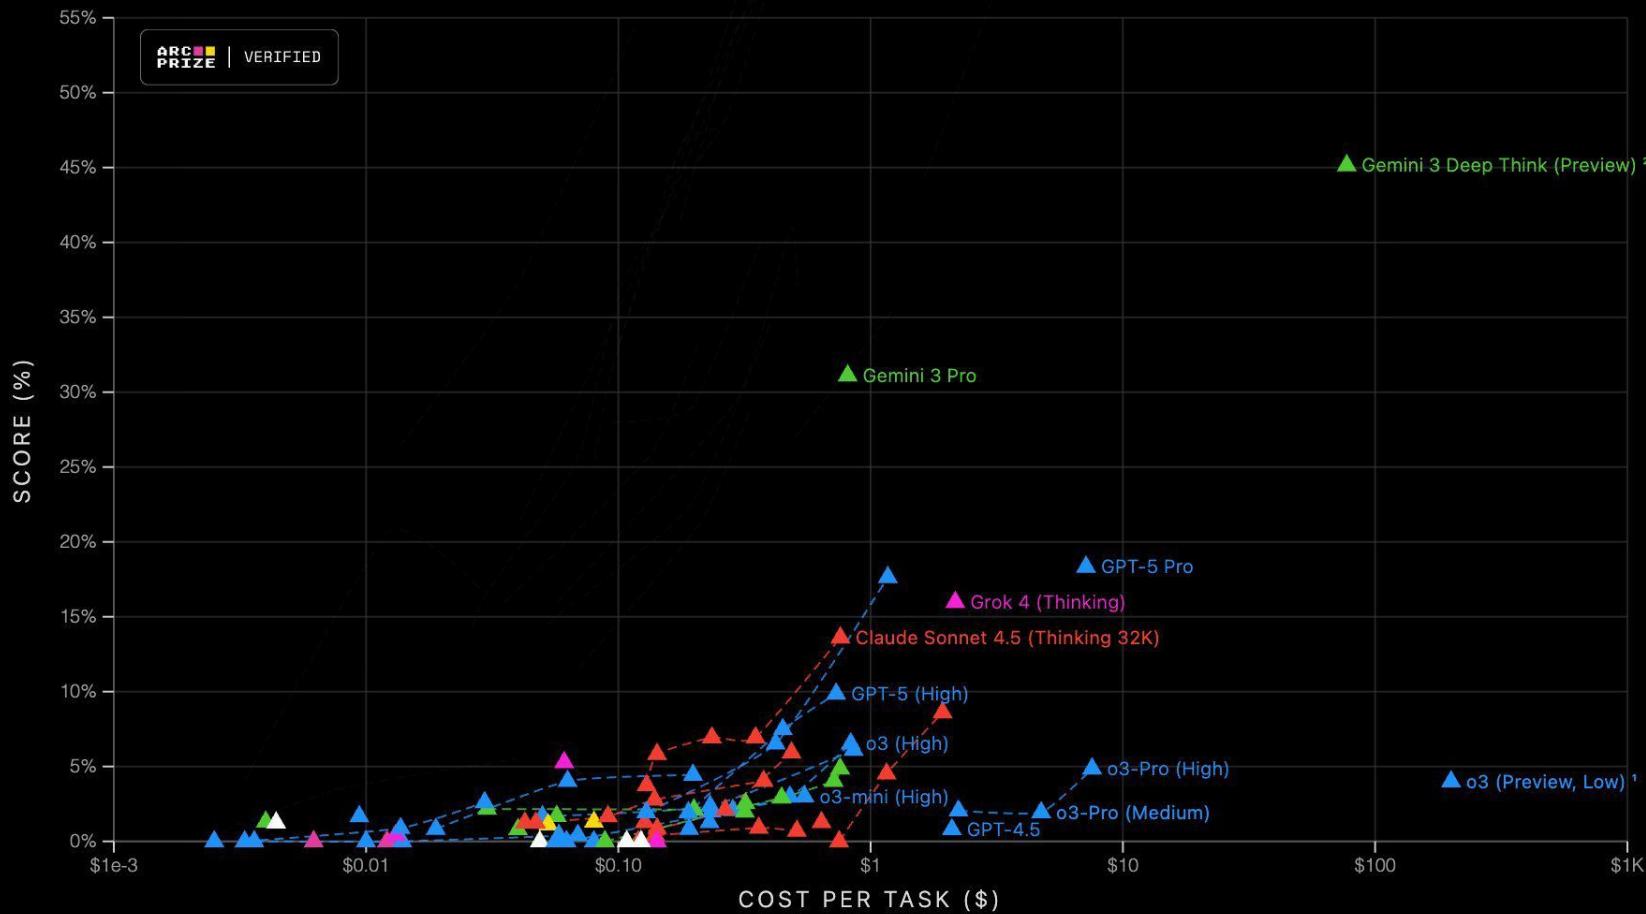
- Codage (Front-end et app rapide)
- Raisonnement plus long
- Conso électrique ?

- 1M de token en entrée
- Superforme sur des tests
- Antigravity ☺
- Intégré à Google Search (US)
- Simulation, visualisation, interactions

Benchmark	Description	Gemini 3 Pro	Gemini 2.5 Pro	Claude Sonnet 4.5	GPT-5.1
Humanity's Last Exam	Academic reasoning No tools With search and code execution	37.5% 45.8%	21.6% —	13.7% —	26.5% —
ARC-AGI-2	Visual reasoning puzzles ARC Prize Verified	31.1%	4.9%	13.6%	17.6%
GPQA Diamond	Scientific knowledge No tools	91.9%	86.4%	83.4%	88.1%
AIME 2025	Mathematics No tools With code execution	95.0% 100%	88.0% —	87.0% 100%	94.0% —
MathArena Apex	Challenging Math Contest problems	23.4%	0.5%	1.6%	1.0%
MMMU-Pro	Multimodal understanding and reasoning	81.0%	68.0%	68.0%	76.0%
ScreenSpot-Pro	Screen understanding	72.7%	11.4%	36.2%	3.5%
CharXiv Reasoning	Information synthesis from complex charts	81.4%	69.6%	68.5%	69.5%
OmniDocBench 1.5	OCR Overall Edit Distance, lower is better	0.115	0.145	0.145	0.147
Video-MMMU	Knowledge acquisition from videos	87.6%	83.6%	77.8%	80.4%
LiveCodeBench Pro	Competitive coding problems from Codeforces, ICPC, and IOI	2,439	1,775	1,418	2,243
Terminal-Bench 2.0	Agentic terminal coding Terminus-2 agent	54.2%	32.6%	42.8%	47.6%
SWE-Bench Verified	Agentic coding Single attempt	76.2%	59.6%	77.2%	76.3%
t2-bench	Agentic tool use	85.4%	54.9%	84.7%	80.2%
Vending-Bench 2	Long-horizon agentic tasks Net worth (mean), higher is better	\$5,478.16	\$573.64	\$3,838.74	\$1,473.43
FACTS Benchmark Suite	Held out internal grounding, parametric, MM, and search retrieval benchmarks	70.5%	63.4%	50.4%	50.8%
SimpleQA Verified	Parametric knowledge	72.1%	54.5%	29.3%	34.9%
MMMLU	Multilingual Q&A	91.8%	89.5%	89.1%	91.0%
Global PIQA	Commonsense reasoning across 100 Languages and Cultures	93.4%	91.5%	90.1%	90.9%
MRCR v2 (8-needle)	Long context performance 128k (average) 1M (pointwise)	77.0% 26.3%	58.0% 16.4%	47.1% not supported	61.6% not supported

For details on our evaluation methodology please see deepmind.google/models/evals-methodology/gemini-3-pro

ARC-AGI-2 LEADERBOARD



What do you want to learn today?



Apply your sk



Explore our Cloud learning catalog

Your home for building AI skills and more. Get hands on with Google Skills Labs, dive into in-depth courses and learn directly from the experts. When ready, prove your new skills with career-boosting credentials

Search catalog



Badge ▾

Format ▾

Level ▾

Language ▾

Duration ▾

Course Skill badge

Classify Images with TensorFlow on Google Cloud

Earn the intermediate Skill Badge by completing the Classify Images with

Course Skill badge

Derive Insights from BigQuery Data

Complete the introductory Derive Insights from BigQuery Data skill badge course to demonstrate skills in the following: Write...

Google lance Skills pour se former à l'IA

Google dévoile Skills, une plateforme gratuite, ouverte à toutes et tous, qui regroupe plus de 3 000 formations autour de l'IA et des technologies associées.





Premier médicament IA en phase de test

Le laboratoire Insilico Medicine a révélé que son composé Rentosertib – découvert grâce à une plateforme d'IA générative – a montré des résultats prometteurs en phase IIa contre la fibrose pulmonaire idiopathique.



News tech

Claude Code & Cursor v2



Claude Code & Cursor v2

- Sonnet & Haiku 4.5
- Subagents
- Skills
- Plugins
- Checkpoint/Rewind
- Extension VS Code
- Claude Agent SDK
- Composer 1
- Nouvelle UX orientée agent
- Manager d'agent
- Intégration navigateur
- Voice mode
- Amélioration des revues

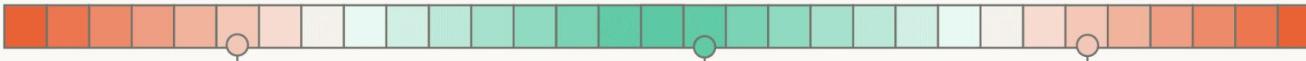
L'ingénierie du contexte

Calibrating the system prompt

Too specific

Just right

Too vague



You are a helpful assistant for Claude's Bakery.
You must respond to the name Claude.
For every user request you MUST FOLLOW THESE STEPS:

1. Identify the user intent as one of the following: ["incident_resolution", "general_inquiry", "order_resubmission", "account_maintenance", "requires_escalation"]
- 2.

- if user intent is "incident_resolution", ask 3 followup questions to gather information, then always call the resolve tool
- if user intent is "general_inquiry", do not ask followup questions and answer in one shot
- if user intent ...
- ...

3. Here is an exhaustive list of cases that should be tagged as "requires_escalation":

- If the intent is incident_resolution but the user is in a different country
- If the user left a physical belonging in the store
- ...

4. Once you've ruled out escalation scenarios you should consider all the tools at your disposal.

5. If the user_request contains an order_id you should tag the user intent as "order_resubmission", unless the user meets 5/7 of the following requirements:

- User is asking for time update
- User is asking for location update
- ...

6. If the user wants to request a new order, but they already have another order in flight, you should follow these 5 steps of the resolution procedure:

- (1) Call check_order tool to see where the current order is
- ...

You are a customer support agent for Claude's Bakery.
You specialize in assisting customers with their orders and basic questions about the bakery. Use the tools available to you to resolve the issue efficiently and professionally.

You have access to order management systems, product catalogs, and store policies. Your goal is to resolve issues quickly when possible. Start by understanding the complete situation before proposing solutions, ask follow-up questions if you do not understand.

Response Framework:

1. Identify the core issue - Look beyond surface complaints to understand what the customer actually needs
2. Gather necessary context - Use available tools to verify order details, check inventory, or review policies before responding
3. Provide clear resolution - Offer concrete next steps with realistic timelines
4. Confirm satisfaction - Ensure the customer understands the resolution and knows how to follow up if needed

Guidelines:

- When multiple solutions exist, choose the simplest one that fully addresses the issue
- If a user mentions an order, check its status before suggesting next steps
- When uncertain, call the human_assistance tool
- For legal issues, health/allergy emergencies, or situations requiring financial adjustments beyond standard policies, call the human_assistance tool
- Acknowledge frustration or urgency in the user's tone and respond with appropriate empathy

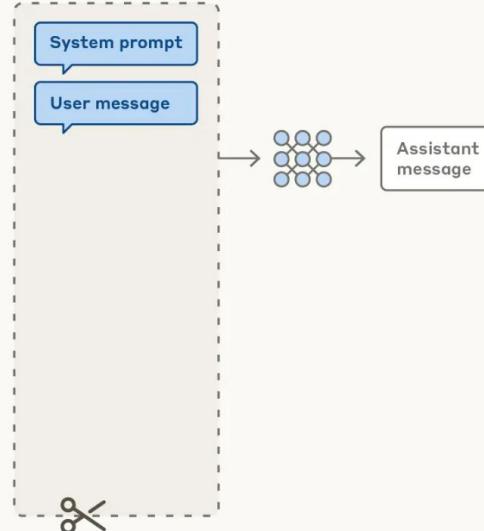
You are a bakery assistant, you should attempt to solve customers issues in a manner consistent with the principles and essence of the company brand. Escalate to a human if needed.

L'ingénierie du contexte

Prompt engineering vs. context engineering

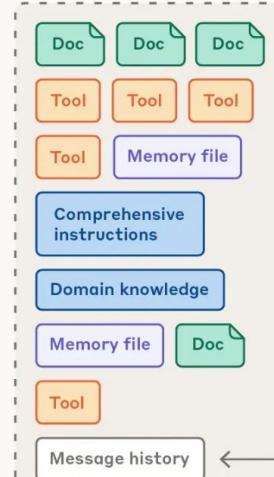
Prompt engineering for single turn queries

Context window

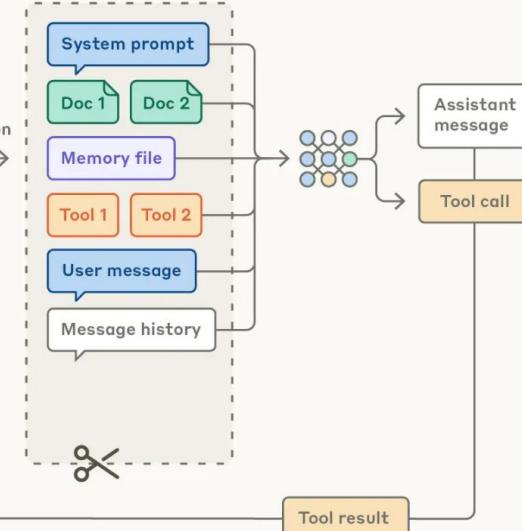


Context engineering for agents

Possible context to give model

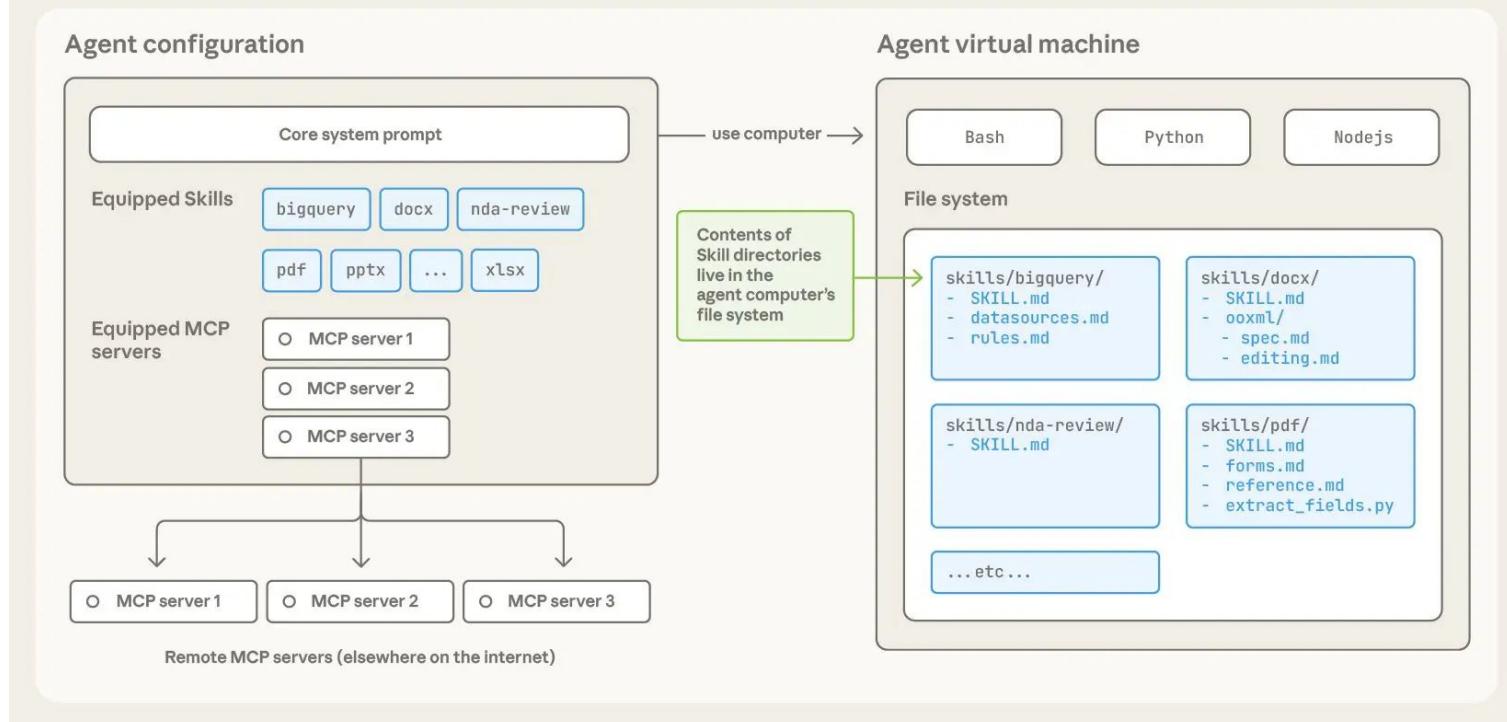


Context window



L'ingénierie du contexte

Agent + Skills + Virtual Machine



Plugins

CLAUDE CODE TEMPLATES

• Ready-to-use configurations for your Claude Code projects

[Blog](#) [Discord](#) [GitHub](#)

downloads 58k ⚡ Star 11k License MIT

• Search (components/settings/templates)

↳ Build your personalized development stack

> *Search components...*

[agents \(163\)](#) [commands \(210\)](#) [settings \(60\)](#) [hooks \(39\)](#) [mcp \(58\)](#) [plugins \(NEW\)](#) [skills \(19\)](#)

category:

All AI Specialists API GraphQL Blockchain Web3 Business Marketing Data AI Database Deep Research Team Development Team Development Tools
 Devops Infrastructure Documentation Expert Advisors FFmpeg Clip Team Game Development Git MCP Dev Team Modernization Obsidian Ops Team OCR Extraction Team
 Performance Testing Podcast Creator Team Programming Languages Realtime Security Web Tools

SORT BY: [Most Downloaded](#) ▾

Add New Agent
Create a new AI specialist agent

Frontend Developer

Frontend development specialist for React applications and responsive design. Use PROACTIVELY for UI components, stat...

Code Reviewer

Expert code review specialist for quality, security, and maintainability. Use PROACTIVELY after writing or modifying ...

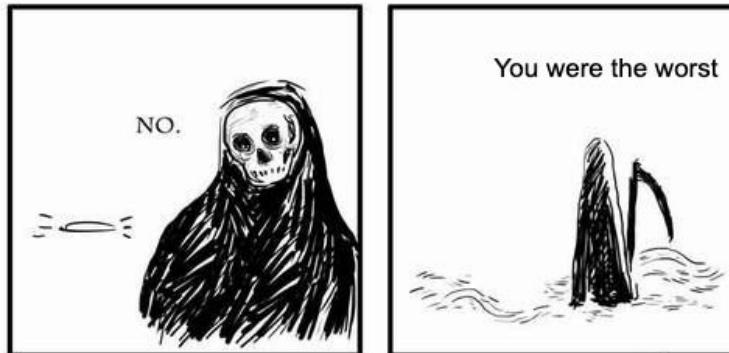
Backend Architect

Backend system architecture and API design specialist. Use PROACTIVELY for RESTful APIs, microservice boundaries, dat...

IDE augmenté



Le RAG est mort, vive le RAG



imgflip.com

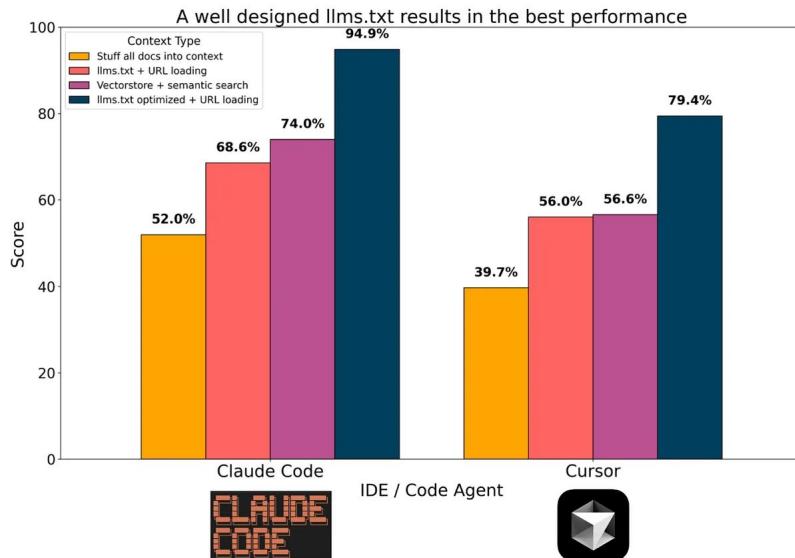
seebangnow

Le RAG est mort, vive le RAG

RAG->Agentic search:

- Hierarchical Structure Preservation
- Table Integrity
- Cross-Reference (eg: “see §3.42.A.c”)
- Footnote

Vraiment ?



Agentic search is the future 🚀

Your entire codebase as Claude's context

License MIT | Node.js 20+ | Documentation | VS Code Extension v0.1.3 | @zilliz/clause-context-core v0.1.3
@zilliz/clause-context-mcp v0.1.3 | Follow @Zilliz | DeepWiki AI Docs | DISCORD

Claude Context is an MCP plugin that adds semantic code search to Claude Code and other AI coding agents, giving them deep context from your entire codebase.

💡 Your Entire Codebase as Context: Claude Context uses semantic search to find all relevant code from millions of lines. No multi-round discovery needed. It brings results straight into the Claude's context.

💰 Cost-Effective for Large Codebases: Instead of loading entire directories into Claude for every request, which can be very expensive, Claude Context efficiently stores your codebase in a vector database and only uses related code in context to keep your costs manageable.

Navigateur nouvelle génération

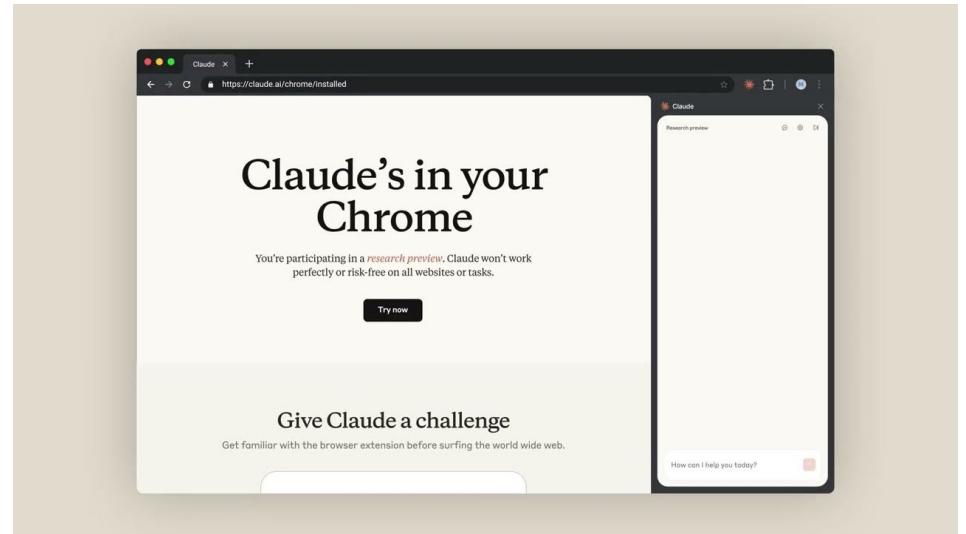
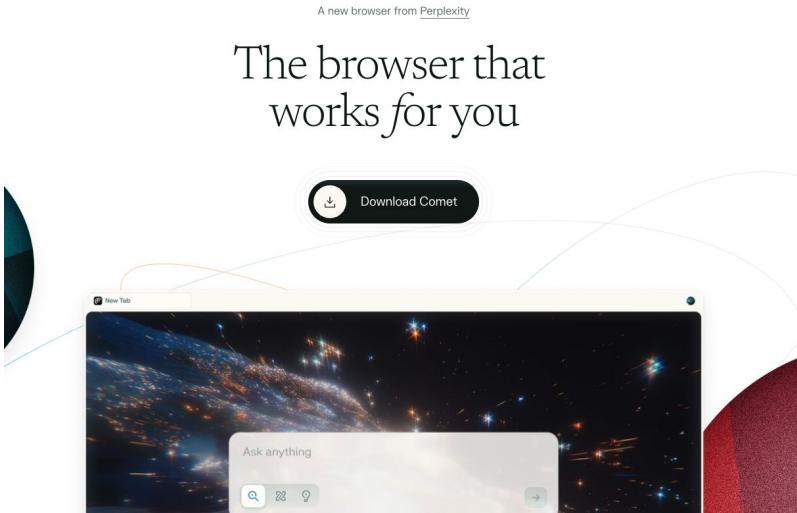


ChatGPT Atlas

AI app store

Apps in ChatGPT

La guerre des navigateurs est lancée



Quid de la sécurité ?

Navigateurs vs opérateurs

manus	ChatGPT
Type & Autonomy	Fully autonomous AI agent; minimal user input needed
Web Interaction	Real-time browsing with live task execution
Task Execution	End-to-end completion—research, reports, bookings
Outputs	Reports, presentations, spreadsheets, bookings
Learning Ability	Learns from tasks and user feedback to improve
Access & Availability	Invite-only preview; open-source coming soon
Strengths	High autonomy, real-time actions, hands-free execution
Weaknesses	Limited access, early-stage development, less tested
Developer	Monica (China)
	OpenAI (USA)
January 2024	November 2022
Release date	

Use case gallery

Learn how Manus handles real-world tasks through step-by-step replays.

Featured Research Life Data Analysis Education Productivity WTF

Trip to Japan in April

Manus not only integrates information for personalized travel planning but also creates a custom handbook for your trip.

Deeply analyze Tesla stocks

Manus performs in-depth stock analysis and designs visually compelling dashboards showcasing results.

Interactive Course on the Momentum Theorem

Manus creates video presentation materials for middle school teachers.

Comparative analysis of insurance policies

Comparing insurance policies? Manus creates clear comparison tables of key information with just a few clicks.

B2B Supplier Sourcing

Manus conducts in-depth research across the entire network to find the most suitable sources for your needs. manus is a truly fair supplier.

Analyze Amazon's financial report

Manus captured market sentiment changes toward Amazon over the past four quarters through research.

List of YC Companies

We needed a list of all B2B companies from YC W25 batch, and Manus efficiently visited the YC website to identify them.

Online store operation analysis

Upload your Amazon store sales data, and Manus delivers actionable insights, detailed visualizations, and recommendations.

Chrome DevTools MCP

npm v0.10.2

`chrome-devtools-mcp` lets your coding agent (such as Gemini, Claude, Cursor or Copilot) control and inspect a live Chrome browser. It acts as a Model-Context-Protocol (MCP) server, giving your AI coding assistant access to the full power of Chrome DevTools for reliable automation, in-depth debugging, and performance analysis.

[Tool reference](#) | [Changelog](#) | [Contributing](#) | [Troubleshooting](#)

Key features

- **Get performance insights:** Uses [Chrome DevTools](#) to record traces and extract actionable performance insights.
- **Advanced browser debugging:** Analyze network requests, take screenshots and check the browser console.
- **Reliable automation.** Uses [puppeteer](#) to automate actions in Chrome and automatically wait for action results.

Expérimentation techniques



Github Action: Claude Code

```
with:
  anthropic_api_key: ${{ secrets.ANTHROPIC_AUTH_TOKEN }}
  use_sticky_comment: true
  track_progress: true
  prompt: |
    REPO: ${{ github.repository }}
    PR NUMBER: ${{ github.event.pull_request.number }}
    AUTHOR: ${{ github.event.pull_request.user.login }}

    Please review this pull request with a focus on:
    - Code quality and best practices
    - Potential bugs or issues
    - Security implications
    - Performance considerations
    - Potential breaking changes
    - Developer experience and documentation

  Check contribution guidelines: ./docs/docs/contributing.md
  Check extending ro guidelines: ./docs/docs/hacking.md
  Check examples/documentation guidelines: ./docs/CLAUDE.md
  Check troubleshooting guidelines: ./docs/docs/troubleshooting/
  If you need more context on the project, read the LLMs documentation: ./docs/static/llms.txt

Note: The PR branch is already checked out in the current working directory.
```

Github Action: Claude Code

Note: The PR branch is already checked out in the current working directory.

Use `mcp__github_inline_comment__create_inline_comment` to highlight specific code issues.
Be welcoming but thorough in your comments. Use inline comments for code-specific feedback.

Write short and concise comments/reviews. Write comments only if there is a specific issue: don't waste my time commenting on the good stuff. If you have nothing to say, don't say anything. Before posting a new comment (inline or in the PR thread), make sure it doesn't repeat a point already made previously in comments.

```
# security: a prompt injection can be used to update comments in the current PR.  
claude_args: |  
  --mcp-config '{"mcpServers": {"sequential-thinking": {"command": "npx", "args": ["-y", "@modelcontextprotocol/server-sequential-thinking"]}}}  
  --allowedTools "mcp__github_inline_comment__create_inline_comment,mcp__sequential-thinking__sequentialthinking,Bash(gh pr comment:*)",Bash
```

Github Action: Claude Code



claude bot reviewed last month

[View reviewed changes](#)

subscriber.go

Hide resolved

```
149 +     status      int32
150 +     backpressure Backpressure
151 +
152 +     _ [59]byte // padding to prevent false sharing
```



claude bot last month

...

The 59-byte padding (4 bytes for int32 + 59 = 63 bytes total) is just shy of a full cache line. Consider making it 60 bytes to reach 64 bytes total, which is the typical cache line size on most modern CPUs.



Reply...

[Unresolve conversation](#)

samber marked this conversation as resolved.

Spec-driven development

Th
r]

task-master-ai — ~/terminal

↳ Weekly: 38,515 ↑ 14% | Monthly: 138,564 ↑ 10% | 🌐 Active: ~34,641



Star 23.2k

works with



TASKMASTER AI

Stuck in AI loops?

"Absolutely insane. I built a multiplayer game in one go with zero errors." |

Taskmaster is your AI's personal project manager. Organize, research, expand, prioritize, and ship tasks effortlessly. Enjoy permanent context, zero drift & instant clarity. It's free, open-source & API-friendly. Let's vibe productively.

▶ Taskmaster Demo

⬇ Install Instructions

↗ Early Access

Try typing a command:

```
$ help          $ features        $ commands  
$ upvote       $ feedback        $ early-access
```

\$ Enter 'help' to see available commands

Type [help](#) to see commands or [home](#) to return to the main menu

[README](#)[Code of conduct](#)[Contributing](#)[MIT license](#)[Security](#)

Spec Kit

Build high-quality software faster.

An open source toolkit that allows you to focus on product scenarios and predictable outcomes instead of
vibe coding every piece from scratch.

[Create Release](#)

passing

[Stars](#)

50k

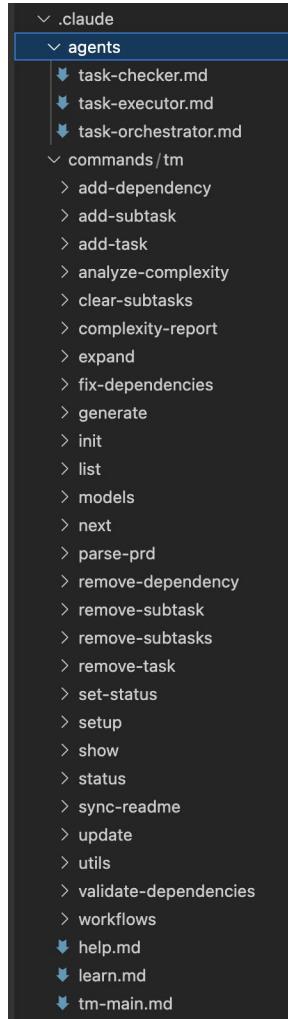
[license](#)

MIT

[docs](#)[GitHub Pages](#)

Installation Options

- ▶ Option 1: MCP (Recommended)
- ▶ Option 2: Using Command Line



\$ task-master init

Technical Specification:
github.com/samber/so v1.0
...

T Technical specification writing • Custom Gem

Here's my plan to tackle that topic. Let me know if you need to make changes.

Go Stream Processing Library

Research Websites

Technical Specification: github.com/samber/so v1.0

Project Name: github.com/samber/so

Project Goal: To provide a robust, flexible, and idiomatic Go library for building stateful and stateless stream processing...

More

Analyze Results

Create Report

Ready in a few mins

Edit plan

Start research

What do you want to research?

+ - Deep Research X

v

mic



Go Stream Processing Library

Contents

Share & Export

Create

X

Technical Assessment of the github.com/samber/so v1.0 Specification

I. Introduction

Purpose: This report provides an expert technical assessment of the version 1.0 specification for the github.com/samber/so library. The analysis evaluates the proposed design, architecture, features, feasibility, and potential challenges based exclusively on the provided technical specification document and supplementary research materials.

Scope: The assessment is strictly confined to the details presented in the [samber/so](https://github.com/samber/so) v1.0 specification document. It analyzes the design's technical soundness, identifies potential risks, complexities, and trade-offs inherent in the proposed approach, and evaluates the feasibility of achieving the stated goals and features within the v1.0 scope. This report assesses the *design document* itself, not an existing implementation.

Methodology: The evaluation involves a systematic analysis of the specified architectural components, including the core abstractions, state management, coordination mechanisms, and transactional protocols. The proposed mechanisms are evaluated against established distributed systems principles and insights derived from the provided research materials. Key risks and complexities are identified, and the analysis directly addresses the specific points of inquiry raised by the user regarding the project's goals, architecture, scalability, fault tolerance, processing guarantees (particularly Exactly-Once Semantics via Two-Phase Commit), connector capabilities and limitations (with a focus on HTTP), inherent technical challenges, and stipulated development requirements.

II. Project Vision and Positioning

Analysis of Core Goal, Target Audience, and Differentiators:

The stated core goal of the [samber/so](https://github.com/samber/so) project is "To provide a robust, flexible, and idiomatic Go library for building stateful and stateless stream processing applications." This goal is ambitious, aiming to deliver a comprehensive stream processing solution within the Go ecosystem. The specification targets Go developers who seek native tooling for stream processing.

Claude Code v2.0.28

Welcome back!



glm-4.6 · API Usage Billing
/Users/samber/project/github.com/samber/so-spec

Recent activity

No recent activity

What's new

Plan mode: introduced new Plan subagent

Subagents: claude can now choose to resume subagents

Subagents: claude can dynamically choose the model used by its subagents

/release-notes for more

> Split the project described in `@.taskmaster/docs/GEMINI-2.5-pro.md` into 10 big milestones and write the PR in `.taskmaster/docs/`
└ Read `.taskmaster/docs/GEMINI-2.5-pro.md` (345 lines)

- I'll analyze the technical assessment document and split the project into 10 big milestones, then create a PRD for Task Master.

+ Reading and analyzing the technical assessment document... (esc to interrupt + ctrl+t to hide todos)

- └
 - Read and analyze the technical assessment document
 - Extract key components and phases from the assessment
 - Create 10 big milestones for the project
 - Write PRD document with the milestones
 - Parse PRD with Task Master to generate tasks

⌄	.taskmaster	
⌄	docs	
↳	000-global-project-instru...	
↳	001-mvp-foundation... M	
↳	001.5-dead-letter-q... M	
↳	002-state-managem... M	
↳	003-windowing-ope... M	
↳	004-universal-coord... M	
↳	005-multi-broker-cr... M	
↳	006-advanced-featu... M	
↳	007-ecosystem-exp... M	
↳	008-production-har... M	
↳	GEMINI-2.5-pro.md	
>	reports	
>	tasks	
>	templates	
↳	CLAUDE.md	
{}	config.json	M
{}	state.json	M

```
tm parse-prd -i .taskmaster/docs/001-mvp-foundation-prd.md -n 40 --append
```

The screenshot shows a terminal window with the following content:

```
by https://x.com/eyaltoledano
Version: 0.27.3 Project: so

tag: master
Parsing PRD file: /Users/samber/project/github.com/samber/so-spec/.taskmaster/docs/001-mvp-foundation-prd.md
generating 40 tasks...
Appending to existing tasks...
[DEBUG] Parsing PRD file: /Users/samber/project/github.com/samber/so-spec/.taskmaster/docs/001-mvp-foundation-prd.md, Force: false, Append: true,
[INFO] Tag 'master' is empty or doesn't exist. Creating/updating tag with new tasks.
[DEBUG] ✓ JSON schema validation enabled
- Parsing PRD and generating tasks...
[INFO] Calling AI service to generate tasks from PRD...
[INFO] generateObjectService called {"role": "main", "commandName": "parse-prd", "outputType": "cli", "projectRoot": "/Users/samber/project/github.com/samber/so-spec/.taskmaster/docs/001-mvp-foundation-prd.md", "modelUsed": "gemini-2.5-pro", "maxTokens": 65536}
[DEBUG] New AI service call with role: main
[DEBUG] Applying model-specific max_tokens (65536) for gemini-2.5-pro. Effective limit: 65536
[INFO] Attempt 1/3 calling generateObject (Provider: gemini-cli, Model: gemini-2.5-pro, Role: main)
- Parsing PRD and generating tasks...
[DEBUG] Gemini CLI generateObject completed successfully for model: gemini-2.5-pro
[INFO] generateObject succeeded for role.main (Provider: gemini-cli) on attempt 1
[INFO] AI Usage Telemetry: {"timestamp": "2025-10-28T04:53:41.142Z", "userId": "1234567890", "commandName": "parse-prd", "modelUsed": "gemini-2.5-pro", "tokensGenerated": 40, "tokensUsed": 40}
✓ Tasks generated successfully!
[DEBUG] Successfully saved 40 tasks to /Users/samber/project/github.com/samber/so-spec/.taskmaster/tasks/tasks.json

Successfully generated 40 new tasks. Total tasks in /Users/samber/project/github.com/samber/so-spec/.taskmaster/tasks/tasks.json: 40
```

Next Steps:

1. Run `task-master list` to view all tasks
2. Run `task-master expand --id=<id>` to break down a task into subtasks

tm list

```
>tag: master
Listing tasks from: /Users/samber/project/github.com/samber/so-spec/.taskmaster/tasks/tasks.json
```

Project Dashboard

Tasks Progress:  0% 0/40

Done: 0 Cancelled: 0 Deferred: 0

In Progress: 0 Review: 0 Pending: 40 Blocked: 0

Subtasks Progress:  0% 0/8

Completed: 0/8 Cancelled: 0 Deferred: 0

In Progress: 0 Review: 0 Pending: 8 Blocked: 0

Priority Breakdown:

- High priority: 20
- Medium priority: 17
- Low priority: 3

Dependency Status & Next Task

Dependency Metrics:

- Tasks with no dependencies: 1
- Tasks ready to work on: 1
- Tasks blocked by dependencies: 39
- Most depended-on task: #2 (9 dependents)
- Avg dependencies per task: 1.4

Next Task to Work On:

ID: 1 - Initialize Project Structure and Core Definitions

Priority: high Dependencies: None

Complexity: N/A

ID	Title	Status	Priority	Dependencies	Complexity
1	Initialize Project Structure and Core Definitions	pending	high	None	N/A
2	Define Core Component Interfaces	pending	high	1	N/A
3	Implement DSL Builder and Topology Interfaces	pending	high	2	N/A
4	Implement Core Runtime Executor	pending	high	3	N/A
5	Implement Runtime Task Management	pending	high	4	N/A
6	Implement In-Memory State Store	pending	high	2	N/A
7	Implement All Map Operator Variants	pending	medium	2	N/A
8	Implement All Filter Operator Variants	pending	medium	2	N/A
9	Implement All FlatMap Operator Variants	pending	medium	2	N/A
10	Implement Stateful Count Operator	pending	medium	6, 9	N/A

tm analyze-complexity

```
>tag: master
Analyzing task complexity from: /Users/samber/project/github.com/samber/so/.taskmaster/tasks/tasks.json
Output report will be saved to: /Users/samber/project/github.com/samber/so/.taskmaster/reports/task-complexity-report.json
Analyzing task complexity and generating expansion recommendations...
[INFO] Reading tasks from /Users/samber/project/github.com/samber/so/.taskmaster/tasks/tasks.json...
[INFO] Found 40 total tasks in the task file.
AI service call complete. Parsing response...
[INFO] Parsing complexity analysis from text response...
[WARN] Missing analysis for 30 tasks: 3, 4, 6, 7, 8, 9, 14, 15, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 30, 31, 32, 33, 35, 36, 37, 38, 39, 40
Missing analysis for 30 tasks: 3, 4, 6, 7, 8, 9, 14, 15, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 30, 31, 32, 33, 35, 36, 37, 38, 39, 40
[INFO] Adding default analysis for task 3
[INFO] Adding default analysis for task 4
[INFO] Adding default analysis for task 6
[INFO] Adding default analysis for task 7
[INFO] Adding default analysis for task 8
[INFO] Adding default analysis for task 9
[INFO] Adding default analysis for task 14
[INFO] Adding default analysis for task 15
[INFO] Adding default analysis for task 17
[INFO] Adding default analysis for task 18
[INFO] Adding default analysis for task 19
[INFO] Adding default analysis for task 20
[INFO] Adding default analysis for task 21
[INFO] Adding default analysis for task 22
[INFO] Adding default analysis for task 23
[INFO] Adding default analysis for task 24
[INFO] Adding default analysis for task 25
[INFO] Adding default analysis for task 26
[INFO] Adding default analysis for task 27
[INFO] Adding default analysis for task 28
[INFO] Adding default analysis for task 30
[INFO] Adding default analysis for task 31
[INFO] Adding default analysis for task 32
[INFO] Adding default analysis for task 33
[INFO] Adding default analysis for task 35
[INFO] Adding default analysis for task 36
[INFO] Adding default analysis for task 37
[INFO] Adding default analysis for task 38
[INFO] Adding default analysis for task 39
[INFO] Adding default analysis for task 40
[INFO] Writing complexity report to /Users/samber/project/github.com/samber/so/.taskmaster/reports/task-complexity-report.json...
[SUCCESS] Task complexity analysis complete. Report written to /Users/samber/project/github.com/samber/so/.taskmaster/reports/task-complexity-report.json
Task complexity analysis complete. Report written to /Users/samber/project/github.com/samber/so/.taskmaster/reports/task-complexity-report.json

Current Analysis Summary:
-----
Tasks analyzed in this run: 40
High complexity tasks: 0
Medium complexity tasks: 32
Low complexity tasks: 8
Research-backed analysis: No
```

tm complexity-report

Complexity Distribution

Low (1-4): 8 tasks (20%)
Medium (5-7): 32 tasks (80%)
High (8-10): 0 tasks (0%)

ID	Title	Score	Subta...	Expansion Command
34	Implement Kafka Offset Management Strategy	● 7	5	task-master expand --id=34 --num=5 --prompt="Expand this task to implement advanced Kafka offset management. Subtasks include 1. Implementing the logic for 'Seek' and 'CommitInterval' to 'KafkaSourceConfig'. 2. Implementing internal state in the 'KafkaSource' to buffer pending offsets. 3. Implementing commit logic within the 'Poll' method. 4. Implementing the time-based commit logic using a background goroutine and a final commit."
5	Implement Runtime Task Management	● 5	5	task-master expand --id=5 --num=5 --prompt="Expand this task into subtasks for implementing the runtime task. Subtasks include 1. Implementing the 'NewTask' constructor. 2. Implementing the 'Run' method (operator -> sink). 3. Implementing error handling. 4. Creating unit tests using mock operators."
3	Implement DSL Builder and Topology Interfaces	● 5	3	task-master expand --id=3 --num=3 --prompt="Break down this task with a focus on implement dsl builder and topology interfaces."
4	Implement Core Runtime Executor	● 5	3	task-master expand --id=4 --num=3 --prompt="Break down this task with a focus on implement core runtime executor."
6	Implement In-Memory State Store	● 5	3	task-master expand --id=6 --num=3 --prompt="Break down this task with a focus on implement in-memory state store."
7	Implement All Map Operator Variants	● 5	3	task-master expand --id=7 --num=3 --prompt="Break down this task with a focus on implement all map operator variants."
8	Implement All Filter Operator Variants	● 5	3	task-master expand --id=8 --num=3 --prompt="Break down this task with a focus on implement all filter operator variants."
9	Implement All FlatMap Operator Variants	● 5	3	task-master expand --id=9 --num=3 --prompt="Break down this task with a focus on implement all flatmap operator variants."
14	Implement Simple Mock Connectors (stdin/stdout)	● 5	3	task-master expand --id=14 --num=3 --prompt="Break down this task with a focus on implement simple mock connectors."
15	Implement Rich Mock Connectors (In-Memory)	● 5	3	task-master expand --id=15 --num=3 --prompt="Break down this task with a focus on implement rich mock connectors." (In-Memory)

```
tm expand --all --research --complexity-report .taskmaster/reports/complexity-report.json
```

```
💡 Telemetry
```

```
AI Usage Summary:
Command: expand-task
Provider: gemini-cli
Model: gemini-2.5-pro
Tokens: 2172 (Input: 1210, Output: 962)
Est. Cost: $0.000000
```

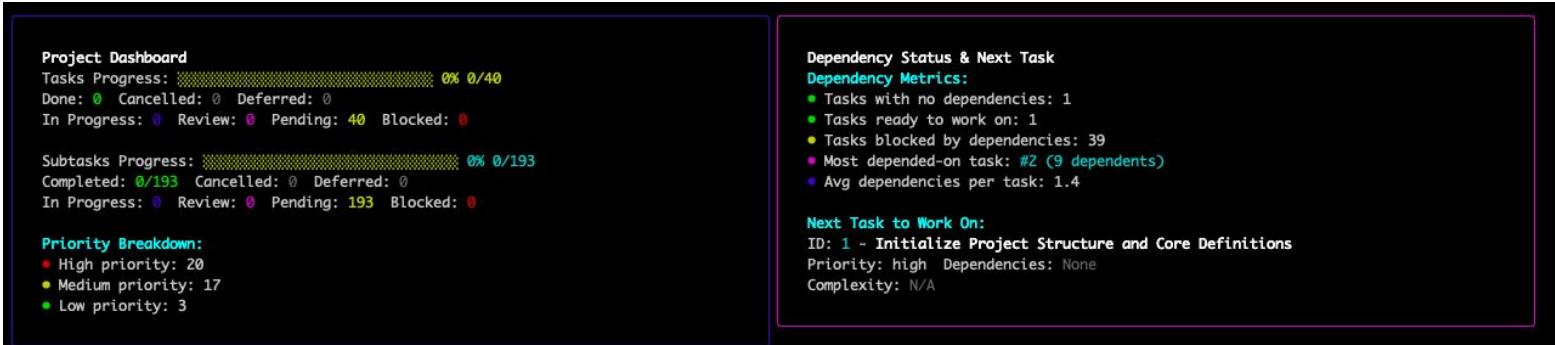
```
[INFO] Successfully expanded task 40.
[INFO] Expansion complete: 39 expanded, 1 failed.
```

```
Expansion Summary:
- Attempted: 40
- Expanded: 39
- Skipped: 0
- Failed: 1
```

```
[DEBUG] User prompt (first 500 chars): ## IMPORTANT: Codebase Analysis Required
You have access to powerful codebase analysis tools. Before generating subtasks:
1. Use the Glob tool to explore relevant files for this task (e.g., "**/*.js", "src/
2. Use the Grep tool to search for existing implementations related to this task
3. Use the Read tool to examine files that would be affected by this task
4. Understand the current implementation state and patterns used

Based on your analysis:
```

tm list --with-subtasks



ID	Title	Status	Priority	Dependencies	Complexity
1	Initialize Project Structure and Core Definitions	o pending	high	None	N/A
└ 1	Create Top-Level Project Directories	o pending	medium	None	--
└ 2	Create Core Package and Define Enums	o pending	medium	None	--
└ 3	Define the Generic Record Struct	o pending	medium	None	--
└ 4	Create Placeholder Go Files for All Packages	o pending	medium	None	--
└ 5	Initialize Go Module and Verify Build	o pending	medium	None	--
2	Define Core Component Interfaces	o pending	high	1	N/A
└ 1	Define Base Component Interface with Lifecycle Hooks	o pending	medium	None	--
└ 2	Define the Generic Source Interface	o pending	medium	2.1	--
└ 3	Define the Generic Sink Interface	o pending	medium	2.1	--
└ 4	Define the Generic Operator Interface	o pending	medium	2.1	--
└ 5	Define the StateStore Interface	o pending	medium	None	--
3	Implement DSL Builder and Topology Interfaces	o pending	high	2	N/A
└ 1	Implement the DSL Builder API	o pending	medium	None	--
└ 2	Implement the Topology API	o pending	medium	None	--
└ 3	Integrate with the StateStore	o pending	medium	None	--
└ 4	Handle Complex Topology Transformations	o pending	medium	None	--
└ 5	Optimize Performance for Large Topologies	o pending	medium	None	--



Claude Code v2.0.28

glm-4.6 · API Usage Billing

/Users/samber/project/github.com/samber/so

> Implement the next task.

- I'll help you implement the next task. Let me first check if there's a Task Master setup in this project and see what tasks are available.

task-master-ai - get_tasks (MCP)(projectRoot: "/Users/samber/project/github.com/samber/so")

```
└ {  
    "data": {  
        "tasks": [  
            ... +1051 lines (ctrl+o to expand)  
    }  
}
```

- Now let me check what the next task to work on should be:

task-master-ai - next_task (MCP)(projectRoot: "/Users/samber/project/github.com/samber/so")

```
└ {  
    "data": {  
        "nextTask": {  
            ... +114 lines (ctrl+o to expand)  
    }  
}
```

- Perfect! I can see the next task is to initialize the project structure and core definitions. Let me set up a todo list and start working on this task.

- Let me start by marking the first task as in progress and setting the main task as in progress as well:

task-master-ai - set_task_status (MCP)(id: "1", projectRoot: "/Users/samber/project/github.com/samber/so", status: "in-progress")

```
└ {  
    "data": {  
        "message": "Successfully updated task 1 status to \"in-progress\""  
    }  
    ... +15 lines (ctrl+o to expand)  
}
```

+ Creating project directory structure... (esc to interrupt · ctrl+t to hide todos · 23s · + 713 tokens)

```
└ □ Create Project Directory Structure  
  □ Initialize Go Module and Basic Files  
  □ Define Core Data Structures and Enums  
  □ Create Core Interface Definitions  
  □ Create Package Placeholder Files  
  □ Create CLI Tool Placeholders  
  □ Create Example Directory Structure  
  □ Verify Project Structure and Initial Build
```

“Implement the next task”

Project Dashboard

Tasks Progress:  3% 1/40

Done: 1 Cancelled: 0 Deferred: 0

In Progress: 0 Review: 0 Pending: 39 Blocked: 0

Subtasks Progress:  15% 8/52

Completed: 8/52 Cancelled: 0 Deferred: 0

In Progress: 0 Review: 0 Pending: 44 Blocked: 0

Priority Breakdown:

- High priority: 20
- Medium priority: 17
- Low priority: 3

Dependency Status & Next Task

Dependency Metrics:

- Tasks with no dependencies: 0
- Tasks ready to work on: 2
- Tasks blocked by dependencies: 37
- Most depended-on task: #2 (9 dependents)
- Avg dependencies per task: 1.4

Next Task to Work On:

ID: 2 - Define Core Component Interfaces

Priority: high Dependencies: 1

Complexity: N/A

ID	Title	Status	Priority	Dependencies	Complexity
1	Initialize Project Structure and Core Definitions	✓ done	high	None	N/A
└ 1	Create Project Directory Structure	✓ done	medium	None	--
└ 2	Initialize Go Module and Basic Files	✓ done	medium	None	--
└ 3	Define Core Data Structures and Enums	✓ done	medium	None	--
└ 4	Create Core Interface Definitions	✓ done	medium	None	--
└ 5	Create Package Placeholder Files	✓ done	medium	None	--
└ 6	Create CLI Tool Placeholders	✓ done	medium	None	--
└ 7	Create Example Directory Structure	✓ done	medium	None	--
└ 8	Verify Project Structure and Initial Build	✓ done	medium	None	--
2	Define Core Component Interfaces	○ pending	high	1	N/A
└ 1	Define Source Interface with Generic Types	○ pending	medium	None	--
└ 2	Define Sink Interface with Generic Types	○ pending	medium	None	--

CLI vs MCP

Don't create too many tasks/subtasks up-front.

Iterate.

Claude Code used

mcp_tm_create_task(id, title, description)

instead of

mcp_tm_expand_task(id)

Learnings

Abandoned

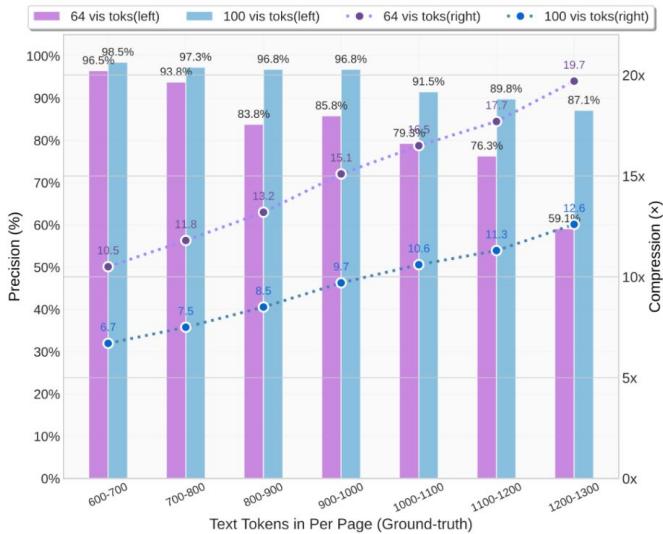
Task status tracking is useless

Just use flat files + Coding Agent

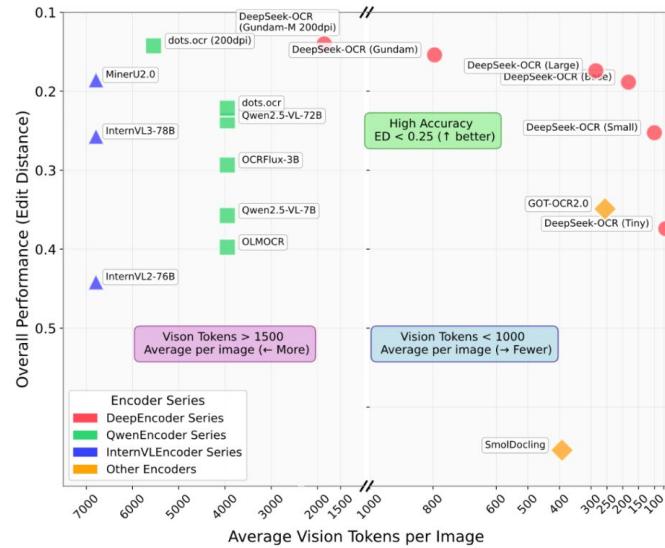
Really loved working on flat files vs traditional plans

Deepseek-OCR

DeepSeek-OCR: Contexts Optical Compression



(a) Compression on Fox benchmark



(b) Performance on Omnidocbench

Explore the boundaries of visual-text compression.

Deepseek-OCR



Samuel Berthe @SamuelBerthe · Nov 10



Deepseek, October 2025:

"Hey guys, look at our new deepseek-ocr model!"

Deepseek CTO in 6 months 🤝

DID YOU
USE GOOGLE
BOOKS ?



Deepseek-OCR

vs

Mistral-OCR

Self-hosting vs cloud

<https://www.alphxiv.org/models/deepseek/deepseek-ocr>

Deepseek-OCR

vs

Mistral-OCR

Conclusion

The benchmark analysis between DeepSeek and Mistral OCR systems reveals significant insights that are crucial for stakeholders considering these solutions for their optical character recognition needs. Through rigorous, multi-metric evaluations utilizing diverse real-world datasets, a comprehensive understanding of each system's performance was achieved.

DeepSeek OCR demonstrated a marginally superior performance in terms of Character Error Rate (CER), achieving an average CER of 3.2% across datasets like Fox and OmniDocBench. Meanwhile, Mistral OCR was not far behind, with a CER of 3.8%. This slight edge suggests DeepSeek's enhanced capability in recognizing subtle text variations, particularly in complex cursive and printed handwriting styles.

However, when considering system efficiency and speed, Mistral OCR showed a 20% faster processing time, which is a significant advantage in high-volume document processing environments. Despite its slightly higher error rates, Mistral's efficiency may present a compelling case for businesses where speed is critical.

Mistral-OCR & annotations

PYTHON TYPESCRIPT CURL JSON SCHEMA

```
from pydantic import BaseModel, Field

# BBOX Annotation response formats
class Image(BaseModel):
    image_type: str = Field(..., description="The type of the image.")
    short_description: str = Field(..., description="A description in english describing the image.")
    summary: str = Field(..., description="Summarize the image.")
```

Returns an output formatted
in markdown

PYTHON TYPESCRIPT CURL

```
import os
from mistralai import Mistral, DocumentURLChunk, ImageURLChunk, ResponseFormat
from mistralai.extra import response_format_from_pydantic_model

api_key = os.environ["MISTRAL_API_KEY"]

client = Mistral(api_key=api_key)

response = client.ocr.process(
    model="mistral-ocr-latest",
    document=DocumentURLChunk(
        document_url="https://arxiv.org/pdf/2410.07073"
    ),
    bbox_annotation_format=response_format_from_pydantic_model(Image),
    include_image_base64=True
)
```

Mistral-OCR

Up to 100 pages :-(

Footnotes KO

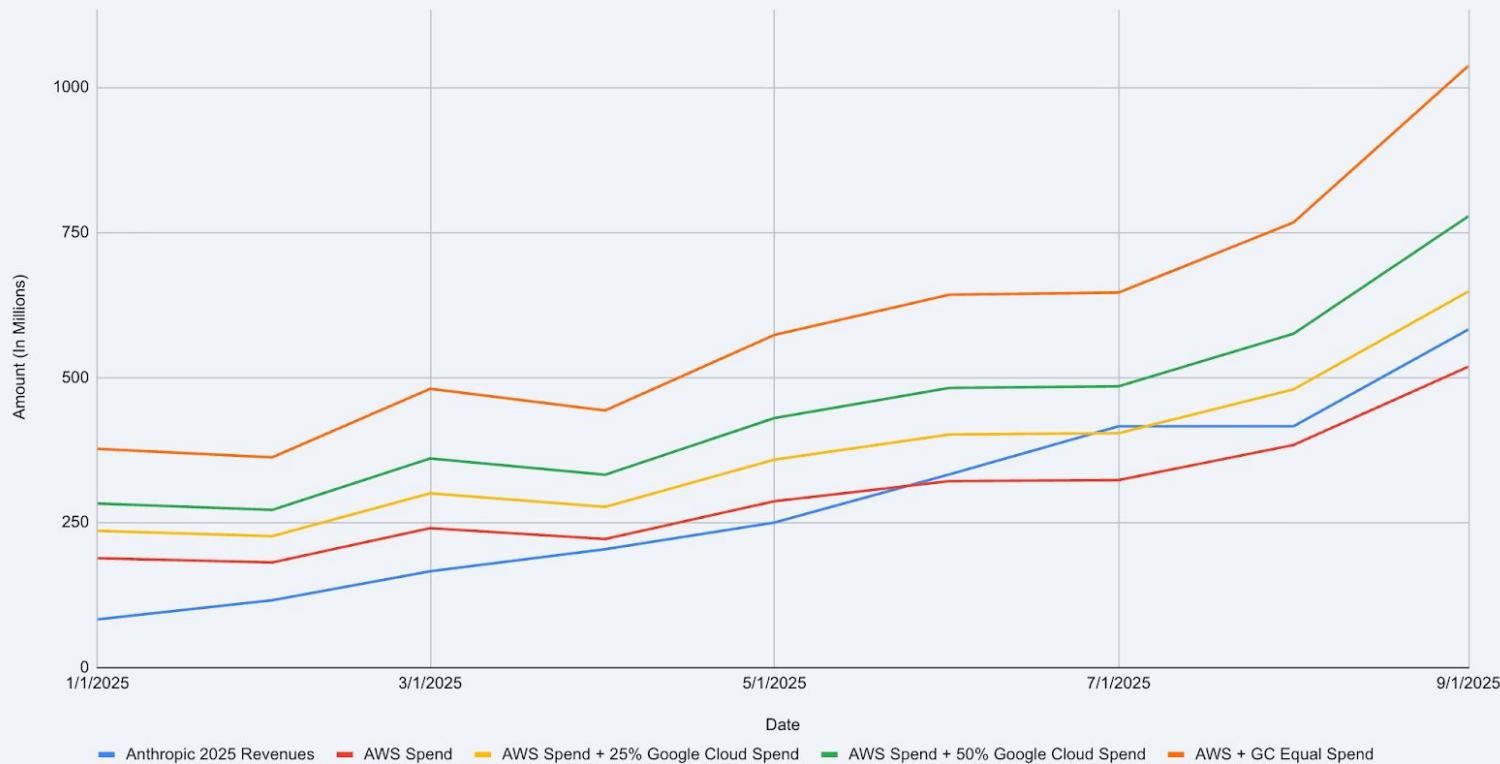
Limited annotations

Cost of Anthropic + Cursor

<https://www.wheresyoured.at/costs/>

Anthropic

Anthropic's Revenue Versus Cloud Spend Through September 2025



Anthropic

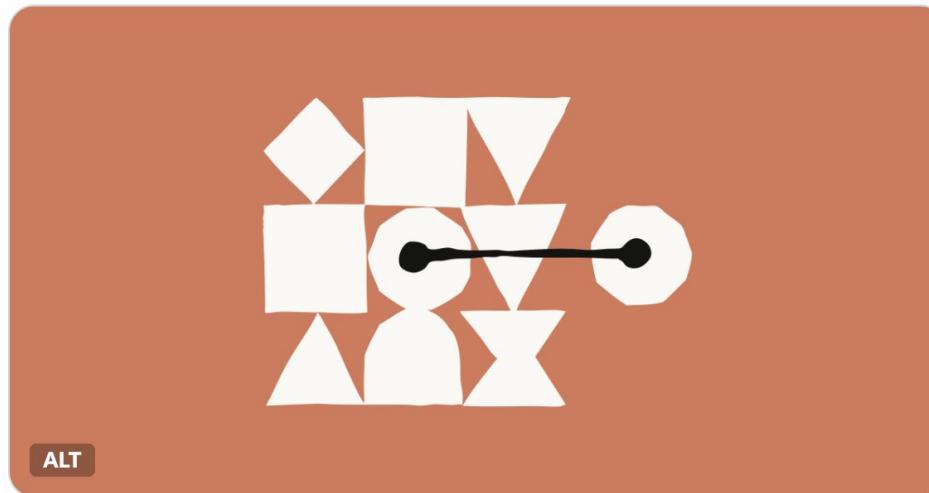
← Post



Anthropic @AnthropicAI

...

We're rolling out new weekly rate limits for Claude Pro and Max in late August. We estimate they'll apply to less than 5% of subscribers based on current usage.



8:23 PM · Jul 28, 2025 · 2.3M Views

Cursor

Mai 2025 : ~ 6,19 M\$

Juin 2025 : ~ **12,67 M\$**

Juillet : ~ **15,5 M\$**

Août : ~ 9,67 M\$

Septembre : ~ 12,91 M\$

Cursor

JUNE 16, 2025

Introducing the Ultra Plan

In collaboration with the model providers, we're introducing a \$200 / mo tier for power users.



Posted By Michael

1 minutes read

We're excited to roll out an option to purchase Ultra, a \$200 / mo plan with 20x more usage than Pro.

While the vast majority of Cursor users are well-served by our Pro plan, this change was highly requested by power users seeking more predictability than usage-based pricing would offer.

Ultra is made possible by multi-year partnerships from OpenAI, Anthropic, Google, and xAI. Their support was instrumental in offering this volume of compute at a predictable price.

Cursor

Series C and Scale

Jun 6, 2025 by Cursor Team

Today, we're announcing new funding to improve Cursor, \$900 million at a \$9.9 billion valuation from Thrive, Accel, Andreessen Horowitz, and DST.

We're also happy to share that Cursor has grown to over \$500 million in ARR and is used by over half of the Fortune 500, including NVIDIA, Uber, and Adobe.

This scale will help us push the frontier of AI coding research.

We started Anysphere, the lab behind Cursor, to build a better way to code. We're excited that many of the pieces are in place for us to execute on this. ■

Past, Present, and Future

Nov 13, 2025 by Cursor Team

When we raised our seed round nearly two years ago, we wrote:

In the next few years, we'd like to build a code editor that is more helpful, delightful, and fun than the world has ever seen. Cursor should be a place where it's impossible to write bugs. An editor where you whip up 2,000-line PRs with 50 lines of pseudo code. A tool where you get any codebase question answered instantly. Perhaps even an interface where the source code itself starts to melt away.

It's been exciting to see the early innings of this vision take shape.

We're obsessed with the magical moments in the history of programming with AI. Internally, we often talk about how high the ceiling is for how great Cursor can become, and how much work still remains to get there.

Today, we're pleased to announce a new round of financing: our Series D of \$2.3B at a \$29.3B post-money valuation. We're excited to deepen our work with existing investors, including Accel, Thrive, Andreessen Horowitz, and DST, and welcome new partners Coatue, NVIDIA, and Google.

Hardware

ASIC: GCP

The image shows the Google Cloud navigation bar. The products section is highlighted with a blue underline. Other menu items include Overview, Solutions, Pricing, and Resources. To the right are search, documentation, support, console, and user profile icons. A 'Contact Us' button is also present.

Cloud Tensor Processing Units (TPUs)

[Overview](#)

Options

How It Works

Common Uses

Run large-scale AI training workloads

Fine-tune foundational AI models

Serve large-scale AI inference workloads

Cloud TPU in GKE

Cloud TPU in Vertex AI

Generate a solution

Pricing

[Cloud TPU tutorials, quickstarts and docs](#)

Cloud TPU versions

Cloud TPU version	Description	Availability
Ironwood	Our most powerful and efficient TPU yet, for the largest scale training and inference	Ironwood TPU will be generally available in Q4, 2025
Trillium	Sixth-generation TPU. Improved energy efficiency and peak compute performance per chip for training and inference	Trillium is generally available in North America (US East region), Europe (West region), and Asia (Northeast region)
Cloud TPU v5p	Powerful TPU for building large, complex foundational models	Cloud TPU v5p is generally available in North America (US East region)

ASIC: AWS

AWS Trainium

Get high performance for deep learning and generative AI training while lowering costs

Get started with Trainium using AWS Neuron

Why Trainium?

AWS Trainium chips are a family of AI chips purpose built by AWS for AI training and inference to deliver high performance while reducing costs.

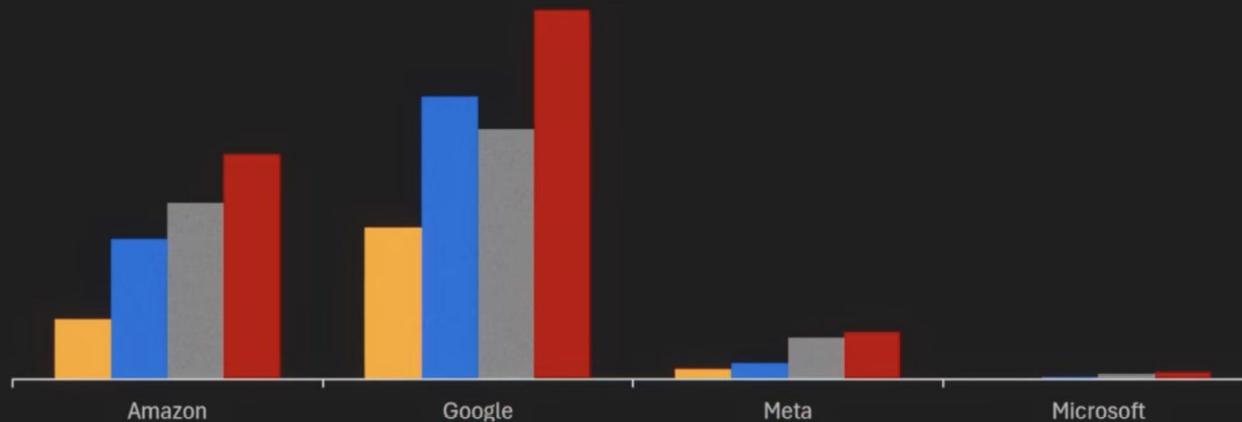
The first-generation AWS Trainium chip powers [Amazon Elastic Compute Cloud \(Amazon EC2\) Trn1 instances](#), which have up to 50% lower training costs than comparable Amazon EC2 instances. Many customers, including Databricks, Ricoh, NinjaTech AI, and Arcee AI, are realizing performance and cost benefits of Trn1 instances.

AWS Trainium2 chip delivers up to 4x the performance of first-



UNIT SHIPMENTS FOR HYPERSCALER AL ASIC

■ 2023 ■ 2024 ■ 2025 ■ 2026



Google lifetime TPU shipments: 5-7 million units

Amazon lifetime Trainium shipments: 3-5 million units



Microsoft

Microsoft CEO says the company doesn't have enough electricity to install all the AI GPUs in its inventory - 'you may actually have a bunch of chips sitting in inventory that I can't plug in'

News





“Prompt to Video: comment l'IA redéfinit la publicité”



Yannis Sulfourt
Designer et formateur

Maxime Bodereau
Art Director @ Seelab



19 Novembre 2025 à 19h00
54 Bd de Berlin, 44000 Nantes

chez



[sfɛir]

lonestone

Slides dispo sur:

<https://github.com/genai-nantes-meetup/meetups/>



SHIFT

Le Hackathon Gen AI

27 au 29 Mars

