



# “Comment créer une équipe d'agents IA en 2026”



Globetrotters.ai



**Benoit Hediard**

Founder @Globetrotters.ai | Ex-CTO &  
Co-Founder @Agorapulse | Angel Investor



15 janvier 2026 à 19h00



4 Rue Voltaire, 44000 Nantes

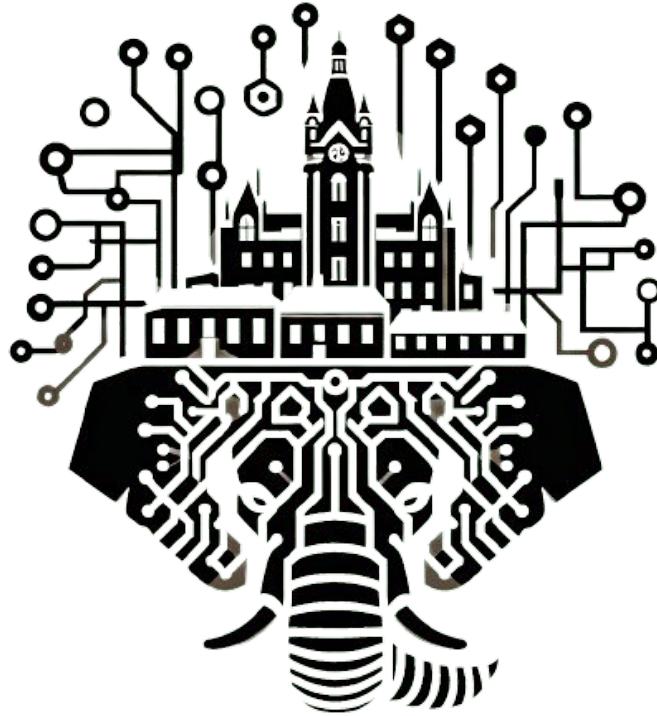
chez [icilundi](#)



techtown lonestone RCA  
externatic zenika [sfɛir]

# GenAI Nantes

- Microphone icon: 15 événements / an
- Pirate flag icon: 1 hackathon (Shift)
- Construction worker icon: 2 workshops
- Heart icon: 1 communauté de 1000p\*



\* 10.000p selon le syndicat des llamas 

**Qui vient pour la première fois ?** 😍



## You have won the Super Organizer badge – congratulations!

You have been recognized as one of the organizers who have generated the highest levels of participation and community engagement on Meetup in recent months.

Your profile, your avatar, your groups and your events will now display the Super Organizer badge – highlighting your impact as an exceptional community leader. Incredible work!

C'est parti

The image shows a Meetup profile card for "GenAI Nantes". The card features a large circular badge in the center containing a stylized illustration of a city skyline with circuit board patterns. To the right of the badge is a smaller blue circular badge with a white gear and person symbol. Below the badge, the group name "GenAI Nantes" is displayed in bold black text. Underneath the group name, there are two lines of information: "Nantes, FR" with a location pin icon, and "Membre Meetup depuis : avr. 2025" with a calendar icon. At the bottom of the card, a blue button-like element says "Super Organisateur". In the bottom right corner of the card, there is a smaller callout bubble containing the text "Animé par GenAI N. Super Organisateur" next to a small profile picture.

# Merci aux partenaires ! ❤

💛 Gold



❤ Silver

lonestone

techtown

[sfɛir]

z zenika

x externalic



ÉDITEUR DE LOGICIELS  
POUR LES EXPERTS-COMPTABLES

**SHIFT**

**TIME  
TO  
CUSTOM**

27 - 29 Mars  
2026



**Qui est codeur ?**



**Qui est no-codeur ?**



**Autre ?**



**Qui a déjà utilisé un agent ?** 🍪

**Qui a déjà créé un agent ?** 

**Qui recrute dans la GenAI ?** 

**Qui cherche un job dans la GenAI ?**



**Qui souhaite s'associer dans la GenAI ?** 

# Schedule

1- News

2- Comment créer une équipe d'agents IA en 2026

3- Enjoy

# **News !**

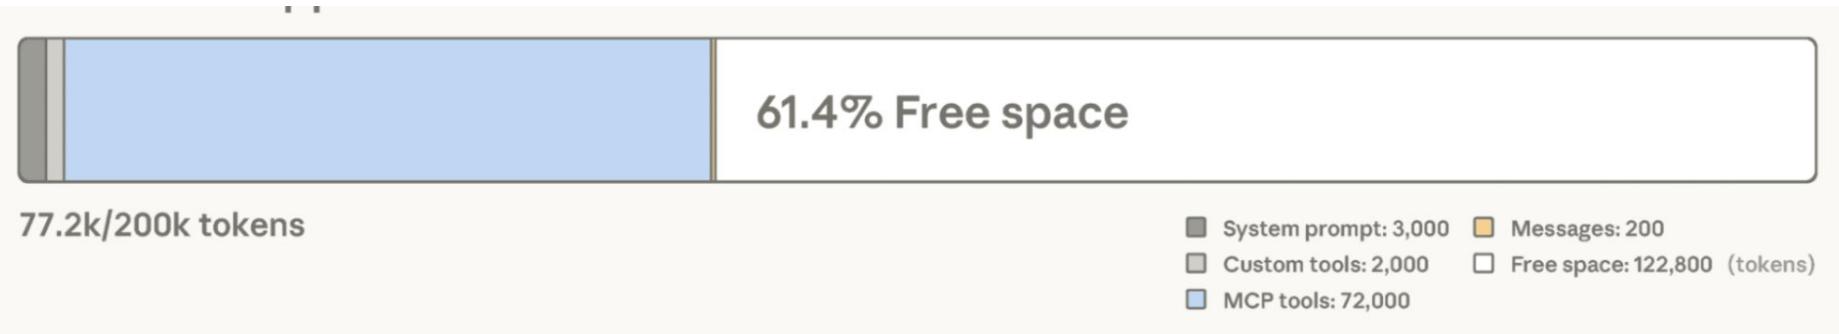
# **Hardware**

# Quelques rappels

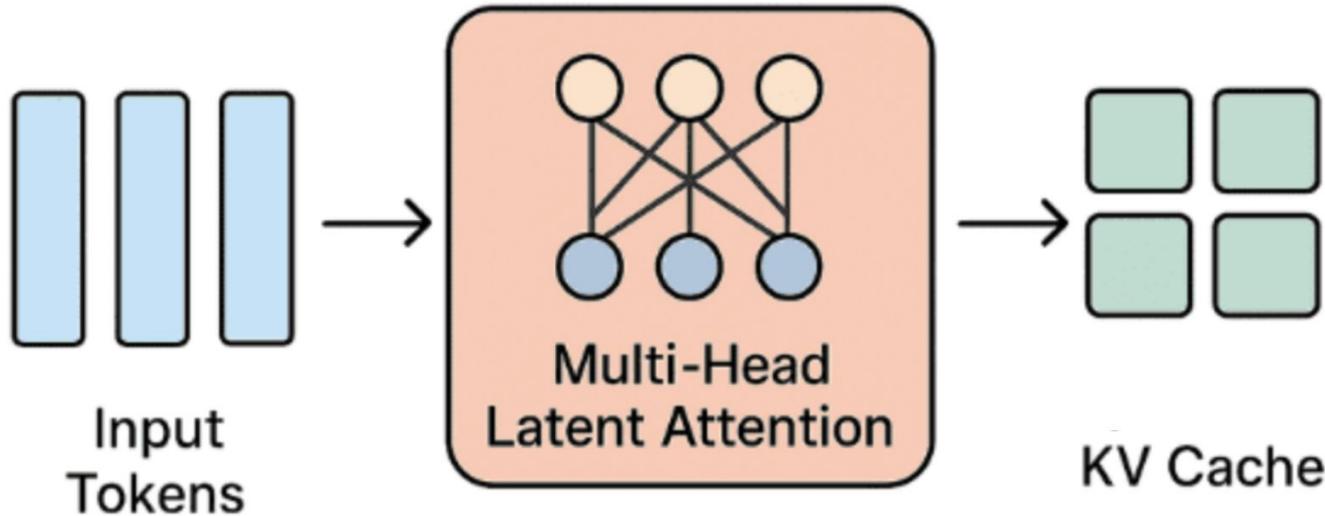
Les agents ont besoin de contexts énormes

MCP GitHub: 25k tokens

Multi-turn conversations



# Quelques rappels



# Quelques rappels

## K/V-cache size:

- Llama-3.1-405B FP32: **1 MB** / tok
- Qwen3-32B FP32: **0.32 MB** / tok
- Qwen3-32B **int8**: **0.08 MB** / tok
- Qwen3-**0.6**B int8: **0.028 MB** / tok

# Quelques rappels

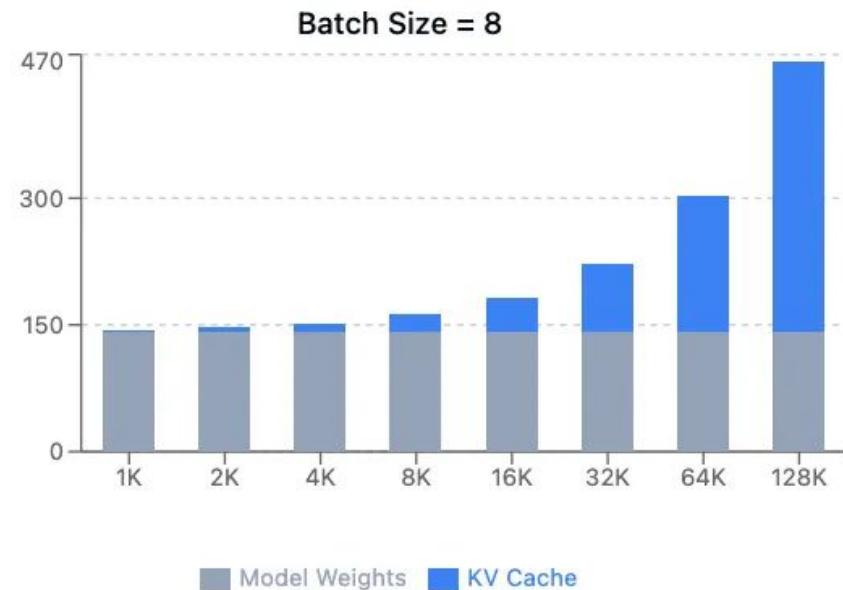
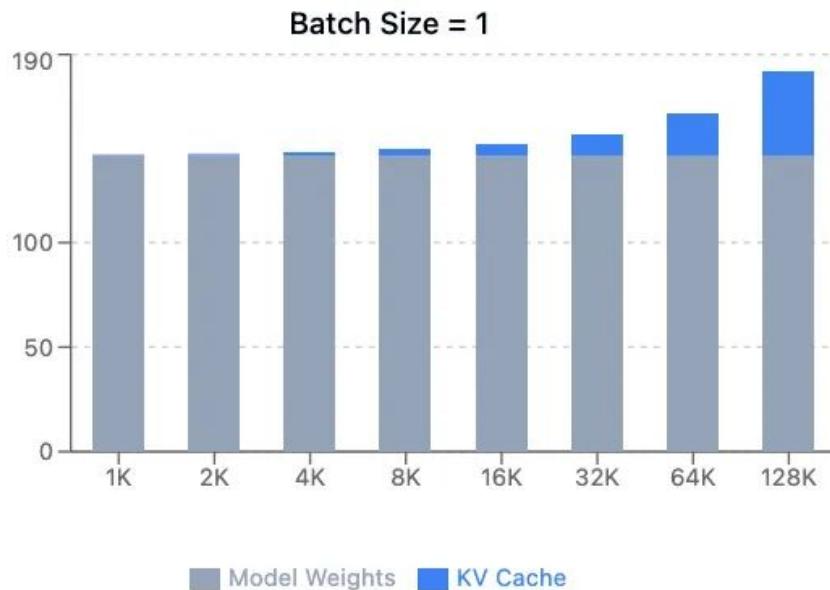
**100 kB per K/V cache  
x 10k tokens (Claude System Prompt)**

**= 1 GB**

# Quelques rappels

## LLaMA 3.3 70B: KV Cache vs Model Weights (142GB)

Memory Composition at Different Sequence Lengths



# Quelques rappels

## Prefix-caching

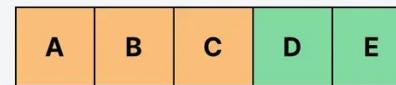
### Request #1

A chat between a curious user and an artificial intelligence assistant. The assistant gives helpful, detailed, and polite answers to the user's questions.  
User: Hello!

### Request #2

A chat between a curious user and an artificial intelligence assistant. The assistant gives helpful, detailed, and polite answers to the user's questions.  
User: How are you?

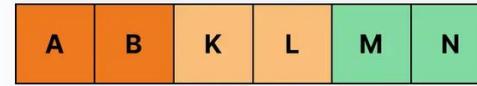
### Request #1



### Request #2

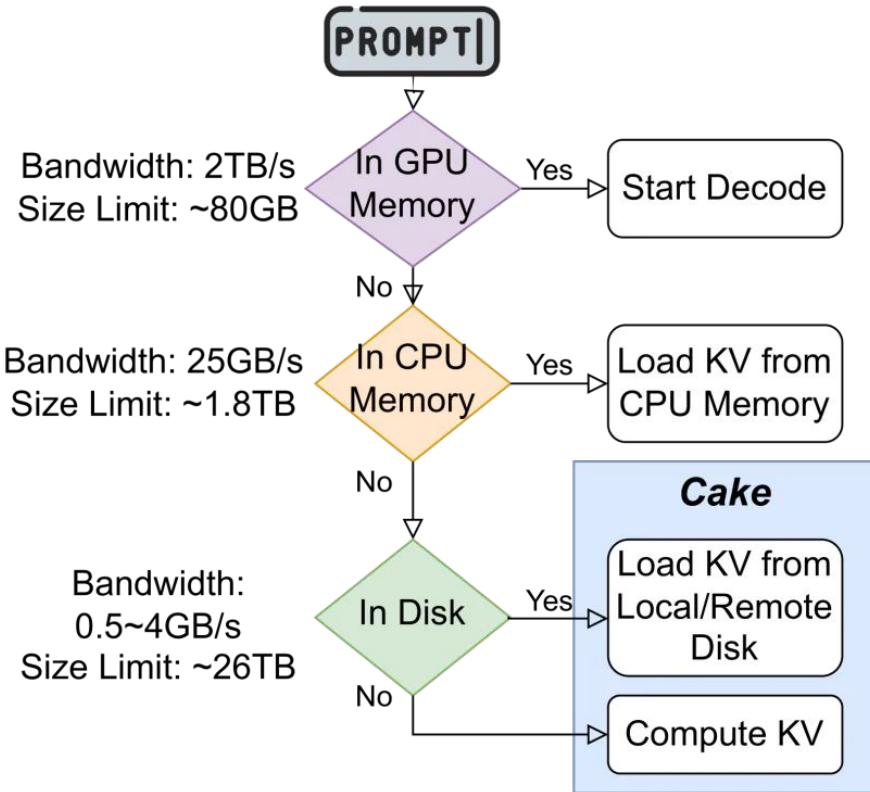


### Request #3



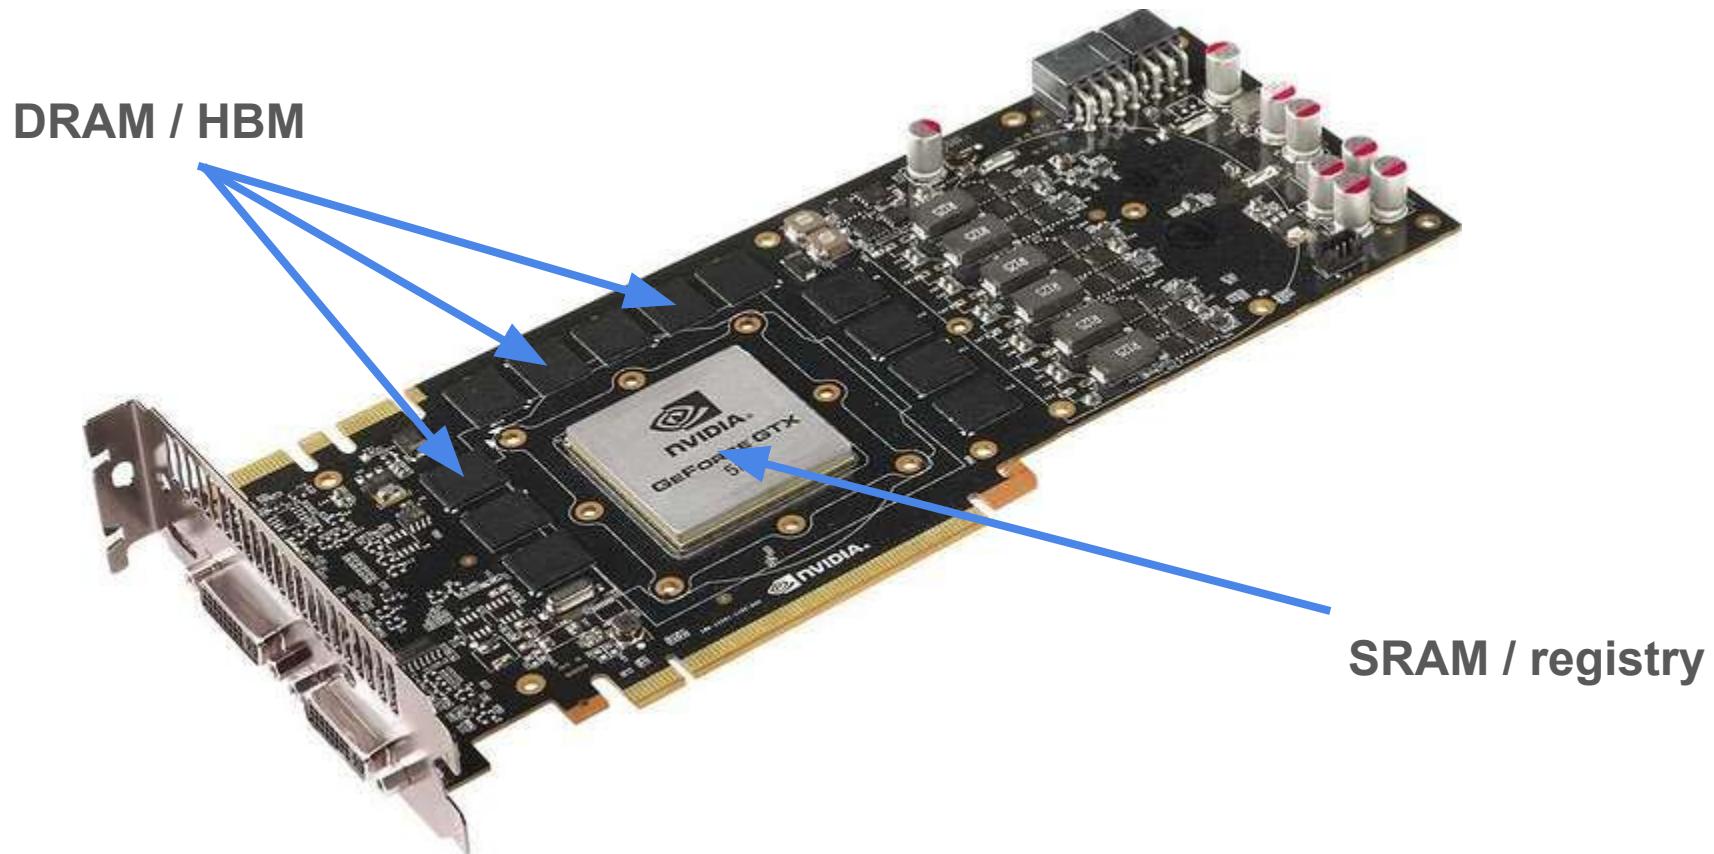
■ Prefill ■ Prefill (Cached) ■ Decode

# Quelques rappels



**Persistent K/V cache & Offloading**  
**(SSD is slower than GPU memory)**

# Quelques rappels



# Comparaisons

	<b>SRAM</b>	<b>DRAM / HBM</b>	<b>Bandwidth DRAM</b>
<b>H100</b>	50 MB	80 GB	3.35 TB/sec
<b>TPU v5</b>	?	95 GB	2.7 TB/sec

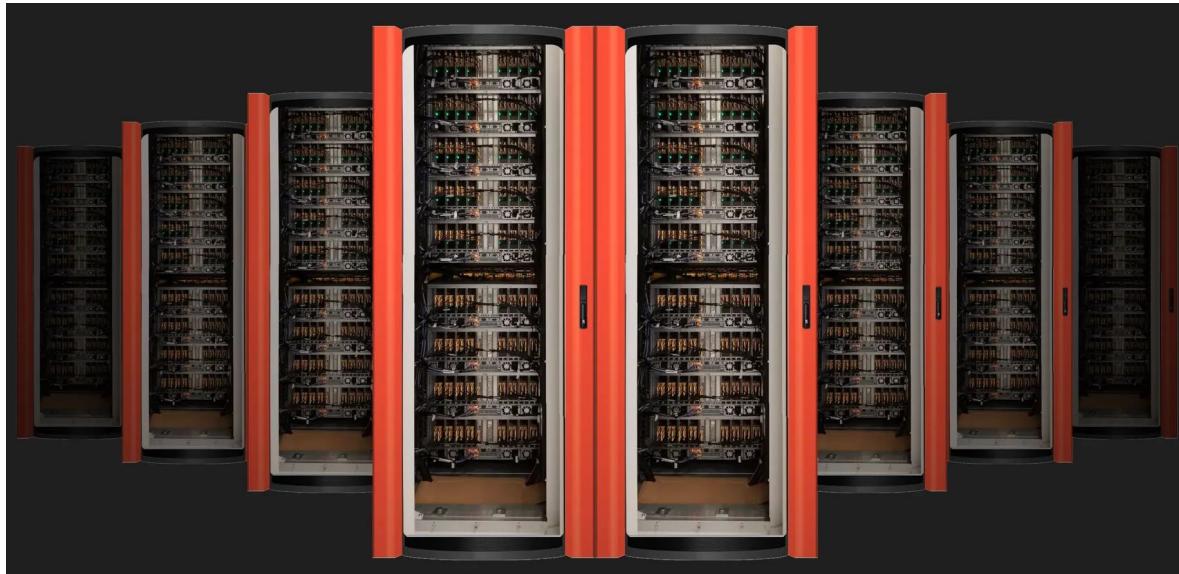
# Comparaisons

	<b>SRAM</b>	<b>DRAM / HBM</b>	<b>Bandwidth DRAM</b>
<b>H100</b>	50 MB	80 GB	3.35 TB/sec
<b>TPU v5</b>	?	95 GB	2.7 TB/sec
<b>B200</b>	200 MB	192 GB	8 TB/sec
<b>TPU v7</b>	?	192 GB	7.37 TB/sec

# Comparaisons

	<b>SRAM</b>	<b>DRAM / HBM</b>	<b>Bandwidth DRAM</b>
<b>H100</b>	50 MB	80 GB	3.35 TB/sec
<b>TPU v5</b>	?	95 GB	2.7 TB/sec
<b>B200</b>	200 MB	192 GB	8 TB/sec
<b>TPU v7</b>	?	192 GB	7.37 TB/sec
<b>Groq LPU</b>	230 MB	0 GB	/

# Groq



Models 70B:

- 1000 GPU
- 10 racks
- 500+ tok/sec

# Groq

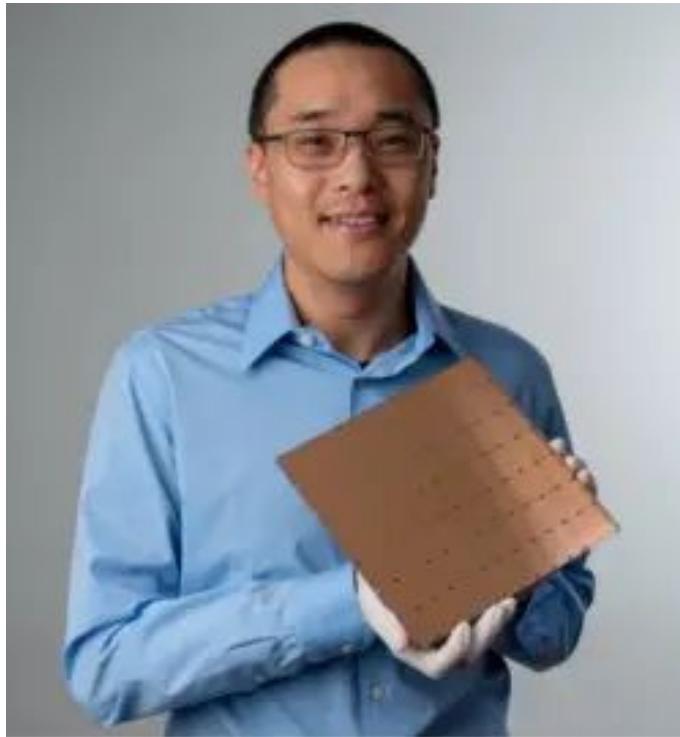


Nvidia acquisition  
20 G\$

# Comparaisons

	<b>SRAM</b>	<b>DRAM / HBM</b>	<b>Bandwidth DRAM</b>
<b>H100</b>	50 MB	80 GB	3.35 TB/sec
<b>TPU v5</b>	?	95 GB	2.7 TB/sec
<b>B200</b>	200 MB	192 GB	8 TB/sec
<b>TPU v7</b>	?	192 GB	7.37 TB/sec
<b>Groq LPU</b>	230 MB	0 GB	/
<b>Cerebras</b>	44 GB	0 GB	/

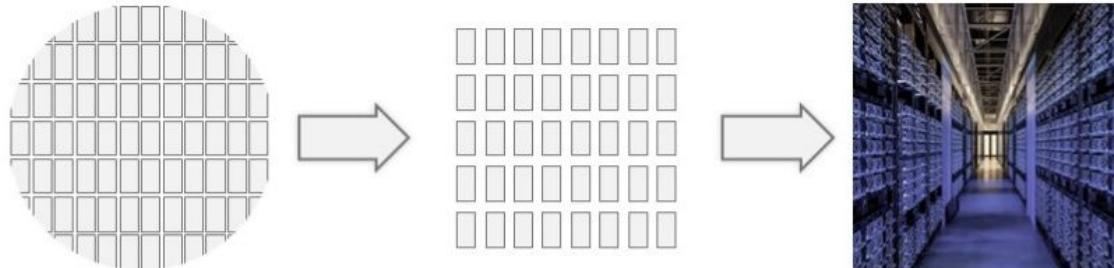
# Cerebras



Models 70B:

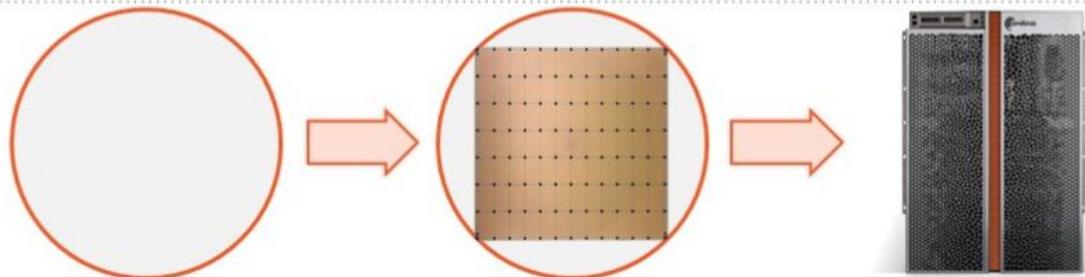
- 2500+ tok/sec

**Traditional GPUs require a complex network of hundreds to thousands of small chips for GenAI**



Traditionally, silicon wafers are cut and packaged into individual chips and servers, then reassembled again in the data center through complex and slow networking. Hundreds to thousands of small chips need to be tied together and programmed to work together to provide enough compute and memory to run modern GenAI workloads.

**Cerebras delivers the massive compute and memory needed for GenAI on a single chip**



The Cerebras WSE-3 is the largest square that can be cut out of a silicon wafer. The massive compute and memory resources required for GenAI are all kept on a single piece of silicon, leading to 52x more cores, 880x more on-chip memory, 7,000x more memory bandwidth, and 3,715x more interconnect bandwidth than the NVIDIA H100. Packaged in the CS-3 system, it fits into a standard data center rack.

# OpenAI signs \$10 billion computing deal with Nvidia challenger Cerebras

By Zaheer Kachwala and Krystal Hu

January 15, 2026 3:34 AM GMT+1 · Updated 10 hours ago



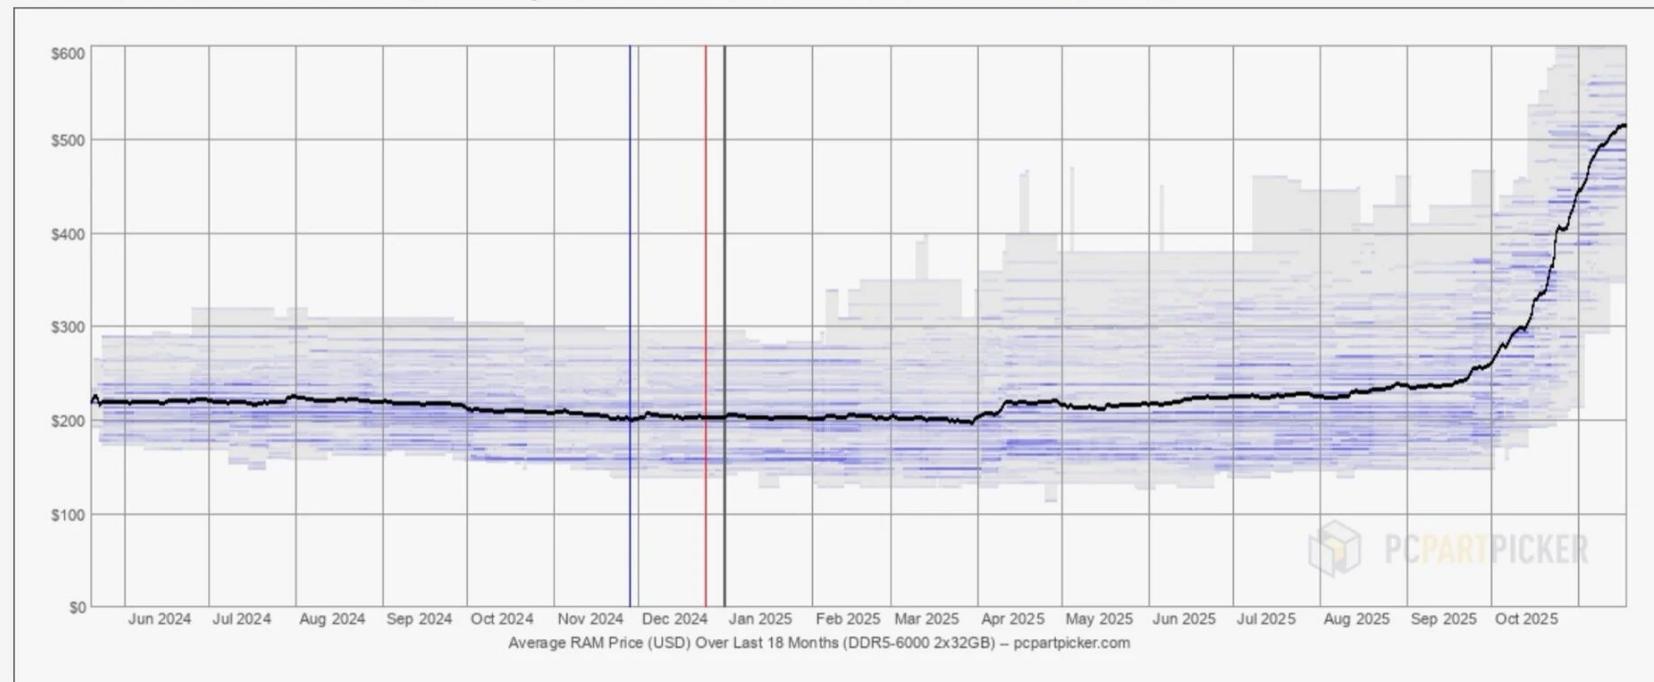
OpenAI logo is seen in this illustration taken May 20, 2024. REUTERS/Dado Ruvic/Illustration/File Photo [Purchase Licensing Rights](#)

**Bottleneck of LLMs:**

**Memory capacity and  
bandwidth**

# Explosion du prix de la DRAM

DDR5-6000 2x32GB (Average price in USD over last 18 months)



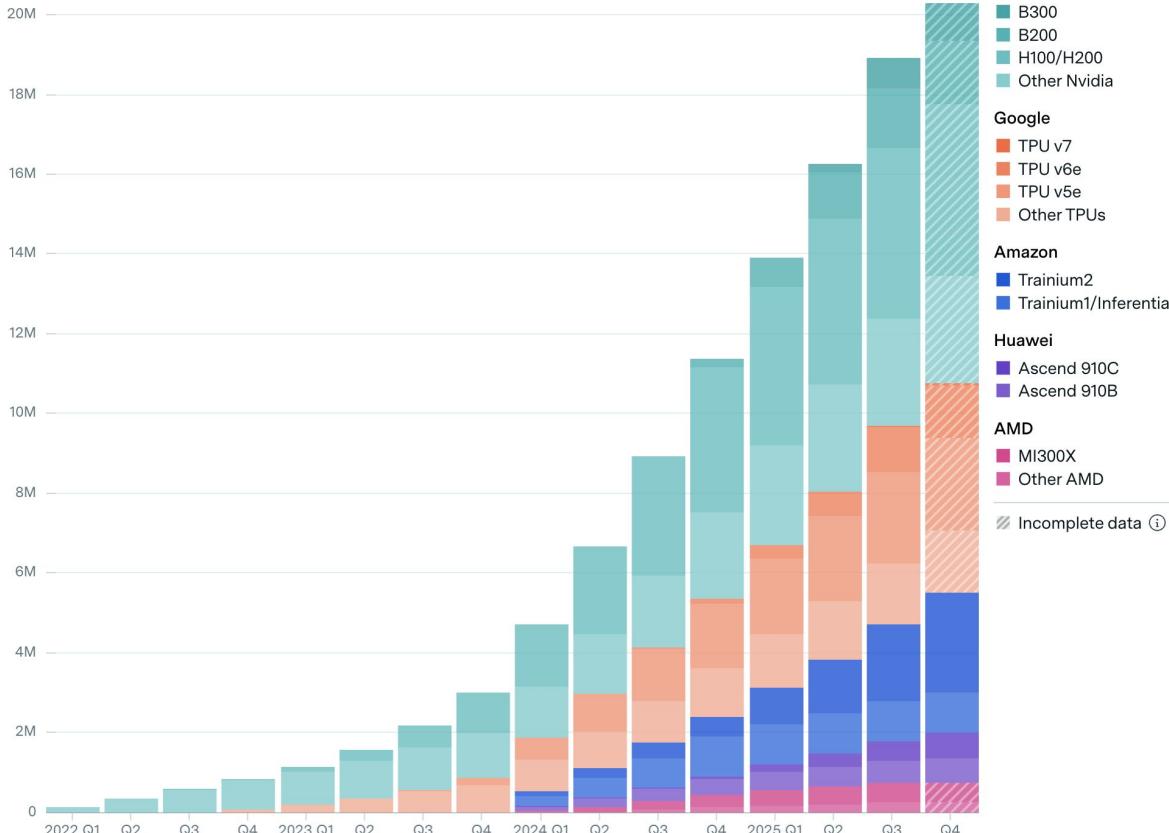
# Explosion du prix de la DRAM



# AI Chip Sales

Graph Settings

Cumulative number of units ⓘ



Global AI computing capacity is doubling every 7 months

(cumulative)

<https://epoch.ai/data/ai-chip-sales>

# AI Chip Sales

Graph Settings

Cumulative number of units ⓘ

20M

18M

16M

14M

12M

10M

8M

6M

4M

2M

0

2022 Q1 Q2 Q3 Q4 2023 Q1 Q2 Q3 Q4 2024 Q1 Q2 Q3 Q4 2025 Q1 Q2 Q3 Q4

## Nvidia

- B300
- B200
- H100/H200
- Other Nvidia

## Google

- TPU v7
- TPU v6e
- TPU v5e
- Other TPUs

## Amazon

- Trainium2
- Trainium1/Inferentia

## Huawei

- Ascend 910C
- Ascend 910B

## AMD

- MI300X
- Other AMD

Incomplete data ⓘ

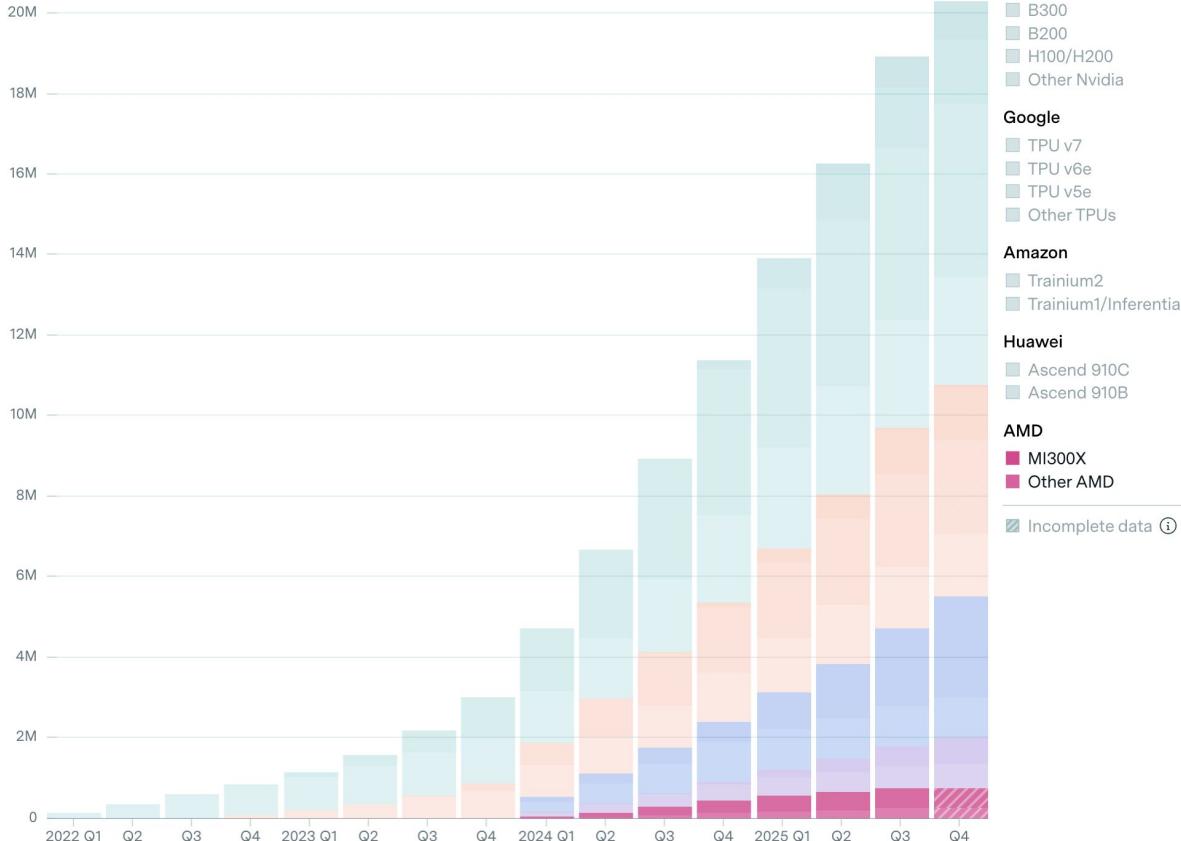
# China

<https://epoch.ai/data/ai-chip-sales>

# AI Chip Sales

Graph Settings

Cumulative number of units ⓘ



AMD

# AI Chip Sales

Graph Settings

Cumulative number of units ⓘ

20M

18M

16M

14M

12M

10M

8M

6M

4M

2M

0

2022 Q1 Q2 Q3 Q4 2023 Q1 Q2 Q3 Q4 2024 Q1 Q2 Q3 Q4 2025 Q1 Q2 Q3 Q4

## Nvidia

- B300
- B200
- H100/H200
- Other Nvidia

## Google

- TPU v7
- TPU v6e
- TPU v5e
- Other TPUs

## Amazon

- Trainium2
- Trainium1/Inferentia

## Huawei

- Ascend 910C
- Ascend 910B

## AMD

- MI300X
- Other AMD

Incomplete data ⓘ

# Google TPU v7

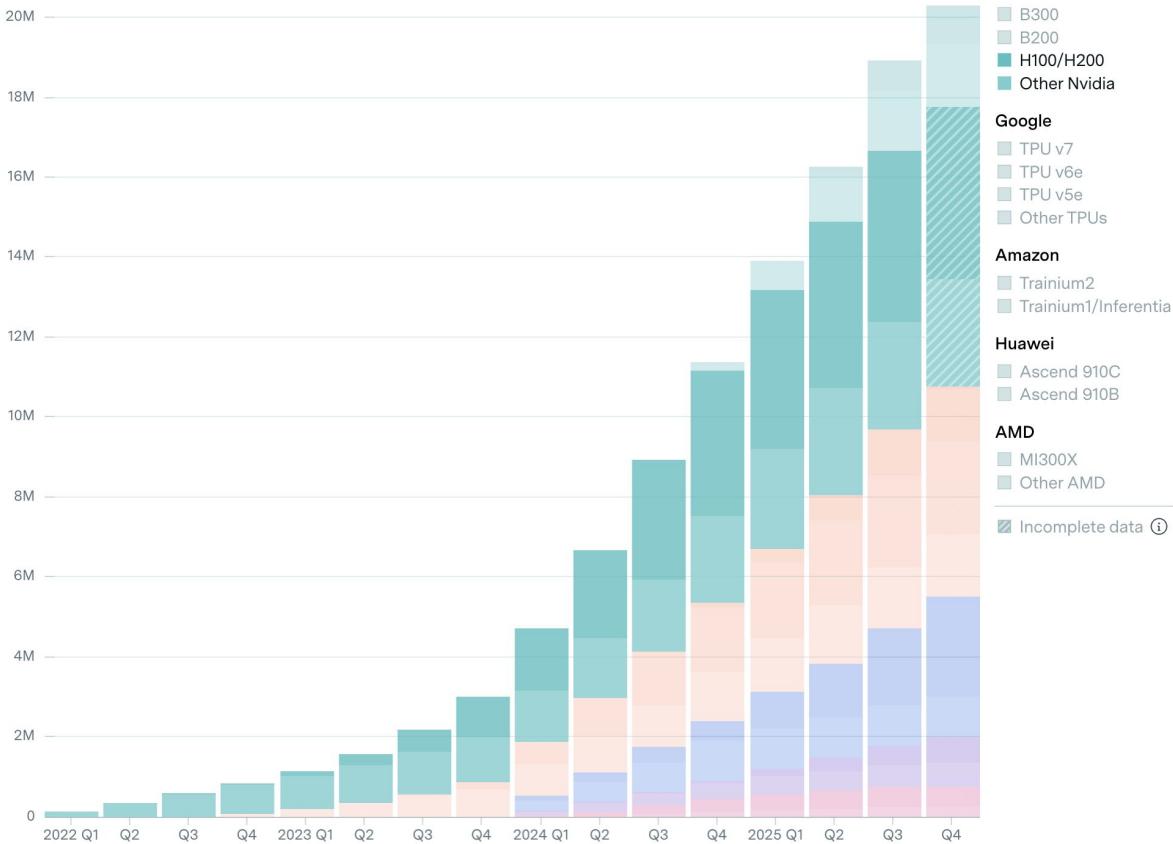
Révolutionnaire, mais marginal

<https://epoch.ai/data/ai-chip-sales>



# AI Chip Sales

Cumulative number of units ⓘ



Graph Settings

**Nvidia GPU pre-ChatGPT**

“Volta” generation  
-> V100 (32GB HBM)

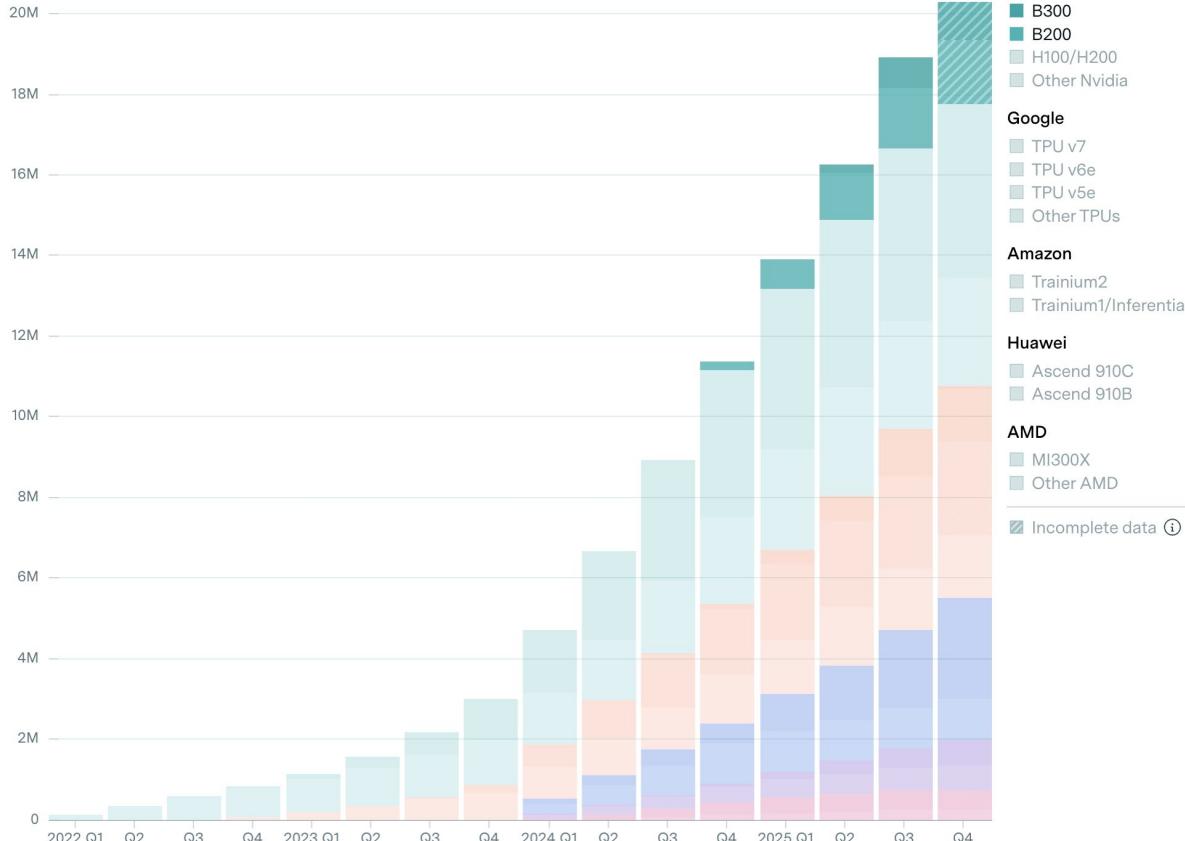
“Ampere” generation  
-> A100 (80GB HBM)

“Hopper” generation  
-> H100 (80GB HBM)  
-> H200 (141GB HBM)

# AI Chip Sales

Graph Settings

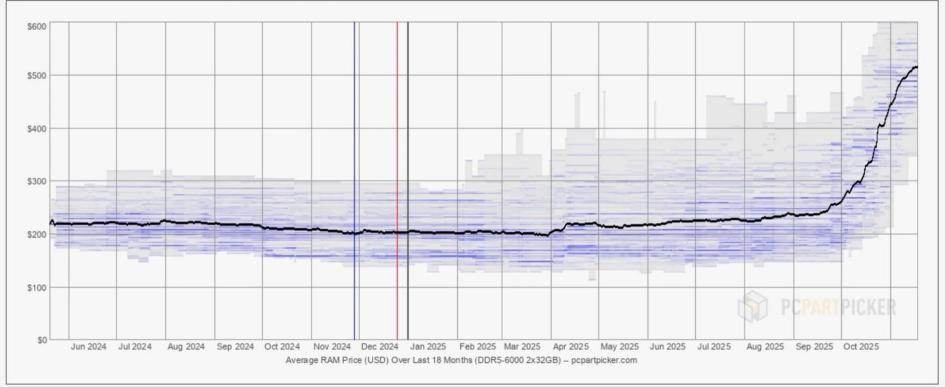
Cumulative number of units ⓘ



## Nvidia GPU post-ChatGPT

“Blackwell” generation  
-> B200 (192GB HBM)  
-> B300 (270GB HBM)

## DDR5-6000 2x32GB (Average price in USD over last 18 months)



↗  
Collapse



NASDAQ: MU  
Micron Technology



# Explosion du prix de la DRAM

N TOOLS



Download PDF

---

December 3, 2025 at 11:40 AM EST

## Micron Announces Exit from Crucial Consumer Business

BOISE, Idaho, Dec. 03, 2025 (GLOBE NEWSWIRE) -- Micron Technology, Inc. (Nasdaq: MU), a leader in innovative memory and storage solutions, today announced its decision to exit the Crucial consumer business, including the sale of Crucial consumer-branded products at key retailers, e-tailers and distributors worldwide.

# Explosion du prix de la DRAM

## DRAM crisis may last until 2028 – SK Hynix sees no quick recovery

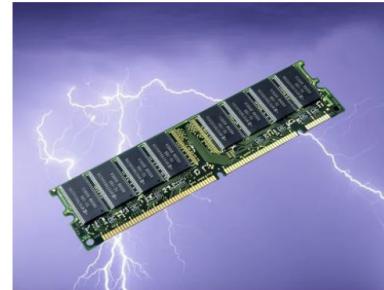
Internal analyses by SK Hynix suggest that the current DRAM shortage could persist until at least 2028 – although some forecasts paint a far more optimistic picture.

Marius Müller ([translated by](#) Marius Müller), Published 12/16/2025



[Business](#) [Leaks / Rumors](#) [Desktop](#) [Gaming](#) [Social Media](#)

The memory crisis, largely driven by surging demand from AI data centers, is reportedly intensifying – [Samsung has allegedly doubled contract prices](#) for DDR5 RAM, while manufacturers such as Dell and Lenovo are said to be [limiting mid-range laptops](#) to just 8 GB of memory. The long-anticipated Steam Machine could face [delays](#) or launch only as a [barebone system](#) due to the shortage. While many are hoping for a swift resolution, SK Hynix paints a far more pessimistic picture of what lies ahead.



*The South Korean chipmaker SK Hynix expects the memory crisis to persist until 2028. (Image source: Pixabay Oimheidi/ cliffsmith23)*

# Nvidia “Storage Next” project



## SK hynix Targets 100M IOPS NAND by 2027 in Joint Effort with NVIDIA

by Nomad76 | Dec 11th, 2025 15:35 | [Discuss \(3 Comments\)](#)

SK hynix is pushing hard on next-gen AI NAND, and 2027 seems to be its next major milestone. According to a ZDNet report, the company is working with NVIDIA to build ultra-fast AI-focused NAND chips that could hit up to 30x the performance of today's enterprise SSDs. Early samples are planned for late 2026, with second-gen entering mass production by the end of 2027. A big part of that effort is SK hynix AI-N P high-performance SSD architecture aimed at removing I/O bottlenecks in large AI inference workloads. ZDNet says the redesigned NAND and controller are already in proof-of-concept (PoC) testing with NVIDIA. SK hynix is targeting 25 million IOPS on PCIe Gen 6 for first samples next year, and 100 million IOPS for the production version in 2027. For comparison, current enterprise SSDs manage roughly 2-3 million IOPS (high-end models).

# Nvidia “Storage Next” project

~~Von Neumann architecture~~  
~~(CPU<->Memory<->Storage)~~



GPU<->SSD

Expand DRAM with SSD

10m-100m context window in 2027 ?

# DPU BlueField 4

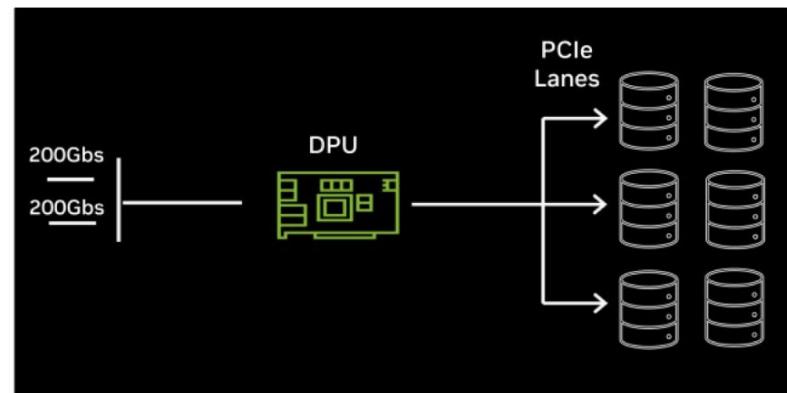
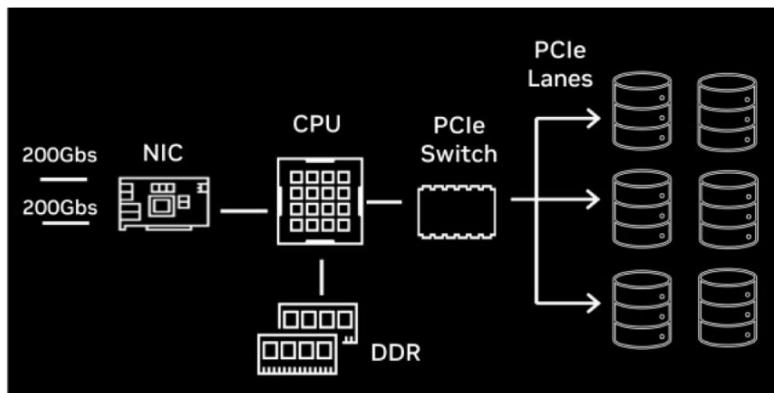


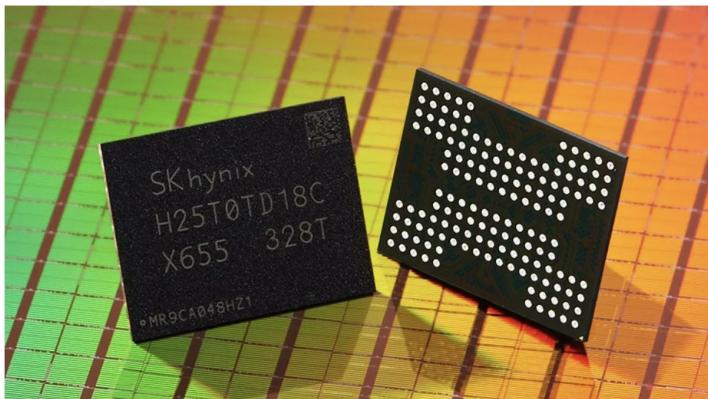
Figure 2 - x86 CPU base JBOF (left) and DPU based JBOF (right)

# US foundries

**SK hynix to build first U.S. packaging plant for HBM — plugs critical hole in U.S. supply chain, \$3.9B investment challenges TSMC and reshapes AI supply chains**

News Analysis By Luke James published December 30, 2025

Memory giant eyes turnkey HBM strategy with Indiana packaging line as U.S. reshoring effort accelerates.



(Image credit: SK hynix)

≡  
MENU

micron



June 12, 2025 at 8:30 AM EDT

## Micron and Trump Administration Announce Expanded U.S. Investments in Leading-Edge DRAM Manufacturing and R&D

*Micron Plans to Invest Approximately \$200 Billion in Semiconductor Manufacturing and R&D in Idaho, New York and Virginia, Enhancing Domestic Memory Supply and Technology Leadership*

*Micron to Build Second Leading-Edge Memory Manufacturing Fab in Idaho, Modernize and Expand Virginia Fab and Bring End-to-End High Bandwidth Memory (HBM) Manufacturing Capabilities to U.S. to Meet Anticipated AI-Driven Demand*

# Future GPU market ?

Cerebras/Groq for ultra-low latency usages (real-time STT, robotics...)

Current GPU architecture for mid to large models

GPUs backed by SSD for extra long CoT/reasoning context, video generation and agentic AI

# US foundries

Taiwan considers TSMC export ban that would prevent manufacturing its newest chip nodes in U.S. — limit exports to two generations behind leading-edge nodes, could slow down U.S. expansion

News By Anton Shilov published December 19, 2025

Largely maintained for now, it is set to go in the future.

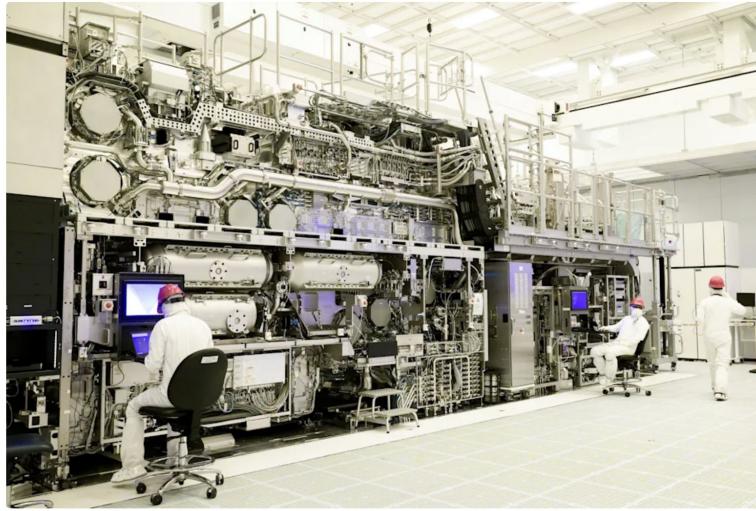


(Image credit: TSMC)

# China vs Taiwan

Bloomberg

**ASML and TSMC can disable chip machines if China invades Taiwan**



ASML and TSMC can disable chip machines if China invades Taiwan · Bloomberg

Conversation:

[https://www.reddit.com/r/taiwan/comments/1cxt96b/asml\\_and\\_tsmc\\_can\\_disable\\_chip\\_machines\\_if\\_china/](https://www.reddit.com/r/taiwan/comments/1cxt96b/asml_and_tsmc_can_disable_chip_machines_if_china/)

# China vs Taiwan

**Exclusive: How China built its  
'Manhattan Project' to rival the West in  
AI chips**

By Fanny Potkin

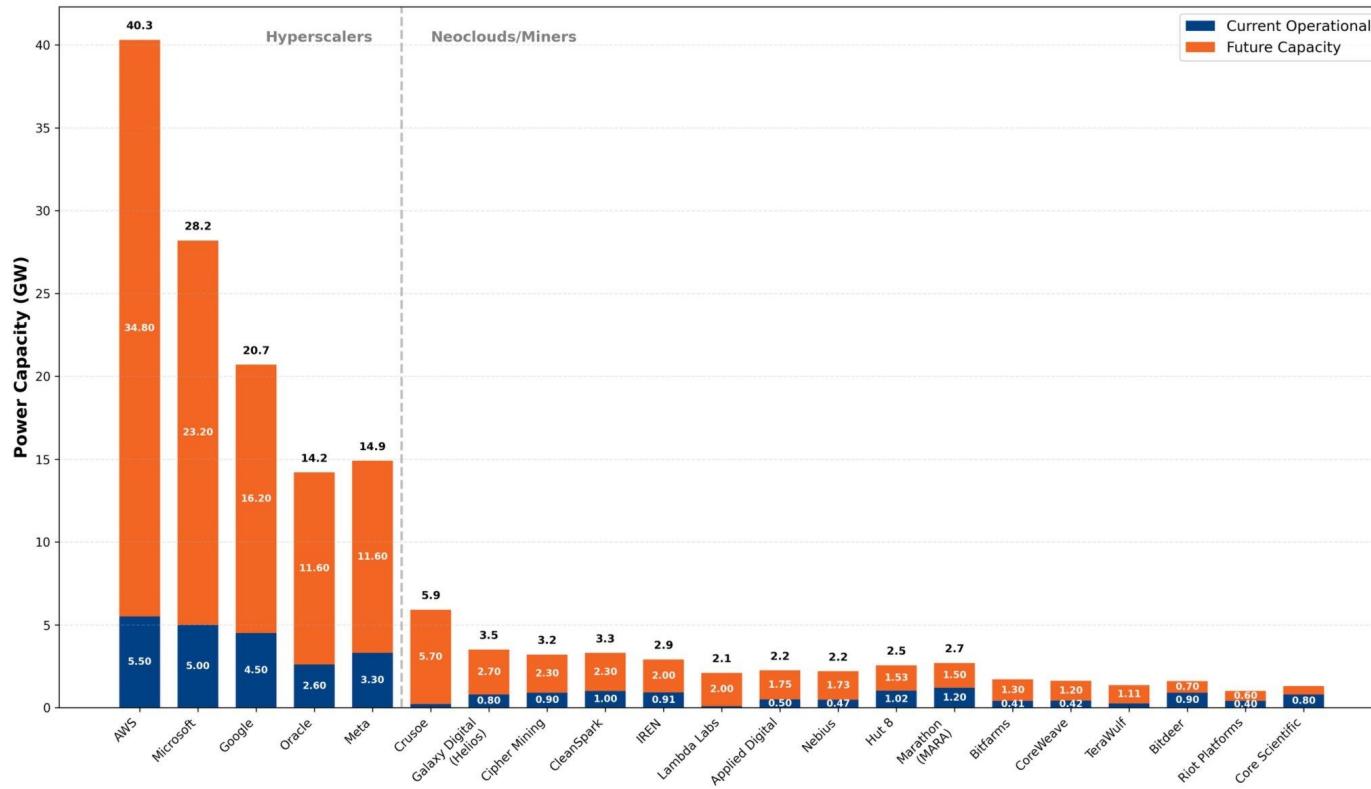
December 18, 2025 5:11 AM GMT+1 · Updated December 18, 2025



A Chinese flag is displayed next to a "Made in China" sign seen on a printed circuit board with semiconductor chips, in this illustration picture taken February 17, 2023. REUTERS/Florence Lo/Illustration/File Photo [Purchase Licensing Rights](#)

<https://www.reuters.com/world/china/how-china-built-its-manhattan-project-rival-west-ai-chips-2025-12-17/>

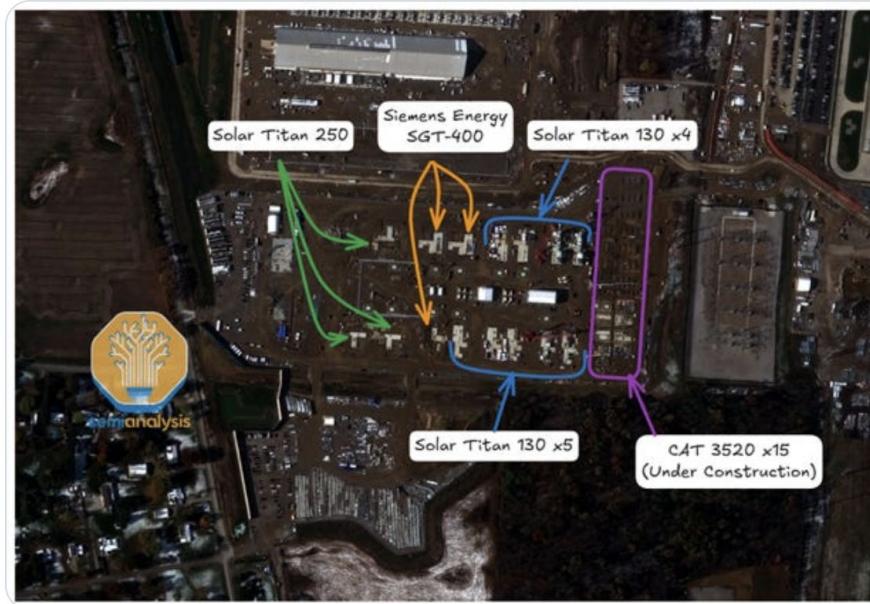
**North America Datacenter Power Capacity: Hyperscalers vs Neoclouds/Bitcoin Miners**  
(Updated November 25, 2025)



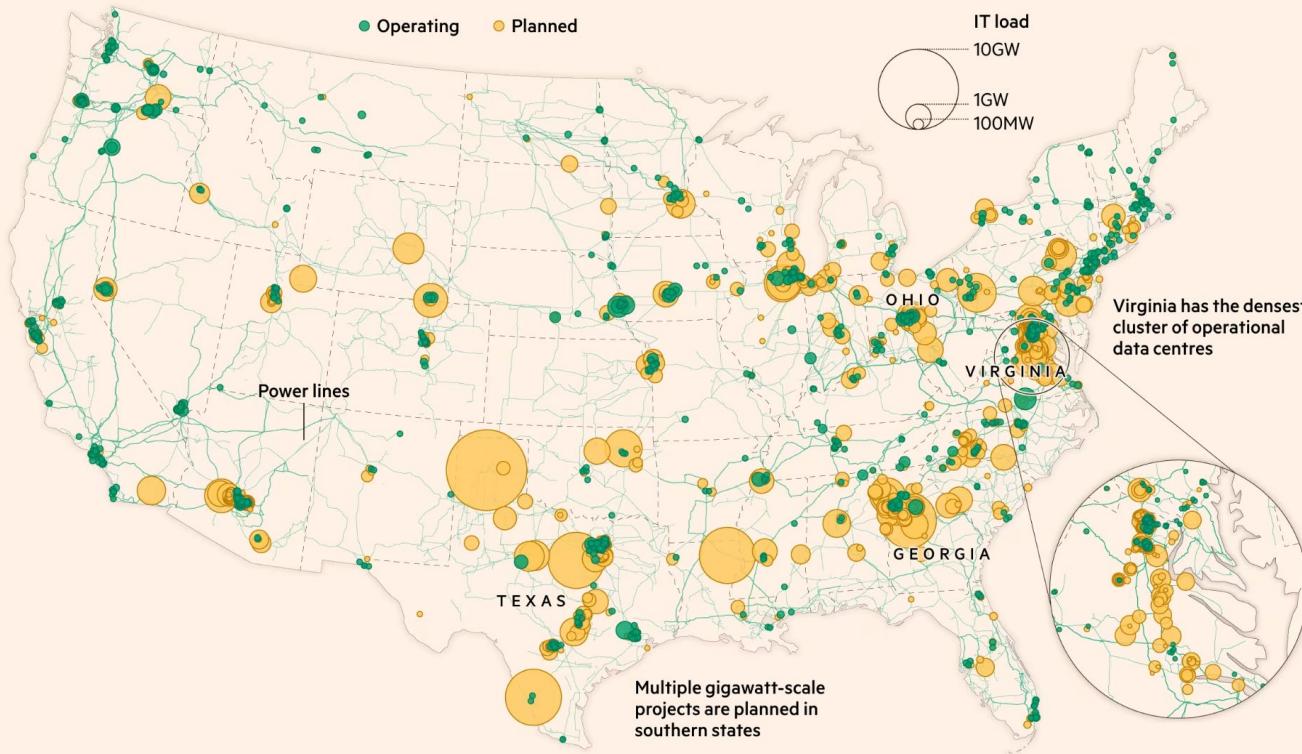


BREAKING: Meta is building 200MW+ of on site gas datacenter power generation using 4 different types of turbines & engines. Meta's Socrates South project is a hybrid fleet with N+1+1 design that considers of

- 3 × Solar Titan 250 IGTs (23 MW)
- 9 × Solar Titan 130 IGTs (16.5 MW)
- 3 × Siemens SGT-400 IGTs (14.3 MW)
- 15 × Caterpillar 3520 fast-start engines (3.1 MW)



New data centres are expected to draw significantly more power



## Datacenter projects

# NATURAL GAS PRODUCTION

by state and province

billion cubic feet of natural gas per day

**Western Canada**  
recently began  
exporting liquefied  
natural gas from the  
West Coast

**Texas** accounted for nearly  
30% of US natural gas production  
in the first half of 2025

The **Marcellus** and **Utica** shale  
plays in Appalachia contain nearly  
half of US natural gas reserves

## New datacenter location

## Vs

## Natural gas production

Source: US Energy Information Administration, Natural Resources Canada

 ORENNIA

# AI data centers may soon be powered by retired Navy nuclear reactors from aircraft carriers and submarines — firm asks U.S. DOE for a loan guarantee to start the project

News By Jowi Morales published December 25, 2025

This is a much cheaper and faster way to get nuclear power.



(Image credit: Pixabay)

# Nvidia @ CES

To be continued...

# Claude Code : L'IA directement dans votre Terminal

Le meilleur outil actuel pour accélérer le développement au sein de  
votre CLI



# L'Impact des Skills : Avant vs Après

## Sans Skill ("AI Slop")

- Police Inter basique
- Fond blanc
- Dégradé violet cliché
- Mise en page prévisible

## Avec Frontend Skill

- Typographie unique (Bricolage Grotesque/JetBrains Mono)
- Mode sombre avec accents Cyan
- Animations d'apparition étagées
- Textures en fond

Les Skills transforment une réponse générique en un résultat expert et adapté au contexte.

# Les Skills : Expertise Pointue et Spécialisée



## Définition du Skill

Un "Skill" est un fichier de configuration qui donne à Claude une expertise pointue sur un domaine précis.



## Le Problème Résolu : "AI Slop"

Le "AI Slop" est évité (design générique, gradients violets, polices comme Inter/Arial).

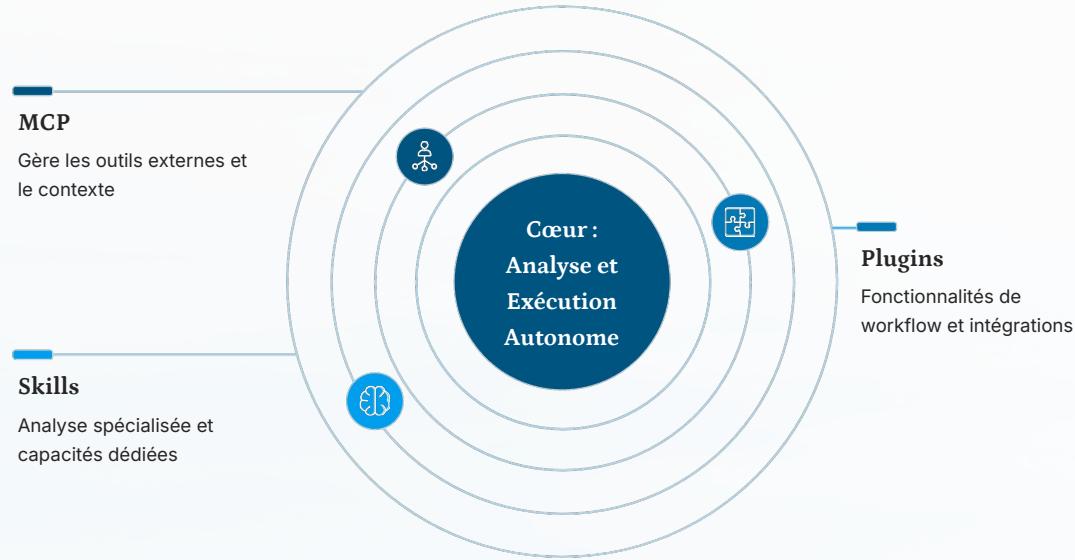


## La Solution du Skill

Force l'IA à utiliser des typographies audacieuses, des palettes de couleurs cohérentes et des animations CSS soignées.

# L'Agent Augmenté : Architecture Complète

Découvrez l'architecture technique qui propulse l'Agent Augmenté, articulant ses composants clés pour une performance optimale et une adaptabilité maximale.



## Le Cœur

Analyse et exécution autonome sur la base de code, formant le moteur principal de l'agent.

## MCP (Model Context Protocol)

Assure la connexion et l'intégration avec des outils externes, tels que Context Seven pour la documentation ou Playwright pour les tests E2E.

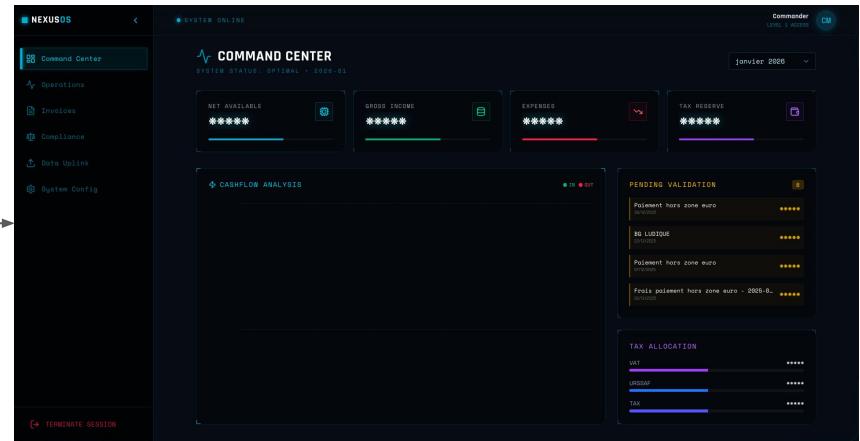
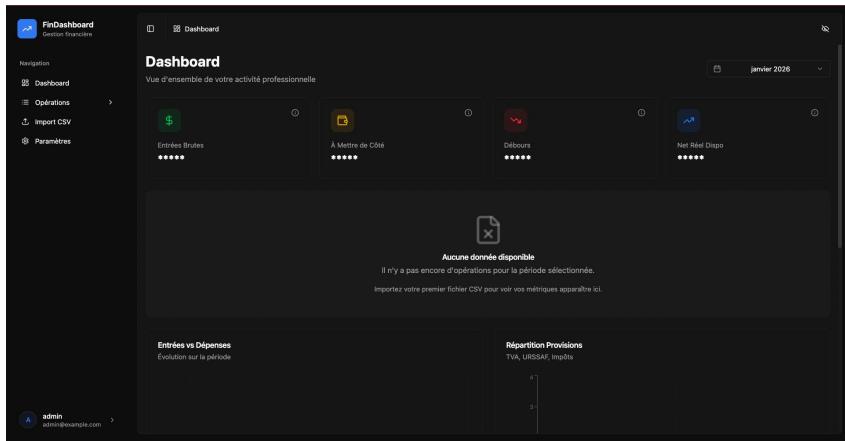
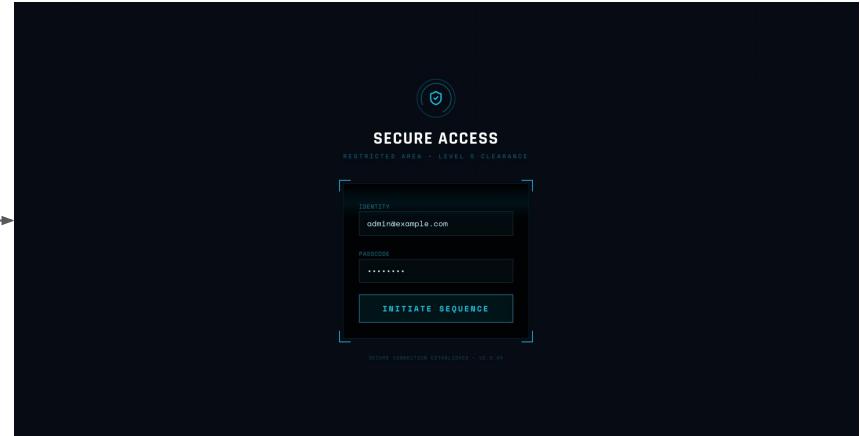
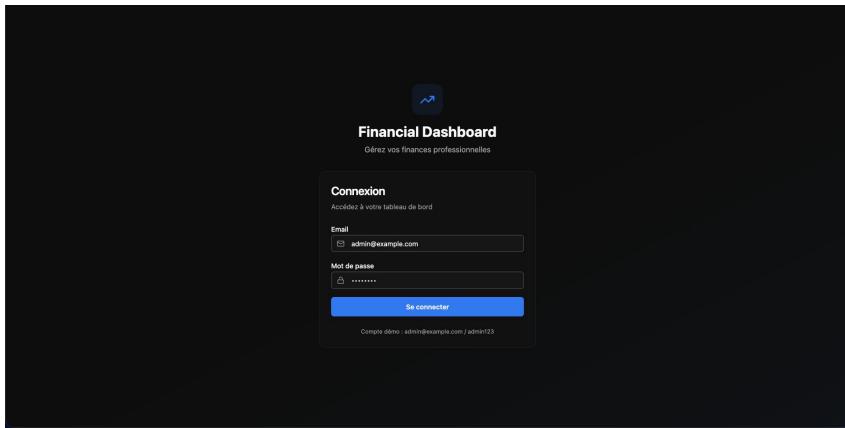
## Plugins

Permettent l'ajout de fonctionnalités de workflow spécifiques, comme la création automatique de tickets Jira ou l'affichage d'éléments graphiques en CLI.

## Skills

Fichiers de compétences personnalisés, essentiels pour spécialiser et affiner l'analyse de l'agent selon les besoins.

# Le résultat



# Le résultat

**FinDashboard** Gestion financière

Navigation

- Dashboard
- Opérations
- Liste
- Factures
- Obligations

Import CSV

Import CSV

Nouveau fichier

Cliquez pour sélectionner un fichier CSV  
Format accepté: .csv

Banque

Sélectionner la banque

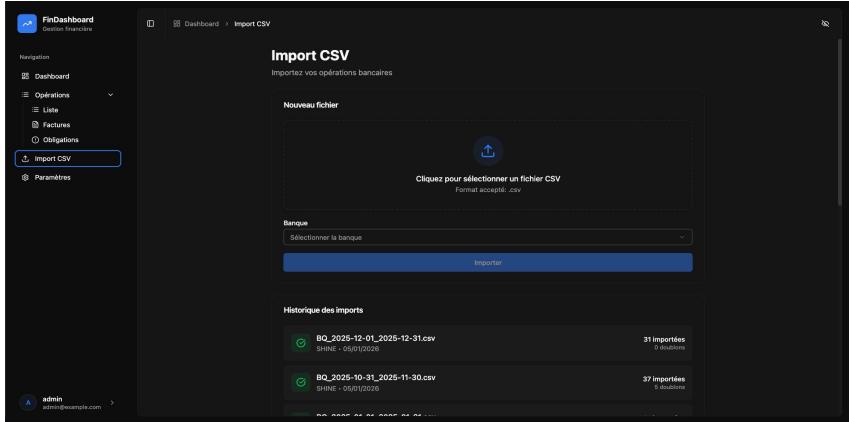
Importer

Historique des imports

BO\_2025-12-01\_2025-12-31.csv SHINE - 05/01/2026 31 importées 0 échouées

BO\_2025-10-31\_2025-11-30.csv SHINE - 05/01/2026 37 importées 0 échouées

admin admin@example.com



**NEXUS OS** SYSTEM ONLINE

Command Center

Operations

Invoices

Compliance

Data Uplink

System Config

DATA UPLOAD TERMINAL

INITIATE FILE TRANSFER

DROP CSV FILE OR CLICK TO BROWSE

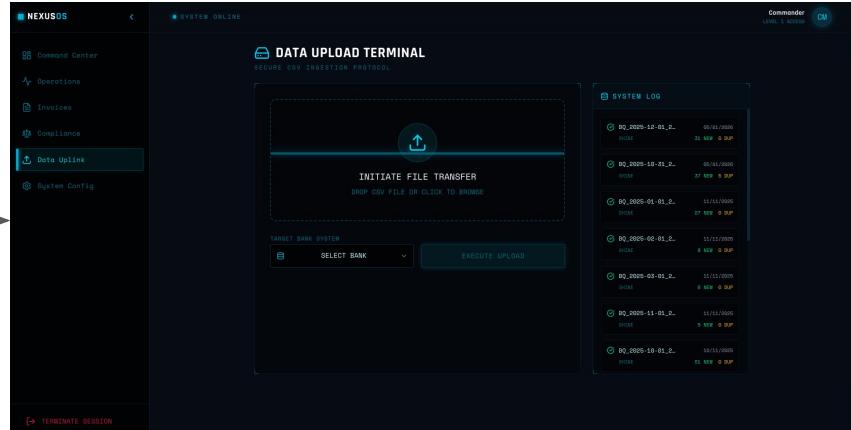
TARGET BANK SYSTEM

SELECT BANK EXECUTE UPLOAD

TERMINATE SESSION

SYSTEM LOG

DATE	LABEL	BANK	AMOUNT	CONTEXT	CATEGORY	ACTIONS
30/12/2025	LE CASTELLI	SHINE	*****	PERSO	achat marchandise	
30/12/2025	Paiement hors zone euro	SHINE	*****	INDETERMINE	achat marchandise	
29/12/2025	RAILWAY	SHINE	*****	PRO	abonnement	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	E-LECLERC MAGASIN	SHINE	*****	PERSO	achat marchandise	



**FinDashboard** Gestion financière

Navigation

- Dashboard
- Opérations
- Liste
- Factures
- Obligations

Import CSV

Opérations

361 opérations

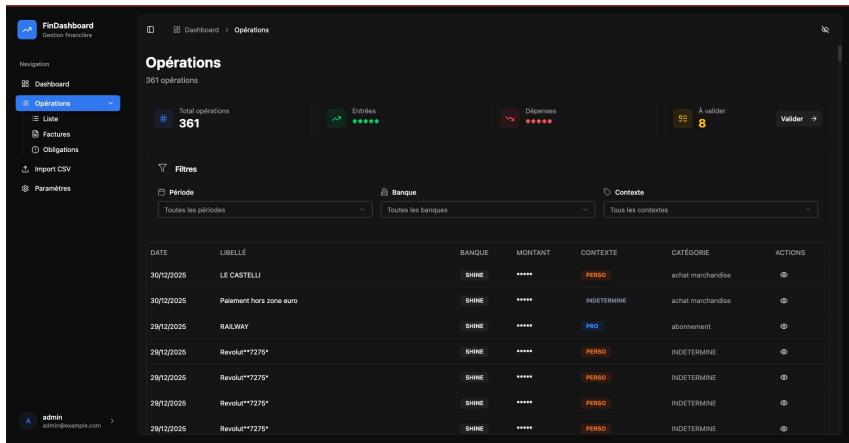
Total opérations # 361 Entrées Dispenses À valider 8 Valider →

Filtres

Période Toutes les périodes Banque Toutes les banques Contexte Tous les contextes

DATE	LIBELLÉ	BANQUE	MONTANT	CONTEXTÉ	CATÉGORIE	ACTIONS
30/12/2025	LE CASTELLI	SHINE	*****	PERSO	achat marchandise	
30/12/2025	Paiement hors zone euro	SHINE	*****	INDETERMINE	achat marchandise	
29/12/2025	RAILWAY	SHINE	*****	PRO	abonnement	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	E-LECLERC MAGASIN	SHINE	*****	PERSO	achat marchandise	

admin admin@example.com



**NEXUS OS** SYSTEM ONLINE

Command Center

Operations

Invoices

Compliance

Data Uplink

System Config

DATA GRID

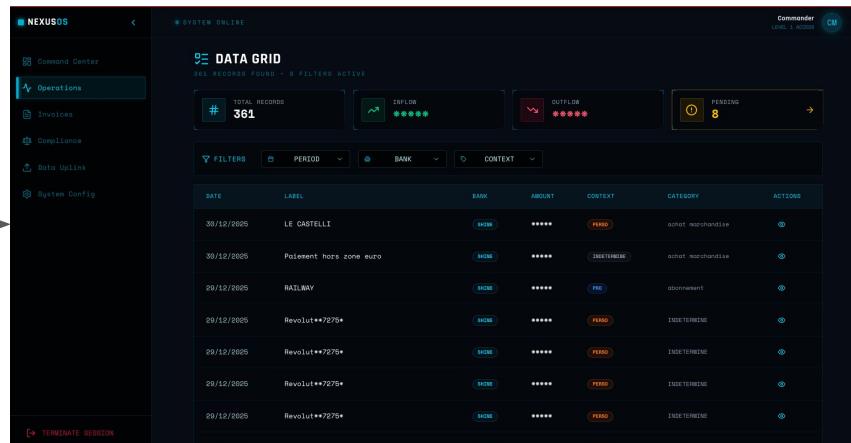
361 TOTAL RECORDS 0 FILTERS ACTIVE

# TOTAL RECORDS 361 INFLOW \*\*\*\* OUTFLOW \*\*\*\* PENDING 8 →

FILTERS PERIOD BANK CONTEXT

DATE	LABEL	BANK	AMOUNT	CONTEXT	CATEGORY	ACTIONS
30/12/2025	LE CASTELLI	SHINE	*****	PERSO	achat marchandise	
30/12/2025	Paiement hors zone euro	SHINE	*****	INDETERMINE	achat marchandise	
29/12/2025	RAILWAY	SHINE	*****	PRO	déboursent	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	Revolut**727*	SHINE	*****	PERSO	INDETERMINE	
29/12/2025	E-LECLERC MAGASIN	SHINE	*****	PERSO	achat marchandise	

TERMINATE SESSION





# “Comment créer une équipe d'agents IA en 2026”



Globetrotters.ai



**Benoit Hediard**

Founder @Globetrotters.ai | Ex-CTO &  
Co-Founder @Agorapulse | Angel Investor



**15 janvier 2026 à 19h00**



4 Rue Voltaire, 44000 Nantes

chez icilundi



techtown

xexternatic

lonestone

RCA

Slides dispo sur:

<https://github.com/genai-nantes-meetup/meetups/>



**Learning expedition le 9-10 février @ “No code & IA Day”**  
**Contacter [lise@lacantine.co](mailto:lise@lacantine.co) pour un coupon**



10 FÉVRIER  
2026

## AI Day 2026

STATION F, PARIS, FRANCE

Le rendez-vous incontournable des 2 000 leaders IA & NoCode qui transforment l'IA en Europe, co-construit avec l'IA-cluster PR[AI]RIE.

Devenir partenaire

Je veux une place

Sous le haut patronage de

Monsieur Emmanuel MACRON

Président de la république



# “Comment Google Antigravity Transforme le vibe-coding”



DeepMind



**Guillaume Vernade**

Developer Advocate @ Google Gemini &  
Deepmind



11 Février 2026 à 19h00



4 Rue Voltaire, 44000 Nantes

chez [icilundi](#)



techtown lonestone RCA  
externatic  zenika [sfɛir]



# Workshop: “L’IA générative appliquée aux RH”



**Simon Timssale**

Consultant et formateur GenAI

**Yannis Sulfourt**

Consultant et formateur GenAI



4 Mars 2026 à 19h00



2 Pl. Louis Daubenton, 44100 Nantes

chez



techtown lonestone RCA  
xternatic zenika [sfΞir]



## “Before Shift”



**18 Mars 2026 à 19h00**



4 Rue Voltaire, 44000 Nantes

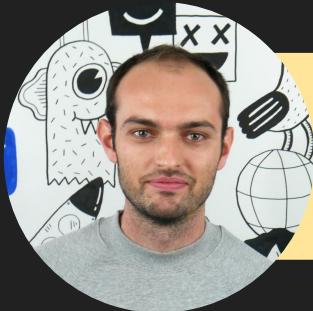
chez  icilundi



techtown lonestone RCA  
xternatic  zenika [sfɛir]



## “Agentic AI for Support”



**Baptiste Jamin**  
CEO @ Crisp



**15 Avril** 2026 à 19h00



4 Rue Voltaire, 44000 Nantes

chez [icilundi](#)



**techtown** **lonestone** **RCA**  
**xexternatic** **zenika** [sfΞir]



# “Réinventer l'architecture cloud pour l'IA en temps réel”



**Sacha Morard**  
CEO @ Edgee  
Ex CTO @ Le Monde



**6 Mai** 2026 à 19h00



4 Rue Voltaire, 44000 Nantes

chez [icilundi](#)



**techtown** **lonestone** **RCA**  
**xexternatic** **z** **zenika** **[sfΞir]**

# GenAI

Nantes

“?”



**2 Juin** 2026 à 19h00



4 Rue Voltaire, 44000 Nantes

chez



techtown lonestone RCA  
externatic zenika [sfɪr]

**SHIFT**

**TIME  
TO  
CUSTOM**

27 - 29 Mars  
2026

